

Measurement of $\nu_\mu + \bar{\nu}_\mu$ Charged-Current Single Pion Production Differential Cross Sections on Argon with the MicroBooNE Detector in the NuMI Beam

Analysis Note — Full FHC, RHC, and Combined Exposure

Draft 1.0 — 2026-06-09

Ishanne Pophale, Jarek Nowak

Lancaster University

Abstract

We present a measurement of $\nu_\mu + \bar{\nu}_\mu$ charged-current single charged-pion ($\text{CC}\pi^\pm$) differential cross sections on argon-40 using the full MicroBooNE NuMI exposure, in three configurations: forward horn current (FHC; Runs 1, 2, 4, 5; 8.857×10^{20} protons on target, POT), reverse horn current (RHC; Runs 1, 2, 3, 4a, 4b, 4c; 1.108×10^{21} POT), and their combination (1.994×10^{21} POT). The analysis is blind: results are extracted on Poisson-thrown central-value fake data. Events are selected requiring a reconstructed muon candidate track and exactly one contained pion candidate track emerging from a fiducial neutrino vertex, with zero reconstructed electromagnetic showers and a muon–pion opening angle less than 2.6 rad. In the measured phase space the selection achieves $\approx 55\%$ signal purity and $\approx 15\%$ efficiency (Table 10). Differential cross sections are extracted for each configuration via Wiener-SVD unfolding in five kinematic variables: muon momentum p_μ , pion momentum p_π , muon polar angle $\cos\theta_\mu$, pion polar angle $\cos\theta_\pi$, and the muon–pion opening angle $\theta_{\mu\pi}$. Each configuration is normalised with its own integrated $\nu_\mu + \bar{\nu}_\mu$ flux (FHC 6.81159 , RHC 6.44646 , combined $6.60865 \times 10^{-10} \text{ cm}^{-2} \text{ POT}^{-1}$; new-G4 GEANT4 values above 60 MeV, Sec. 2); the RHC value is lower because its beam is $\bar{\nu}_\mu$ -dominant. The resulting Wiener-SVD integrated cross sections are of order $0.8\text{--}1.2 \times 10^{-38} \text{ cm}^2/\text{Ar}$ (Table 23), with RHC consistent with (or slightly above) FHC; an earlier apparent RHC deficit was a per-run POT normalization bug affecting the multi-file RHC runs, now corrected. The observed yield is $\approx 0.75\times$ the GENIE v3 prediction, consistent with the known over-prediction of $\text{CC}\pi^\pm$ production in current neutrino interaction generators. The extraction is validated by a same-sample self-closure that recovers the truth to within $\approx 6\%$ for every observable, and cross-checked against the D’Agostini iterative-Bayesian method: its flat, iteration-stable integral ($\sim 20\text{--}40\%$ above Wiener-SVD, config-dependent) confirms that the observable-to-observable spread of the Wiener-SVD integrated cross sections is the additional-smearing matrix A_C , not a physical or acceptance effect (Sec. 8.11). All results use the XSECANALYZER framework on the resolution-optimised binning of Sec. 14.

Contents

1	Introduction	1
2	Beam and Neutrino Flux	1
2.1	NuMI beam and flux composition	1
2.2	Data exposure	2
2.3	Integrated flux	2
3	Monte Carlo Samples	3
4	Signal Definition	4
5	Fiducial Volume	6
6	Event Selection	6
6.1	BDT-based particle identification	6
6.2	Cut definitions	7
6.3	Kinematic reconstruction	8
7	Cut-Flow and Selection Performance	8
7.1	Key N−1 distributions	11
7.2	Pion proton-rejection variables	12
7.3	Final-cut data/MC distributions	12
7.4	Background decomposition	14
7.5	Event displays	14
8	Unfolding	15
8.1	The extraction chain, stage by stage	15
8.2	Signal extraction	17
8.3	Response matrix	17
8.4	Regularised unfolding: the bias–variance trade-off	17
8.5	Wiener-SVD unfolding	18
8.6	Reconstructed spectra and response matrices	19
8.7	Reconstruction resolution	22
8.8	The pion-momentum bias and why p_π is the worst-conditioned observable	22
8.9	Model comparison: χ^2 against truth, tune and generators	27
8.10	Proton-tagged χ^2 and the A_C correlation	28
8.11	Cross-check methods	29
9	Full-exposure results: FHC, RHC, and combined	30
9.1	Total flux-averaged cross section (cut-and-count)	34
9.2	Angular variable: $\cos \theta_\mu$ versus θ_μ	35
9.3	Background-control sidebands	40
9.4	Control regions for the proton-tagged subsample	42
10	Proton-tagged subsample: hadronic invariant mass and transverse kinematic imbalance	43
10.1	Selection performance and cut-flow	44
10.2	Additional smearing in the proton-tagged observables	50
11	Systematic Uncertainties	51
12	Known Issues and Required Follow-up	52

13 Analysis Pipeline	52
14 Binning Optimisation, Systematic Breakdown, and Threshold Validation	53
14.1 Momentum-threshold validation	53
14.2 Binning optimisation	53
14.3 Systematic breakdown by source	53
14.4 Anatomy of the flux term: genuine flux vs. unfolding amplification	54
14.5 Multi-run overlay and a resonance-model closure	55
14.6 Consistency with the sibling ν_e CC 1π measurement	56
15 Summary	57
A Data/MC Normalisation	58
A.1 Unfolding covariance ($\sim 26\%$ effect)	58
A.2 EXT (cosmic) subtraction for fake data	58
A.3 Opening-angle binning (review item C4)	59
A.4 Generator normalisation (NuWro vs GENIE)	59
A.5 Software trigger (verified to be a no-op)	59
B NuWro Fake-Data Closure and Cross-Pipeline Validation	59
B.1 Extraction-bias investigation and algorithm equivalence	61
C Observable Binning	62
D Selection Cut Parameter Summary	63

1 Introduction

Charged-current single charged-pion ($\text{CC}\pi^\pm$) production,

$$\nu_\mu(\bar{\nu}_\mu) + \text{Ar} \longrightarrow \mu^\mp + \pi^\pm + X, \quad (1)$$

is among the most frequent neutrino interaction channels at energies of 1–5 GeV and constitutes a significant background for oscillation analyses at the same energy scale, most notably for ν_e appearance measurements at experiments such as NOvA [1] and T2K [2], as well as the future DUNE [3]. Despite its importance, the cross section on nuclear targets remains poorly constrained: major neutrino interaction generators (GENIE v3, NuWro, NEUT) disagree by 20–30% in absolute rate and show significant shape differences for individual kinematic variables.

MicroBooNE is a single-phase liquid argon time-projection chamber (LArTPC) at Fermilab that operated from 2015–2021, providing the world’s highest-statistics ν_μ -on-argon dataset below 5 GeV during its NuMI running period. This note presents a simultaneous measurement of five differential cross sections in the $\text{CC}\pi^\pm$ channel using the full NuMI exposure in FHC, RHC, and combined configurations, employing Pandora pattern recognition [4] and Wiener-SVD unfolding [5].

2 Beam and Neutrino Flux

2.1 NuMI beam and flux composition

The NuMI beam is produced by 120 GeV protons from the Fermilab Main Injector striking a graphite target. Charged secondary hadrons are focused by two magnetic horns whose polarity selects the sign of the focused mesons. In Forward Horn Current (FHC) mode the horns focus positive pions and kaons, giving a beam with predominantly ν_μ content; in Reverse Horn Current (RHC) mode the polarity is reversed to focus negatives, giving a predominantly $\bar{\nu}_\mu$ beam. MicroBooNE is located approximately 110 mrad off-axis from the NuMI beam axis, reducing the peak neutrino energy to the 0.5–2.5 GeV range and narrowing the energy spectrum. Because MicroBooNE sits off-axis and at comparatively low energy, the horn sign-selection is imperfect and both modes carry a sizeable wrong-sign component — most notably RHC, whose wrong-sign ν_μ fraction is large.

The flux composition at the MicroBooNE site is taken from the new-G4 central-value flux histograms (`flux_newg4.root`), integrated over $E_\nu > 60$ MeV to exclude an unphysical sub-threshold pile-up near $E_\nu = 0$ (removed by the flux group’s `RemoveFluxPeak` step; below the reaction threshold and irrelevant to the measurement). The muon-flavour split of these central-value histograms matches the author’s authoritative integrated fluxes exactly (Sec. 2.3):

Table 1: NuMI flux composition at the MicroBooNE site by neutrino type, for the FHC and RHC horn modes (new-G4 central-value flux `flux_newg4.root`, integrated over $E_\nu > 0.1$ GeV). FHC is ν_μ -dominant; RHC is $\bar{\nu}_\mu$ -dominant with a large wrong-sign ν_μ fraction. The combined row weights each mode by its data exposure (FHC 8.857, RHC 11.082×10^{20} POT; Table 2), giving a near-balanced $\nu_\mu/\bar{\nu}_\mu$ mix. The signal is charge-blind ($|\nu_{\text{pdg}}| = 14$), so both ν_μ and $\bar{\nu}_\mu$ contribute in every mode.

Horn mode	ν_μ	$\bar{\nu}_\mu$	ν_e	$\bar{\nu}_e$
FHC	63.4%	34.0%	1.7%	0.9%
RHC	44.2%	53.3%	1.4%	1.2%
Combined	53.0%	44.4%	1.5%	1.1%

Flux corrections from the Package to Predict the FluX (PPFX) [6] are applied per event via the central-value weight `ppfx_cv` stored in the analysis ntuples. (The absolute flux normalisation

used for the cross-section in Sec. 2.3 is still taken from the previous flux histograms; moving that normalisation onto the new-G4 flux is a separate, tracked update.)

2.2 Data exposure

The analysis is *blind*: the real beam-on data are not used. Each configuration is extracted on Poisson-thrown central-value fake data, generated and scaled *per run* to that run’s data POT (Table 2). The FHC measurement uses the full FHC exposure (Runs 1, 2, 4 and 5); RHC uses Runs 1, 2, 3, 4a, 4b and 4c; the combined measurement uses all fourteen overlay files. Every run is normalised independently to its own data exposure, so the prediction carries the data-POT mixture of run periods and the framework is ready for real per-run beam-on data (each fake-data sample is a drop-in replacement for its beam-on counterpart).

Table 2: Run-period exposure accounting (final, full FHC+RHC). “MC POT” is each period’s native simulated exposure; “Data POT” is the per-run exposure to which the (blind, per-run) fake data is scaled. *Each run is normalised independently* by its own factor, $\text{Scale} = (\text{run Data POT})/(\text{run MC POT})$, so the prediction is the data-POT-weighted mixture of run periods — the normalisation a real per-run beam-on dataset produces. This replaces an earlier single per-mode pooling factor: because the per-run MC/data POT ratios vary by up to a factor of three (Scale ranges 0.051–0.141 within FHC alone), a single factor would weight the runs by their MC statistics rather than by the data exposure. The RHC Run 4 row sums the 4a/4b/4c sub-files and Run 3 sums aa–ae; sub-files sharing a run’s data POT are scaled by the group MC total so they pool within that run. The combined measurement applies each run’s own factor. Run 5 uses the reweighted-PPFX ν_μ overlay. The scheme is validated by the FHC closure (unfolded vs. smeared truth $\chi^2 = 1.6/7$, $p = 0.98$).

Run period	MC POT ($\times 10^{20}$)	Data POT ($\times 10^{20}$)	Scale
Run 1 (FHC)	23.282	3.283	0.1410
Run 2 (FHC)	24.934	1.268	0.0509
Run 4 (FHC)	28.335	2.075	0.0732
Run 5 (FHC)	19.300	2.231	0.1156
FHC total	95.85	8.857	<i>per run</i>
Run 1 (RHC)	8.997	0.605	0.0673
Run 2 (RHC)	57.865	2.591	0.0448
Run 3 (RHC, aa–ae)	55.185	5.003	0.0907
Run 4 (RHC, 4a–4c)	32.586	2.883	0.0885
RHC total	154.63	11.082	<i>per run</i>
Combined (FHC+RHC)	250.49	19.939	<i>per run</i>

The full exposure substantially reduces the data-statistical term (Tables 28–30), leaving the measurement flux- and detector-systematic limited. When the analysis is unblinded, the real beam-on POT (and gate counts) replace the data-POT values above.

2.3 Integrated flux

The absolute flux normalisation uses the new-G4 GEANT4 NuMI flux (`flux_newg4.root`), integrated over $E_\nu > 60$ MeV (the flux group’s `RemoveFluxPeak` cut, which removes an unphysical ~ 25 – 30 MeV artifact spike). The signal is charge-blind ($|\nu_{\text{pdg}}| = 14$, covering both $\nu_\mu \rightarrow \mu^- \pi^+$ and $\bar{\nu}_\mu \rightarrow \mu^+ \pi^-$), so the normalisation uses the combined $\nu_\mu + \bar{\nu}_\mu$ flux $\Phi = \Phi_{\nu_\mu} + \Phi_{\bar{\nu}_\mu}$. Because RHC is $\bar{\nu}_\mu$ -dominant with a different absolute flux than FHC, each configuration uses its *own* integrated flux:

Table 3: Authoritative integrated ν_μ and $\bar{\nu}_\mu$ fluxes per POT (new-G4 flux, $E_\nu > 60$ MeV), and the combined $\nu_\mu + \bar{\nu}_\mu$ value used as the per-configuration normalisation. The combined row is the data-POT-weighted average of the two modes (weights 8.857 : 11.082; Table 2). The FHC value 6.81159×10^{-10} is the previously confirmed normalisation; RHC and combined were previously normalised with the FHC flux and are corrected here (raising the RHC cross section by 5.7% and the combined by 3.1%).

Config	Φ_{ν_μ} ($10^{-10} \text{ cm}^{-2} \text{ POT}^{-1}$)	$\Phi_{\bar{\nu}_\mu}$ ($10^{-10} \text{ cm}^{-2} \text{ POT}^{-1}$)	$\Phi_{\nu_\mu + \bar{\nu}_\mu}$	$\bar{\nu}_\mu$ frac.
FHC	4.43515	2.37644	6.81159	34.9%
RHC	2.92348	3.52298	6.44646	54.6%
Combined	3.59497	3.01368	6.60865	45.6%

The mean $\bar{\nu}_\mu$ energy is 0.500 GeV (FHC) and 0.417 GeV (RHC), the lower RHC value reflecting the softer wrong-sign-enhanced spectrum. These fluxes are set per configuration in `integrated_numu_flux_in_FV()` (`include/XSecAnalyzer/FiducialVolume.hh`). The correction relative to the superseded flux (`uboone_numi_flux_histograms.root`, whose combined value 2.3725×10^{-9} was $\sim 3.48\times$ too large and made every extracted cross section $\sim 3.5\times$ too small) was validated four independent ways — standalone NuWro and GENIE on the old vs. corrected flux, the flux author’s numbers, and a GENIE closure — documented in `generator_predictions/`.

The number of argon target nuclei in the fiducial volume is computed from first principles as

$$N_{\text{Ar}} = \frac{\rho_{\text{LAr}} \cdot V_{\text{FV}} \cdot N_A}{A_{\text{Ar}}} \approx 8.71 \times 10^{28}, \quad (2)$$

and the per-configuration normalisation factor is $N_{\text{norm}} = \Phi_{\text{mode}} \cdot N_{\text{POT,mode}} \cdot N_{\text{Ar}}$, evaluated with each mode’s flux and data POT (Table 2).

3 Monte Carlo Samples

Monte Carlo events are generated with GENIE v3 [7] on argon-40 and overlaid with real cosmic-ray data (the “overlay” technique) so that simulated events experience in-situ detector noise and cosmogenic activity. Table 4 summarises the input samples.

Table 4: Input samples (final, full FHC+RHC). “ n ” is the number of overlay files; “MC POT” is the summed native simulated exposure ($\times 10^{20}$). The blind Poisson-thrown fake data stands in for the (unused) beam-on data; the detector variations use the native Run-4 FHC and RHC samples. The per-file lists and POT/trigger values are in the `configs/file_properties_numi_*` tables.

Sample	Composition	n	MC POT
ν_μ overlay (FHC)	Runs 1, 2, 4 (pandora) + Run 5 (reweighted-PPFX)	4	95.85
ν_μ overlay (RHC)	Runs 1, 2, 3, 4a, 4b, 4c (pandora)	10	154.63
Dirt MC	GENIE NuMI dirt overlay (Run 1)	1	16.74
EXT (beam-off)	cosmic data, per run, horn-mode-independent	7	—
Detector variations	Run-4 FHC + Run-4 RHC, CV + 8 knobs each	18	—
Fake data (blind)	Poisson-thrown central value, per configuration	—	—
Beam-on data	blinded — not used until unblinding	—	—

Event weights. Each MC event is weighted by

$$w = \text{Scale}[\text{run}] \times w_{\text{PPFX}} \times w_{\text{tune}}, \quad (3)$$

where $\text{Scale}[\text{run}] = (\text{run Data POT})/(\text{run MC POT})$ is the *per-run* POT factor of Table 2: each run period is scaled independently to its own data exposure, so the prediction reproduces the data-POT mixture of runs rather than the MC-statistics mixture (see the caption of Table 2). Operationally this is exactly what the framework does natively — each overlay file carries its run’s data POT via a matching per-run beam-on (here, per-run fake-data) sample, and is scaled by $\text{run_bnb_pot}/\text{summed_pot}$. w_{PPFX} is the PPFX central-value flux correction (`ppfx_cv`) and w_{tune} is the GENIE/GEANT4 reweighting correction (`weightTune`). Events with non-finite, zero, or unphysically large weights ($w > 30$) are assigned unit weight. At unblinding the per-run fake-data samples are replaced by the corresponding per-run beam-on files (same runs, same POT), so the normalisation is already correct.

EXT scaling. The same combined beam-off (cosmics) sample is attached to every run and scaled by that run’s beam-on/beam-off trigger ratio, $\text{bnb_trigs}[\text{run}]/\text{ext_gates}$. Because the beam-on/beam-off ratio is run-independent, the per-run contributions sum to the mode-total scale; for the full FHC exposure,

$$\text{Scale}_{\text{EXT}}^{\text{FHC}} = \frac{\sum_{\text{run}} \text{bnb_trigs}[\text{run}]}{\text{ext_gates}} = \frac{22\,667\,109}{3\,821\,593} = 5.93. \quad (4)$$

(The framework requires each run carrying a beam-on sample to carry a beam-off sample as well; the single beam-off file is therefore attached per run through symlinks so that run bookkeeping is complete without duplicating the data.) The uncertainty on this ratio is a leading systematic on the EXT background.

Software trigger. Data and EXT files are pre-filtered to require the software trigger (`swtrig == 1`). MC events failing `swtrig = 0` are removed at the NeutrinoSlice cut stage. In practice this has no measurable effect at the current working point: all MC events passing the topological cut are verified to have `swtrig == 1` (46 810/46 810 events confirmed numerically).

4 Signal Definition

The signal is defined at generator level and is used only to construct the selection efficiency and response matrix. It is *never* applied as a reconstruction cut. The definition is *charge-blind*: both $\nu_{\mu} \rightarrow \mu^{-}\pi^{+}$ and $\bar{\nu}_{\mu} \rightarrow \mu^{+}\pi^{-}$ are treated as signal on the same footing.

Neutrino flavour

$$|\nu \text{ PDG}| = 14 \text{ } (\nu_{\mu} \text{ or } \bar{\nu}_{\mu}).$$

Interaction type

Charged-current (`ccnc == 0`).

True vertex

Inside the fiducial volume (Section 5).

Primary multiplicity

Exactly one muon ($|\text{PDG}| = 13$), exactly one charged pion ($|\text{PDG}| = 211$), zero neutral pions ($|\text{PDG}| = 111$), zero kaons. Any number of protons is allowed.

Kinematic thresholds

$p_{\mu} > 0.15 \text{ GeV}/c$, $p_{\pi} > 0.175 \text{ GeV}/c$, and $\theta_{\mu\pi} < 2.6 \text{ rad}$. These define the *measured phase space*, chosen from the efficiency turn-on curves (Section 7). The momentum efficiency turns on sharply at $\sim 0.15\text{--}0.20 \text{ GeV}/c$ (set by the 10 cm muon and 20 cm pion track-length cuts); below that it is only $\sim 1\text{--}3\%$ with a strongly model-dependent correction. The opening-angle

upper bound (2.6 rad) excludes the large-angle region, which is dominated by cosmic and dirt background ($\sim 7\%$ purity in the $\theta_{\mu\pi} > 2.513$ rad bin). No *lower* opening-angle cut is applied: the small-angle region is high-purity ($\sim 58\%$), so cutting it would only remove signal (it is low-efficiency but not low-purity). The same cut is applied at reco level so the signal and selection acceptances match. The first p_μ/p_π bins and the $\theta_{\mu\pi}$ upper edge start/end at these thresholds.

(Earlier iterations used $p_\mu, p_\pi > 0.1$ GeV/ c with the full $0-\pi$ angle; that left near-dead efficiency bins at low momentum and at the large-angle extreme. The phase space above raises the efficiency from 12.1% to 15.6% at $\sim 55\%$ purity — see Section 7.)

Pion threshold and the sibling ν_e measurement. The companion MicroBooNE NuMI ν_e CC 1π analysis defines its pion signal by kinetic energy, $\text{KE}_\pi > 40$ MeV, equivalent to $p_\pi > 0.113$ GeV/ c . Our $p_\pi > 0.175$ GeV/ c ($\text{KE}_\pi > 84$ MeV) is therefore a higher threshold. Lowering it toward the ν_e value is possible — the signal definition is a truth-level choice and the unfolding maps reco to truth — but the pion momentum reconstruction (track range under the muon hypothesis, mass-corrected) degrades rapidly below 0.175 GeV/ c :

Table 5: Pion momentum reconstruction versus true p_π (correctly reconstructed signal pions). Below the 0.175 GeV/ c threshold the bias and resolution blow up and the reconstructed fraction halves.

True p_π [GeV/ c]	Bias $\langle (p^{\text{reco}} - p^{\text{true}})/p^{\text{true}} \rangle$	RMS resolution	Reco fraction
0.080–0.113	+131%	73%	3.4%
0.113–0.150	+66%	56%	3.4%
0.150–0.175	+30%	51%	4.3%
0.175–0.250 (threshold)	−8%	31%	10.0%
0.250–0.400	−30%	21%	10.4%

At the threshold the resolution is 31%; in the $[0.113, 0.175]$ GeV/ c region it is 51–56% with a +30–66% bias and only ~ 3 –4% reconstructability. Extending the measurement to the ν_e threshold would thus add bins carrying large, strongly model-dependent unfolding systematics. For consistency with the ν_e measurement the threshold could be lowered to 0.113 GeV/ c with the $[0.113, 0.175]$ GeV/ c region kept in a single coarse bin to contain those systematics, or both phase spaces reported; this is a systematics-versus-cross-analysis-consistency trade rather than a reconstruction limitation.

The low-momentum bin, added. We have implemented this extension for p_π : a $[0.113, 0.175]$ GeV/ c bin is prepended to the pion-momentum binning, lowering the measured phase space to the ν_e threshold. The truth-level signal below 0.175 GeV/ c is recovered from the stored signal-component branches (the reconstruction reproduces `MC_Signal` above 0.175 GeV/ c to 0.4%, the residual being the unstored kaon veto); 3323 true signal events populate the new bin. Its selection efficiency is 2.2% — an order of magnitude below the 15–17% of the standard bins (Table 6), as anticipated from the reconstruction turn-on. Despite this, the Wiener-SVD extraction remains well behaved with the bin included: the M_A^{RES} fake-data closure gives $\chi^2/\text{ndf} = 1.41/6$ ($p = 0.97$) and the flux term is unchanged (17% reco, 18.4% unfolded, amplification 1.08 $\times$). The bin therefore adds to the measurement, at the price of a large per-bin uncertainty driven by the low efficiency.

The generator predictions have been regenerated on the six-bin scheme (the ppi histogram uses $p_\pi > 0.113$ while the other observables retain $p_\pi > 0.175$, mirroring the analysis). GENIE, NuWro and NEUT all provide the low bin; GiBUU could not be re-binned as its generator-level event record was not retained, so it is omitted from the extended p_π comparison. Against the

(fake) data the generators give $\chi^2/\text{ndf} = 11.1/6$ (GENIE, $p = 0.09$), $7.0/6$ (NuWro, $p = 0.33$) and $6.3/6$ (NEUT, $p = 0.39$): NEUT and NuWro describe the extended spectrum well, while GENIE is mildly disfavoured, consistent with the known GENIE over-prediction of low-momentum $\text{CC}\pi^\pm$ production.

Table 6: Selection efficiency per true p_π bin with the low-momentum bin added. The new $[0.113, 0.175]$ GeV/c bin sits at 2.2%, an order of magnitude below the standard bins.

True p_π [GeV/c]	N_{true}	N_{reco}	Efficiency
0.113–0.175 (new)	3323	74	2.2%
0.175–0.250	3368	526	15.6%
0.250–0.320	2190	367	16.8%
0.320–0.420	2831	479	16.9%
0.420–0.550	2499	410	16.4%
0.550–1.000	2933	429	14.6%

5 Fiducial Volume

The fiducial volume (FV) is defined in `selection/FV_new.h`:

$$x \in [10.0, 246.0] \text{ cm}, \quad y \in [-101.0, +101.0] \text{ cm}, \quad z \in [10.0, 986.0] \text{ cm}, \quad (5)$$

with the dead-wire region $z \in [675.1, 775.1]$ cm excluded. The total FV volume is

$$\begin{aligned} V_{\text{FV}} &= (246 - 10) \times 202 \times (976 - 100) \text{ cm}^3 \\ &= 236 \times 202 \times 876 \text{ cm}^3 \approx 4.17 \times 10^7 \text{ cm}^3. \end{aligned} \quad (6)$$

A wider *containment* region (2 cm inset from the physical TPC walls) is used to require that the pion candidate track endpoint is contained:

$$x \in [2.0, 254.35] \text{ cm}, \quad y \in [-113.53, +107.47] \text{ cm}, \quad z \in [2.1, 1034.9] \text{ cm}. \quad (7)$$

6 Event Selection

The event selection is implemented in the framework selection class `src/selections/CC1mu1piXp.cxx` (which produces the response matrices and spectra for all results in this note); the standalone `selection/ccpi_selection.C` pipeline used for the cross-pipeline validation of Appendix B implements the byte-identical cut sequence. The selection comprises 11 sequential cuts.

6.1 BDT-based particle identification

Three gradient-boosted decision trees (TMVA), trained on the MicroBooNE overlay MC, supply the calorimetric particle identification used in the muon- and pion-candidate cuts. All three share a common set of per-track inputs: the three Bragg-peak dE/dx likelihoods under the proton, muon and MIP hypotheses (`trk_bragg_{p,mu,mip}_v`), the log-likelihood-ratio PID score (`trk_llr_pid_score_v`), the Pandora track-vs-shower score (`trk_score_v`), and the space-charge-corrected track end-point coordinates (`trk_sce_end_{x,y,z}_v`):

- **MIP BDT** (`bdt_mip`; weights `dataset_MIP_BDT_no_len`): separates minimum-ionising tracks (muons, pions) from stopping protons. Both the muon and the pion candidate are required to have `bdt_mip` > -0.1 (Cuts 4–5).
- **Pion BDT** (`bdt_pi`; weights `dataset_pion_BDT_no_len`): a dedicated pion-versus-proton discriminant applied to the pion candidate, `bdt_pi` > -0.1 (Cut 5).

- **Muon BDT** (weights `dataset_muon_BDT`): the same inputs plus the track length (`trk_len_v`). Rather than a threshold, it *ranks* the muon-candidate tracks — the highest-scoring track is chosen as the muon.

The two candidate-PID BDTs deliberately omit track length so the score does not bias the momentum spectra. The weight files are resolved from `$CC1MU1PIXP_BDT_WEIGHTS_DIR`, falling back to the in-tree `booster_decision_tree/` copy. A fourth score, `bdt_cosmic`, is available for cosmic rejection but is disabled by default (no benefit in an A/B test; the Pandora topological score already provides the cosmic rejection at Cut 3).

6.2 Cut definitions

Cut 0. EventsInTrueFV. All events. Reference denominator for truth-level studies.

Cut 1. NeutrinoSlice. ≥ 1 reconstructed track (`ntracks > 0`), `bdt_cosmic > 0.0` (currently disabled; no benefit found in A/B test), and `swtrig == 1`.

Cut 2. InFiducialVol. The SCE-corrected reconstructed neutrino vertex lies inside the FV.

Cut 3. Topological. `topological_score > 0.67`. This Pandora-produced multivariate score distinguishes cosmic-ray topologies from neutrino-like topologies and provides the primary cosmic rejection: EXT events are reduced by a factor of ~ 33 between Cut 2 and Cut 3.

Cut 4. MuonCandidate. At least one primary PFP track satisfies the muon PID criteria of Table 7. The longest qualifying track is the muon candidate.

Cut 5. ContainedPion. Exactly one primary track (other than the muon candidate) satisfies the pion PID criteria of Table 8 and has its endpoint inside the containment volume. The track with the highest LLR PID score is the pion candidate.

Cut 6. MuonIn3Planes. The muon candidate has ≥ 1 hit on each of the U, V, Y planes.

Cut 7. PionIn3Planes. The pion candidate has ≥ 1 hit on each of the U, V, Y planes.

Cut 8. ShowerCut. Zero primary-level EM showers, unless exactly one shower is present with fewer than 50 Y-plane hits and a vertex distance > 10 cm (consistent with a delta-ray or low-energy nuclear photon).

Cut 9. OpeningAngle. The reconstructed μ - π opening angle satisfies $\theta_{\mu\pi} < 2.6$ rad ($\approx 149^\circ$), rejecting the large-angle cosmic/dirt background. No lower bound is applied (the small-angle region is high-purity). The same cut is imposed on the signal definition (Section 4).

Cut 10.

Nonprotons (final cut). Fewer than 4 primary tracks have LLR PID > 0.1 (non-proton-like), and the total number of primary reconstructed tracks is less than 5.

Table 7: Muon candidate PID thresholds (applied at Cut 4).

Variable	Threshold	Description
<code>trk_score_v</code>	> 0.80	Track-like PFP
Vertex distance	≤ 4.0 cm	Attached to ν vertex
Track length	> 10.0 cm	Minimum track length
LLR PID (<code>trk_llr_pid_score_v</code>)	> 0.20	Muon-like
MIP BDT (<code>bdt_mip</code>)	≥ -0.10	Minimum-ionising-particle

Table 8: Pion candidate PID thresholds (applied at Cut 5). Endpoint containment requires the track endpoint inside the containment volume. `trk_bragg_pion_v`: pion-hypothesis Bragg-peak likelihood score (rejects stopping protons).

Variable	Threshold	Description
LLR PID	> 0.10	Pion / non-proton-like
Vertex distance	< 4.0 cm	Attached to ν vertex
Track length	> 20.0 cm	Minimum track length
MIP BDT (<code>bdt_mip</code>)	> -0.10	MIP-like
Pion BDT (<code>bdt_pi</code>)	> -0.10	Pion-like (vs. proton)
Bragg-pion score	≥ 0.08	Proton rejection
Endpoint contained	yes	Inside containment volume

6.3 Kinematic reconstruction


Muon momentum. A *hybrid* estimator, switched on containment of the muon-candidate track: the range-based momentum `trk_range_muon_mom_v` is used when the track ends inside the detector, and the multiple-Coulomb-scattering momentum `trk_mcs_muon_mom_v` when it exits (range is meaningless for an exiting track). The reported p_μ is therefore `candidate_muon_mom_reco`, not either branch alone. In the selected FHC sample the split is 37% contained (range) and 63% exiting (MCS), i.e. the measurement is *MCS-dominated*: the NuMI muons are energetic and the MicroBooNE TPC is small, so most of them leave it. This matters for the p_μ resolution and for the detector systematic, since the range and MCS estimators carry different uncertainties.

Pion momentum. Taken from the *muon-hypothesis* range momentum (`trk_range_muon_mom_v`). The muon mass (105.7 MeV) is close to the pion mass (139.6 MeV), so the range-to-momentum conversion is a good proxy, and — crucially — it is a *momentum*, directly comparable to the true p_π it is binned against. An earlier version of this analysis used the proton-hypothesis range *kinetic energy* (`trk_energy_proton_v`) instead. That is a different physical quantity from the momentum used for the true axis with the same bin edges, so the migration matrix was systematically shifted rather than near-diagonal: a true 0.2 GeV/ c pion gives $KE \approx 0.104$ GeV, two bins low, and a true 0.175 GeV/ c pion falls below the first reco bin entirely. Range momentum is meaningful only for contained tracks, which the containment requirement in the selection guarantees.

Angles. Polar angles are defined with respect to the detector z -axis: $\cos \theta_\mu = p_z^\mu / |\mathbf{p}^\mu|$, $\cos \theta_\pi = p_z^\pi / |\mathbf{p}^\pi|$. The opening angle $\theta_{\mu\pi} = \arccos(\hat{u}_\mu \cdot \hat{u}_\pi)$ is computed from reconstructed track directions.

7 Cut-Flow and Selection Performance

Figure 1 shows the full-exposure cut-flow yields for all three configurations as stacked plots, and Table 9 gives the detailed per-category, per-cut breakdown for each.



(a) FHC.

(b) RHC.

(c) Combined.

Figure 1: Cut-flow yields at *full exposure* for FHC, RHC and combined (log scale; MC-prediction stack: signal blue, background orange, EXT grey, dirt green; colourblind-safe Okabe-Ito), signal and background POT-scaled per mode. Data is blind, so this shows the MC-prediction stack; the detailed per-category breakdown is in Table 9.

Table 9: Cut-flow yields at *full exposure* for FHC, RHC and combined (MC prediction; data is blind, so Pred = Data by construction for the Poisson-thrown fake data). Signal and “Bkg” (all non-signal MC: OOFV, NC, CC1 π^0 , CC0 π , CCN π , CCK, CCothers) are from the ν_μ overlays, POT-scaled *per run* (each run by its own $D_{\text{run}}/MC_{\text{run}}$, Table 2); EXT (beam-off cosmic) is scaled by the (per-run summed) data/EXT trigger ratio times the 2% NuMI beam-occupancy factor (FHC 0.98×5.93 , RHC 0.98×6.16 , combined 0.98×12.09); dirt from the GENIE NuMI dirt overlay. Pred. = Signal + Bkg + EXT + Dirt. Eff. is the cumulative signal efficiency relative to the generated signal (the None stage, where all true signal is counted); Pur. = Signal/Pred.

Cut	Signal	Bkg MC	EXT	Dirt	Pred.	Eff. [%]	Pur. [%]
<i>FHC (Runs 1,2,4,5; 8.857×10^{20} POT)</i>							
None	5 580	279 180	23 921 822	36 142	24 242 724	100.0	0.02
InFiducialVol	4 632	62 146	1 020 804	1 890	1 089 473	83.0	0.43
Topological	3 501	28 040	34 481	204	66 225	62.7	5.3
MuonCandidate	3 143	18 980	15 206	137	37 466	56.3	8.4
ContainedPion	1 395	2 094	1 348	25	4 862	25.0	28.7
MuonIn3Planes	1 336	1 975	1 122	17	4 450	23.9	30.0
PionIn3Planes	1 239	1 770	854	7.0	3 871	22.2	32.0
ShowerCut	1 048	898	558	6.2	2 510	18.8	41.7
OpeningAngle	1 040	813	186	2.0	2 041	18.6	51.0
Nonprotons	974	643	128	1.3	1 746	17.5	55.8
<i>RHC (Runs 1,2,3,4a,4b,4c; 1.108×10^{21} POT)</i>							
None	6 130	304 347	24 837 750	28 032	25 176 259	100.0	0.02
InFiducialVol	5 124	69 102	1 059 889	1 466	1 135 581	83.6	0.45
Topological	3 960	31 629	35 801	158	71 548	64.6	5.5
MuonCandidate	3 559	21 443	15 788	106	40 896	58.1	8.7
ContainedPion	1 577	2 649	1 400	19	5 646	25.7	27.9
MuonIn3Planes	1 512	2 500	1 165	13	5 189	24.7	29.1
PionIn3Planes	1 409	2 258	887	5.5	4 559	23.0	30.9
ShowerCut	1 170	1 061	579	4.8	2 815	19.1	41.6
OpeningAngle	1 162	965	193	1.5	2 321	19.0	50.1
Nonprotons	1 081	746	133	1.0	1 961	17.6	55.1
<i>Combined FHC+RHC (1.994×10^{21} POT)</i>							
None	11 710	583 527	48 759 975	64 174	49 419 386	100.0	0.02
InFiducialVol	9 756	131 248	2 080 711	3 357	2 225 072	83.3	0.44
Topological	7 461	59 669	70 282	362	137 773	63.7	5.4
MuonCandidate	6 702	40 423	30 994	243	78 362	57.2	8.6
ContainedPion	2 972	4 743	2 749	44	10 508	25.4	28.3
MuonIn3Planes	2 847	4 475	2 287	30	9 639	24.3	29.5
PionIn3Planes	2 648	4 028	1 742	12	8 430	22.6	31.4
ShowerCut	2 218	1 959	1 137	11	5 325	18.9	41.7
OpeningAngle	2 202	1 778	379	3.5	4 363	18.8	50.5
Nonprotons	2 055	1 389	261	2.3	3 707	17.5	55.4

Summary at the final cut (reference phase space, $p_\mu, p_\pi > 0.1 \text{ GeV}/c$).

- Signal MC (POT-weighted): 324.7 events
- Total background: $242.6 = 198.5 (\nu\text{-bkg}) + 40.1 (\text{EXT}) + 3.8 (\text{dirt})$
- True signal generated (efficiency denominator): 2577–2684 events (observable-dependent)
- **Signal purity: 57.2%**
- **Signal efficiency: 12.1–12.6%**
- **Measured phase space (reported result): efficiency 15.6%, purity 54.9%** (reference phase-space study: selected signal 311.9, generated 1999, background 253.7; Table 10)

Per-run scaling. The yields above and in Table 9 use the per-run POT factors of Table 2 (each

run scaled by its own $D_{\text{run}}/MC_{\text{run}}$), and the EXT by the per-run trigger ratios summing to 5.93 (FHC). Relative to the earlier single per-mode factor this shifts the FHC final-cut signal by -1.9% ($993 \rightarrow 974$ events) and rescales the beam-off by the updated ratio; because the shift is a small, nearly uniform re-weighting of the run mixture, the efficiency and purity move by $< 0.5\%$ and the quoted percentages are unchanged at this precision. The per-observable figures are finalised with the full per-run re-extraction; the FHC p_μ closure (unfolded vs. smeared truth $\chi^2 = 1.6/7$, $p = 0.98$) confirms the scheme is unbiased.

Measured phase space and efficiency turn-on. The reference selection above has near-dead efficiency at low momentum and at the large opening-angle extreme: from the fine truth distributions, the per-bin efficiency is only $\sim 2\text{--}5\%$ for $p_\mu, p_\pi < 0.15 \text{ GeV}/c$ (the 10/20 cm track-length cuts) and the region above the reco $\theta_{\mu\pi} = 2.6$ rad cut is dominated by cosmics ($\sim 7\%$ purity). The measured phase space $p_\mu > 0.15$, $p_\pi > 0.175 \text{ GeV}/c$, $\theta_{\mu\pi} < 2.6$ rad (Section 4) raises and flattens the efficiency to $\sim 13\text{--}20\%$ (Fig. 18). Table 10 compares the selection across the iterations. Two lessons: (i) the 2.6 rad *upper* cut is essential — dropping it admits $3.4\times$ more cosmic background and cuts the purity to 45%; (ii) a *lower* opening-angle cut is counterproductive — the small-angle region is high-purity ($\sim 58\%$), so cutting it lowers the selected signal without improving purity, which is why no lower cut is used.

Table 10: Selection performance across the phase-space iterations (NuWro fake data). “no upper cut” drops the $\theta_{\mu\pi} < 2.6$ cut; “[0.314, 2.6]” adds a lower OA cut. The final choice keeps the 2.6 rad upper cut with no lower cut.

	Reference ($p > 0.1$)	no upper cut ($p > .15/.175$)	[0.314, 2.6] (lower cut)	Final ($\theta < 2.6$)
Generated signal	2684	2055	1916	1999
Selected signal	324.7	320.3	301.8	311.9
Efficiency	12.1%	15.6%	15.8%	15.6%
EXT (cosmic)	40.1	136.2	38.0	40.1
Total bkg	243.6	385.4	252.3	256.4
Purity	57.1%	45.4%	54.5%	54.9%
sig $\times$ purity	185.6	145.4	164.4	171.2

7.1 Key N−1 distributions

Figure 2 shows the N−1 distributions (all cuts applied except the one under study) for the two cuts instrumented per event — the topological score (Cut 3) and the $\mu\text{--}\pi$ opening angle (Cut 9) — for all three configurations. The shapes are consistent across horn modes, as expected since the selection is identical; RHC and combined simply carry the larger exposure. Signal (blue) and background (orange) are stacked and the dashed line marks the applied threshold. The remaining cut variables are shown at the final cut in Sec. 7.3.




Figure 2: N-1 distributions for the topological score (Cut 3, top) and μ - π opening angle (Cut 9, bottom), for FHC, RHC and combined. Signal (blue) and background (orange) stacked; dashed line = applied threshold. Colourblind-safe Okabe-Ito palette.

7.2 Pion proton-rejection variables

Figure 3 shows the two main proton-rejection variables on the pion candidate, evaluated on MC truth-labelled pions ($\pi^\pm$) and protons at the final cut stage, used to tune the working-point thresholds.




Figure 3: Pion-candidate PID distributions for true $\pi^\pm$ (red) and true protons (blue) at the final cut stage. The Bragg-pion score provides the primary proton rejection; χ_p^2 is studied as cross-check.

7.3 Final-cut data/MC distributions

Figures 4 and 5 show the candidate-level distributions at the final selection cut for the four instrumented PID and track-length variables, for all three configurations. Signal (blue) and

background (orange) are stacked (colourblind-safe Okabe-Ito palette).




Figure 4: Muon-candidate LLR PID (top) and track length (bottom) at the final cut, for FHC, RHC and combined. Signal (blue) / background (orange) stacked.




Figure 5: Pion-candidate LLR PID (top) and track length (bottom) at the final cut, for FHC, RHC and combined. Signal (blue) / background (orange) stacked.




Figure 6: Data/MC comparison for topological variables at the final cut.

7.4 Background decomposition

Table 11 shows the truth-level background breakdown at the final cut.

Table 11: Background composition at the final cut (truth-level MC categories, POT-weighted).


Category	Events	Fraction
CC0 π (proton mis-ID as pion)	~ 82	$\sim 41\%$
CCN π (≥ 2 pion events)	~ 67	$\sim 34\%$
OOFV (vertex outside FV)	~ 18	$\sim 9\%$
NC (neutral current)	~ 14	$\sim 7\%$
CC1 π^0	~ 8	$\sim 4\%$
CC other / CC kaon	~ 6	$\sim 3\%$
ν_e / other flavour	~ 4	$\sim 2\%$
EXT (cosmic)	40.1	—
Dirt	3.8	—

The dominant background ($\sim 41\%$) is CC0 π interactions in which a proton is reconstructed as the pion candidate. This mis-identification was verified directly by 2D event displays (truth-labelled by backtracked PDG): the pion candidate track shows a green (proton) truth colour with a stopping Bragg peak, while the genuine muon track is correctly identified.

At the current working point (`trk_bragg_pion_v` ≥ 0.08), the pion candidate is truth-matched to a proton in $\approx 18\%$ of cases. Tighter Bragg-score thresholds (0.20–0.30) were studied but shown not to improve the figure of merit: the pion-to-proton ratio in $CC\pi^\pm$ signal is $\approx 4 : 1$, so rejecting protons removes genuine pions at nearly the same rate.

7.5 Event displays

Figure 7 shows representative 2D truth-labelled event displays for the two dominant backgrounds and the EXT cosmic sample. Tracks are coloured by backtracked truth PDG: μ (blue), $\pi^\pm$ (red), p (green), cosmic/unmatched (grey). The muon candidate is drawn solid-thick; the pion candidate dashed-thick; the reconstructed ν vertex is shown as a star.



(a) CC0 π background: a proton (green, dashed) is reconstructed as the pion candidate. (b) CCN π background: two genuine charged pions present; one is selected as the pion candidate. (c) EXT cosmic event: all tracks are grey (no truth match), faking a neutrino vertex topology.

Figure 7: Truth-labelled 2D event displays (ZX projection) for representative background events at the final cut. Three views (ZX top, ZY side, XY face) are generated per event by `selection/event_display.C`.

8 Unfolding

8.1 The extraction chain, stage by stage

Before the formalism, it is useful to see the measurement assembled in the order it is actually performed. Each stage is shown for all three beam configurations in the flat reco-bin space in which the unfolding operates (every analysis slice concatenated).

Stage 1 — selected events, POT-scaled (Fig. 8). The starting point is the selected reconstructed spectrum, scaled to the full exposure of each configuration with the per-run POT factors of §A. The Monte Carlo is split into signal and background. At the final cut the purity is $\simeq 55\%$ (Table 9), so roughly half of the selected events have to be removed before anything can be interpreted as a signal rate. Because the analysis is blind, the points are the Poisson-thrown fake data rather than beam-on data.

Stage 2 — background subtraction (Fig. 9). The estimated background is subtracted bin by bin, leaving the background-subtracted data alongside the central-value MC signal. The neutrino background is taken from the simulation, the cosmic component from beam-off (EXT) data, and the dirt component from the dedicated overlay; the four control regions of §9.3 are what validate this step. This subtracted spectrum, not the raw selection, is the input to the unfolding.

Stage 3 — unfolding. The subtracted spectrum still folds in the detector response and the selection efficiency, encoded in the response matrix of §8.3. With 50–99% bin migration a direct inversion is unusable (§8.4), so the measurement is regularised with Wiener–SVD. The price is the additional-smearing matrix A_C : what is recovered is $A_C \times (\text{true signal})$, and *every* model must therefore be multiplied by A_C before being compared with the data.

Stage 4 — comparison with models. The unfolded result is then compared with the MicroBooNE tune and the four external generators, all A_C -smeared into the same space (§9). It is worth separating two effects that are easily conflated. A_C suppresses the integral of each prediction, but it does so by a *common* factor: for FHC p_μ it is 0.64, 0.64, 0.59 and 0.67 for GENIE, GiBUU, NEUT and NuWro respectively. The truth-level totals of those same generators span 0.79 to $1.26 \times 10^{-38} \text{ cm}^2/\text{Ar}$, a factor 1.6. A common $\simeq 0.6$ factor cannot produce a $1.6\times$ spread, so the differences between the generators are genuine model differences and not an

artefact of the regularisation. The cut-and-count measurement of §9.1 provides the independent check, since it involves no unfolding and hence no A_C at all.




Figure 8: Stage 1, FHC: selected reconstructed spectra scaled to the full exposure, with the MC split into signal and background. Points are the Poisson-thrown fake data. The RHC and combined equivalents are Figs. 10 and 12.




Figure 9: Stage 2, FHC: background-subtracted data compared with the central-value MC signal — the input to the unfolding. RHC and combined: Figs. 11 and 13.




Figure 10: Stage 1, RHC: selected reconstructed spectra scaled to the full RHC exposure, MC split into signal and background.




Figure 11: Stage 2, RHC: background-subtracted data compared with the central-value MC signal.




Figure 12: Stage 1, combined: selected reconstructed spectra at the combined FHC+RHC exposure.




Figure 13: Stage 2, combined: background-subtracted data compared with the central-value MC signal.

8.2 Signal extraction

After the event selection, the background-subtracted reco-space signal spectrum for observable x is:

$$d_i^{\text{sig}} = d_i^{\text{data}} - b_i^{\nu\text{-bkg}} - b_i^{\text{EXT}} - b_i^{\text{dirt}}, \quad (8)$$

where i runs over reconstructed bins.

8.3 Response matrix

A 2D response matrix A_{ij} (true bin j , reco bin i) is built from selected signal MC using the same binning on both axes. Each column is normalised to unit sum so that $A_{ij} = P(\text{reco in bin } i | \text{true in bin } j)$. The efficiency vector $\varepsilon_j = \sum_i A_{ij}$ is the ratio of selected (`h_smeared_sel`) to generated (`h_true_gen`) signal events per true bin.

8.4 Regularised unfolding: the bias–variance trade-off

Unfolding inverts the folding relation $\vec{m} = A\vec{s} + \vec{b}$ to estimate the true signal spectrum $\vec{s}$ from the measured $\vec{m}$. The naive inverse $A^{-1}(\vec{m} - \vec{b})$ is unbiased but useless: bin migration makes A nearly singular, so its inverse amplifies statistical fluctuations into large, anticorrelated bin-to-bin oscillations. Every practical method therefore *regularises* — trading a small, controlled bias for a large reduction in variance. This analysis uses two complementary methods (a fuller pedagogical treatment is given in the companion note `report/unfolding_note.tex`).

Wiener-SVD (primary). Diagonalising the response with an SVD, each mode k is weighted by the Wiener filter $W_k = \sigma_k^2 / (\sigma_k^2 + N_k)$ — the minimum-mean-squared-error weighting, $\simeq 1$ for well-measured modes and $\rightarrow 0$ for noise-dominated ones — so no regularisation strength is tuned by hand. The price is a *known* linear bias: the unfolded data estimates not $\vec{s}$ but the smeared spectrum $A_C\vec{s}$, where the additional-smearing matrix A_C is published so that every theory prediction is smeared by the same A_C before comparison.

D’Agostini iterative Bayesian (cross-check). Bayes’ theorem inverts the migration, $P(C_j|E_i) = R_{ij}\pi_j / \sum_l R_{il}\pi_l$, and the unfolded estimate is $\hat{s}_j = \varepsilon_j^{-1} \sum_i P(C_j|E_i)(m_i - b_i)$. The MC prior π_j is updated with the result and iterated; the *number of iterations* is the regularisation knob (few $\Rightarrow$ prior-biased, many $\Rightarrow$ noisy). It carries no A_C , so its integral is essentially unbiased and theory is compared raw.

Table 12: Trade-offs of the two unfolding methods. Neither is “more correct”; they make complementary choices about how to trade bias for variance and how to expose the residual bias.

	Wiener-SVD	D’Agostini (iter. Bayes)
Regularisation	Wiener filter (automatic, from S/N)	Iteration count n (chosen)
Residual bias	Explicit A_C , published	Implicit; set by early stopping
Result integral	Suppressed, observable-dependent	Essentially unbiased (physical)
Theory comparison	Smear theory by A_C	Raw theory
Positivity	Not guaranteed	Guaranteed
Uncertainties	Closed-form, linear	Correlated across iters.; needs care

Wiener-SVD is the primary result — for its objective regularisation and its explicit, reportable bias A_C ; D’Agostini serves as the quantitative cross-check (Sec. 8.11), where its flat, iteration-stable integral, $\sim 20\text{--}40\%$ above Wiener-SVD (config-dependent), confirms that the observable-to-observable spread of the Wiener-SVD integrated cross sections (Table 23) is the A_C smearing rather than a physical or acceptance effect.

8.5 Wiener-SVD unfolding

The primary method is Wiener-SVD [5], implemented in `xsec/WienerSVD.h` (`RunWienerSVD_FW`) as a faithful replica of the official `XSecAnalyzer WienerSVDUnfolder`: Cholesky whitening of the data covariance, regularisation inside the SVD of RC^{-1} , and the BNL filter (Eq. 3.24 of Ref. [5]). The method applies a Wiener filter in the SVD basis, providing adaptive regularisation whose strength is set by the data covariance. Because the custom pipeline does not throw systematic universes, the covariance is imported from the framework (the prediction covariance scaled to the custom prediction, plus the custom data Poisson; Section A); a stat-only covariance under-regularises and over-corrects the integral by $\sim 26\%$.

The differential cross section is:

$$\left. \frac{d\sigma}{dx} \right|_i = \frac{\hat{N}_i^{\text{unfold}}}{\varepsilon_i \cdot \Delta x_i \cdot \Phi_{\text{total}} \cdot N_{\text{Ar}}}, \quad (9)$$

where $\hat{N}_i$ is the Wiener-SVD unfolded signal count in true bin i , $\Phi_{\text{total}} = \Phi_{\text{per POT}} \times T_{\text{POT}}$ is the total integrated flux, and N_{Ar} is the number of argon target nuclei in the FV. Equivalently, the events-to-cross-section conversion is $\text{conv} = N_{\text{Ar}} \Phi / 10^{-38}$ (so $\sigma_i = \hat{N}_i / (\varepsilon_i \text{conv})$), with the authoritative flux constant $\phi = 6.81159 \times 10^{-10} \text{ cm}^{-2}/\text{POT}$ ($E > 60 \text{ MeV}$; Sec. 2). At the fake-data POT the inputs are:

Quantity	Value
N_{Ar} (FV, dead-region excluded)	8.710×10^{29}
Flux constant ϕ ($E > 60 \text{ MeV}$)	$6.81159 \times 10^{-10} \text{ cm}^{-2}/\text{POT}$
POT	7.630913×10^{20}
Integrated flux $\Phi = \phi \times \text{POT}$	$5.198 \times 10^{11} \text{ cm}^{-2}$
$\text{conv} = N_{\text{Ar}} \Phi / 10^{-38}$	4.527×10^3

An integrated cross section is computed as the width-weighted sum,

$$\sigma_{\text{int}} = \sum_i \left. \frac{d\sigma}{dx} \right|_i \Delta x_i, \quad (10)$$

and serves as a closure check: all five observables must agree.

8.6 Reconstructed spectra and response matrices

Figure 14 shows the background-subtracted reco-space signal spectra for all five observables, overlaid with the MC signal prediction and the individual background components.




Figure 14: Reconstructed-space spectra at the final cut for all five observables, FHC full exposure (Runs 1+2+4+5): Poisson-thrown central-value fake data (black points), stacked signal (blue), beam background (orange), dirt (green), and EXT (grey). The fake data is drawn from the central-value Monte Carlo and so agrees with the stacked prediction by construction (a reco-level self-closure); the differential results and their closure against the A_C -smeared truth are shown in Fig. 21.




Figure 15: Reconstructed-space spectra, RHC full exposure (Runs 1+2+3+4a+4b+4c); same components as Fig. 14.




Figure 16: Reconstructed-space spectra, combined FHC+RHC (per-mode normalised); same components as Fig. 14.

Figure 17 shows the response matrices (normalised per true bin) for the five observables. The three angular observables and p_μ are near-diagonal; for p_μ the width is set mainly by the MCS estimator, which reconstructs 63% of the selected sample and is intrinsically less precise than range.

p_π is the outlier. Its response is markedly non-diagonal: reconstructed events pile into the first bin from every true bin, so the diagonal band is largely absent. The likely cause is physical rather than algorithmic — charged pions interact hadronically before coming to rest, so the track range under-estimates the momentum, and no range-based estimator can fully recover it. The consequence is that the unfolded p_π result leans more heavily on the regularisation than the other observables, and it is the observable for which a dedicated resolution and binning study would be most valuable.




Figure 17: Response matrices A_{ij} for the five observables. Each column is normalised to unity (migration probability). Colour scale: white = 0, dark blue = 1.

Figure 18 shows the selection efficiency per true bin for all five observables, computed as the ratio of selected to generated signal MC.




Figure 18: Selection efficiency as a function of true observable value, in the measured phase space ($p_\mu > 0.15$, $p_\pi > 0.175 \text{ GeV}/c$, $\theta_{\mu\pi} < 2.6 \text{ rad}$). Efficiency is the ratio of selected to generated signal MC ($\varepsilon_i = \sum_j A_{ij}$, integrated over reco bins). With the momentum thresholds and the 2.6 rad opening-angle cut the efficiency is flat at $\sim 13\text{--}20\%$ across all observables; the former $\sim 5\%$ low-momentum turn-on bins and the cosmic-dominated large-angle region are removed. The lowest point is the smallest opening-angle bin ($\sim 9\%$ efficiency) — it is retained because it is high-purity ($\sim 58\%$), not low-purity.

8.7 Reconstruction resolution

Figure 19 shows the reconstruction resolution of each measured observable, evaluated on selected signal events (GENIE ν_μ overlay, measured phase space). For the momenta the relative resolution $(\text{reco} - \text{true})/\text{true}$ is shown; for the angular observables the absolute residual $(\text{reco} - \text{true})$ is used, since the relative metric diverges as $\cos\theta \rightarrow 0$ / $\theta_{\mu\pi} \rightarrow 0$. The red points are the mean residual per true bin; the quoted numbers are the overall mean (bias) and RMS (resolution).




Figure 19: Reconstruction resolution vs. true value, selected signal. Momenta: relative $(\text{reco} - \text{true})/\text{true}$; angles: absolute $(\text{reco} - \text{true})$. Red points: mean residual per true bin. The muon MCS momentum slightly over-estimates at low p_μ and increasingly under-estimates above ~ 1 GeV/ c ; the range estimator is badly biased for *exiting* muons (range = distance to wall), which is why the analysis uses range only for contained muons (range-or-MCS). The pion panel shows the estimator the analysis actually bins on, not the raw `candidate_pion_mom_reco` branch (§8.8): mean bias -31% , RMS 27% . That average is dominated by high momenta — the bias is only -2% in the first true bin and reaches -65% above 0.8 GeV/ c — because the reconstructed value saturates near 0.34 GeV/ c once pions interact hadronically before stopping. The angular variables are essentially unbiased with narrow resolution.

8.8 The pion-momentum bias and why p_π is the worst-conditioned observable

The p_π response is the least diagonal of any observable in the analysis. This section quantifies why, because the cause is physical and sets a floor on what the measurement can say about the pion momentum spectrum.

Two separate effects, only one of them corrected. The reconstructed pion momentum stored in the ntuple, `candidate_pion_mom_reco`, is `trk_range_muon_mom_v`: a range-based momentum evaluated under the *muon* mass hypothesis and applied to a pion track. Two distinct things follow.

First, the mass hypothesis itself biases the estimate low, because for a given range a pion carries more momentum than a muon. The size is $\sim 17\text{--}19\%$ from the CSDA range scaling, and it is essentially independent of momentum. This effect *is* corrected, though not in the branch: the

reco bin definitions in `ccpi_ppi_bin_config_opt.txt` apply a kinetic-energy-preserving $\mu \rightarrow \pi$ mass swap, $E_\mu \rightarrow T \rightarrow E_\pi \rightarrow p_\pi$, before binning. It works — in the lowest true bin $\langle p^{\text{reco}}/p^{\text{true}} \rangle$ goes from 0.888 with the raw branch to 0.976 with the correction.

Second, and not corrected by anything, charged pions interact hadronically before coming to rest. The reconstructed track then ends at the interaction point rather than the stopping point, so the visible range — and any momentum inferred from it — is truncated. This grows rapidly with momentum and comes to dominate.

The reconstructed momentum saturates. Table 13 and Figure 20 make the consequence plain: as the true momentum rises by a factor of four and a half, the reconstructed momentum barely moves, saturating at ≈ 0.34 GeV/ c . Above that the estimator carries almost no information about how energetic the pion actually was.

Table 13: Pion momentum reconstruction versus truth, selected signal, Run-1 FHC. “raw” is the `candidate_pion_mom_reco` branch; “corr.” additionally applies the $\mu \rightarrow \pi$ mass swap used by the reco bin definitions. The mass correction removes the bias at threshold; nothing removes the saturation above it.

true p_π [GeV/ c]	N	$\langle p^{\text{true}} \rangle$	$\langle p^{\text{corr}} \rangle$	raw/true	corr./true
0.175–0.250	528	0.210	0.205	0.888	0.976
0.250–0.300	270	0.276	0.238	0.790	0.863
0.300–0.350	254	0.325	0.253	0.716	0.779
0.350–0.400	265	0.374	0.270	0.664	0.721
0.400–0.500	403	0.443	0.279	0.581	0.629
0.500–0.600	252	0.548	0.307	0.520	0.560
0.600–0.800	272	0.688	0.326	0.441	0.473
0.800–1.200	164	0.951	0.336	0.329	0.353

What that does to the response matrix. Migrating the five analysis bins with the corrected estimator gives a nearly lower-triangular matrix: the diagonal fraction falls from 94.3% in the first true bin to 36.5, 23.1, 10.4 and finally 6.2% in the last. For the highest true bin ($p_\pi > 0.55$ GeV/ c) only 6.2% of events are reconstructed in the corresponding reco bin while 32.6% fall all the way into the *lowest* one. Every true bin above the first therefore deposits most of its events into reco bin 1, which is exactly the “pile-up into the first bin” seen in Figure 17.

How the BNB CC1 $\pi^\pm$ analysis handles the same problem. The published BNB CC1 $\pi^\pm$ Xp analysis (internal note v0.91) faces an identical estimator — it too takes “the momentum of pion candidates from range” — and its treatment is worth reproducing here, because it is a solved problem there and an open one here.

That analysis does three things we do not. First, it defines a *golden pion*: one that comes to rest and leaves a Bragg peak, as opposed to one that scattered or interacted. Second, it trains a dedicated *golden-pion BDT* and applies it as an extra cut for the pion-momentum differential cross section only, leaving the total and angular results on the generic selection. The BDT separates golden from non-golden pions using track “wiggleness” (the standard deviation of the angular difference between directions of adjacent points along the fitted track, so a direct handle on scattering), the number of reconstructed descendant particles, the track–shower score, and the Bragg likelihood ratio. The cost and benefit are quoted explicitly: efficiency falls from 20.3% to 9.5% and purity is essentially unchanged (51.9 $\rightarrow$ 49.5%), but the golden fraction of the selected signal rises from 36.3% to 67.7%. They apply the same logic to the muon momentum, restricting that cross section to contained muons (8.4% efficiency, 56.1% purity) so it uses range only.




Figure 20: Mean reconstructed against mean true pion momentum, selected signal. The analysis estimator (filled, $\mu \rightarrow \pi$ mass corrected) tracks truth at threshold but flattens out near 0.34 GeV/c: a true 0.95 GeV/c pion reconstructs at the same place as a true 0.55 GeV/c one. The raw `candidate_pion_mom_reco` branch (open) sits systematically below it by the uncorrected mass-hypothesis offset. The saturation is caused by hadronic interaction truncating the track range, and no range-based estimator can recover it.

Third, and most directly comparable, their binning is chosen against an explicit criterion: each bin must hold at least 100 predicted MC events *and* the diagonal element of the smearing matrix must exceed 0.68. The resulting pion-momentum edges are 0.10, 0.16, 0.19, 0.22, 0.60 GeV/c with an overflow bin — fine structure only where the estimator still tracks truth, then a single wide bin from 0.22 to 0.60 and everything above 0.60 in overflow.

Set against that criterion our binning does not survive: the diagonal fractions measured above are 94.3, 36.5, 23.1, 10.4 and 6.2%, so *one* of five bins passes 0.68 and the top three fail it by a wide margin. Our edges (0.175–0.550 in five bins, open above) put four boundaries inside precisely the region where the reconstructed momentum has saturated.

Consequences, and what would improve it.

Consequences, and what would improve it. This is not a bias on the extracted cross section — the response matrix contains it and the unfolding inverts it — but it is the reason the p_π result leans harder on the regularisation than any other observable, and it explains the A_C survey result that the momentum observables retain only 11–21% of their truth-level model discrimination (Table 25) while $\theta_{\mu\pi}$ retains 58–93%. With 6% diagonal strength in the top bin, the measurement above ~ 0.5 GeV/c is effectively inferred from the response model rather than measured.

The BNB note therefore supplies a tested recipe, and *all four* of its BDT inputs are already available to us. The Bragg likelihood ratios (`trk_bragg_pion`, `_mip`, `_mu`, `_p`) and the track-shower score are in the ntuples — the multi-pion pion-ID BDT in this analysis already trains on exactly those — and the descendant counts are `pfp_trk_daughters_v` / `pfp_shr_daughters_v`. Wiggleness is there too, under a different name: the PeLEE ntuples carry `trk_avg_deflection_stdev_v` (with companions `_mean_v` and `_separation_mean_v`), which is precisely the standard deviation

of the angular deflection between adjacent track points. It is fully populated. It was simply not among the branches `ProcessNTuples` copied forward; it now is.

Wiggleness works, but with the opposite sign to the naive expectation. Testing it directly on truth-matched charged pions in the Run-1 FHC overlay — calling a pion “golden-like” when the mass-corrected range momentum recovers the true momentum to within 20%, and “interacting” when it recovers less than 60% — gives a clear separation, but golden pions are the *wigglier* ones, by about a factor two in the mean (0.0048 against 0.0027). That is not a momentum artefact: it survives binning in track length, with the ratio staying at 0.51–0.63 from 5 cm to 400 cm.

The sign makes physical sense on reflection. A pion that stops travels its full range and scatters violently through its final slow centimetres, so the deflection spread averaged along the track is large. A pion that interacts hadronically is cut short while still fast and never reaches that wiggly end, so its track is straighter. High wiggleness is thus a signature of *having stopped* — exactly the population for which a range-based momentum is meaningful.

As a single cut it is already useful: on this sample the golden fraction is 39.9% before any cut, and requiring `trk_avg_deflection_stdev_v` > 0.001 keeps 86.8% of golden pions while rejecting 48.8% of interacting ones, lifting the golden fraction to 53.0%. A tighter cut at 0.002 gives 57.8% efficiency and 55.7% golden. For reference the BNB’s four-variable BDT reaches 67.7% golden at roughly half the generic-selection efficiency, so one variable recovers a good part of the gain and the remaining three — all of which we have — should close much of the rest.

In increasing order of effort, then: rebin p_π against the BNB’s own criterion (diagonal > 0.68), which on the numbers above means roughly two resolved bins below ~ 0.25 GeV/ c plus a wide bin and an overflow, and can be done today; add a golden-pion cut for the p_π cross section built from the Bragg, track-score and descendant variables already available, accepting the factor-two efficiency loss the BNB analysis accepts; or reproduce their BDT in full, which is now a training exercise rather than a production change. None changes the total or angular results, which do not use this estimator. The first is worth doing before the p_π spectrum is quoted as a shape measurement.

Adopted response: a two-bin p_π scheme

The binning has been changed accordingly. The five-bin scheme is replaced by two bins, [0.175, 0.205] GeV/ c and everything above, which raises the migration diagonals from 99/24/13/6/6% to **97/62%**.

The boundary is not a free choice dressed up as an optimisation. The worst diagonal is always the upper bin and it falls monotonically as the boundary rises, so the binding constraint is the *width* of the first bin rather than the diagonal itself:

boundary	first-bin width	diagonal, bin 1	diagonal, bin 2
0.185	10 MeV/ c	95.8%	69.3%
0.200	25 MeV/ c	97.1%	63.8%
0.205	30 MeV/c	97.4%	61.9%
0.250	75 MeV/ c	99.2%	46.4%

0.205 is the narrowest first bin with precedent: 30 MeV/ c matches the narrowest bin used in the BNB CC1 π analysis. Pushing lower would formally clear the 68% criterion but only by making the first bin 10 MeV/ c wide, which is not a useful cross-section bin. An independent three-bin optimisation placed its first boundary at 0.205 as well.

A three-bin variant was built and measured before being discarded. The best achievable (same 30 MeV/ c minimum width) is 97/38/39%, against 97/62% for two bins: the third bin costs about 25 points on the worst bin. Above ~ 0.25 GeV/ c the reconstructed momentum barely varies, so a

second boundary there separates almost nothing and merely divides one poorly-measured region into two equally poor halves. The same ordering holds under a golden-pion cut (67/52/46%), so it is not an artefact of the selection.

Even at 97/62% the upper bin remains below the 68% criterion, which is why the change of scope described below still stands: the two-bin scheme makes the inclusive p_π result honest about its own resolution, but it does not by itself make it a competitive shape measurement.

A golden-pion classifier was trained and tested: it works, and it is not enough

Following the BNB recipe, a golden-pion BDT was trained and evaluated (`macros/golden_pion_train.C`, `macros/golden_pion_compare.C`). The truth label is `mc_end_p < 0.01 GeV/c` — the pion ranged out — joined to reconstruction by matching `backtracked_p` to `mc_p`, which pairs 100% of candidate tracks with no ambiguity. The inputs are the three deflection variables, the four Bragg likelihood ratios, the track-shower score and the descendant count. Training used 126 518 labelled tracks from Runs 1, 2 and 4; the split is by *run*, so all numbers below are on Run 2, which the classifier never saw.

As a classifier it reproduces the BNB benchmark. ROC integral 0.758. At 44.8% efficiency it lifts the golden fraction from 43.1% to 64.9%, against the BNB’s 36.3 $\rightarrow$ 67.7% at roughly 47% efficiency. The variable ranking is dominated by the deflection variables (`wmean`, then `wstd`), with the Bragg ratios contributing essentially nothing — worth noting, since Bragg was expected to carry the stopping signature.

But it does not fix the binning, and cannot. Table 14 gives the migration diagonal in the five analysis bins as the cut is tightened. The first bin is already fine and stays fine; no other bin reaches 0.68 at any working point, even at 15% efficiency.

Table 14: Golden-pion BDT: efficiency, achieved golden fraction, and the p_π migration diagonal per true bin, on the held-out run. The BNB criterion for a usable bin is a diagonal above 0.68.

BDT cut	eff.	golden frac.	b1	b2	b3	b4	b5
none	100%	43.1%	99.2	23.7	13.2	6.4	5.5
> -0.05	57.0%	58.1%	99.1	29.4	19.7	11.2	7.1
> 0.00	44.8%	64.9%	99.1	33.1	24.5	15.4	9.1
$> +0.10$	24.6%	79.4%	99.1	42.9	36.8	29.3	15.7
$> +0.15$	15.2%	84.5%	99.2	46.9	42.2	36.1	19.5
<i>perfect (truth) golden selection</i>			98.9	36.1	24.1	14.2	8.9

The last row is the decisive one, and it must be read against the *efficiency* column rather than against the tightest cut. Perfect truth-level golden selection keeps 43.1% of tracks and gives 36.1, 24.1, 14.2, 8.9% in bins 2–5; the classifier at the matching 44.8% efficiency gives 33.1, 24.5, 15.4, 9.1%. They agree bin by bin. The classifier is therefore already performing as well as perfect knowledge of which pions stopped, and no better classifier exists to be built.

A caveat on the tighter working points. The > 0.15 row looks better than perfect truth selection (46.9, 42.2, 36.1, 19.5%), which is impossible for an honest classifier and is worth understanding rather than quoting. It arises because a hard cut on the classifier preferentially keeps *long* tracks: inside the top true bin, the surviving tracks have a mean length of 120.8 cm against 78.2 cm for all pions, and a mean reconstructed momentum of 0.406 against 0.307 GeV/c. The cut is thus partly a cut on the reconstructed momentum itself, which raises the diagonal by construction while sculpting the spectrum. The apparent gain is circular and the tight working

points should not be used. This is the same failure mode anticipated for track length as a direct input; it enters here through the deflection variables, which correlate with length, even though length itself was excluded from training.

Why, and what is left. Golden pions still reconstruct at $\langle p^{\text{corr}}/p^{\text{true}} \rangle = 0.73$ on this sample. Interaction is therefore not the only problem: the momentum conversion is itself biased. A single flat rescaling does *not* repair it — applying the best global factor drops the first bin from 98.9 to 49.1% and leaves 0/5 bins passing — because the bias is momentum-dependent, exactly as the saturation curve of Figure 20 shows.

Coarser binning helps but does not rescue it either. Collapsing to two bins (0.175, 0.230, ∞) gives 98.7 and 58.3% with perfect golden selection: better, still short of 0.68. Only a single bin — that is, no differential measurement at all — trivially passes.

Conclusion. The comparison is worth stating plainly rather than filing as a partial success. The classifier does what it was built to do and matches the published benchmark, but the p_π differential measurement above ~ 0.23 GeV/ c cannot be made to meet the BNB’s own binning criterion with a range-based estimator, with or without a golden-pion cut. The BNB analysis passes that criterion because its resolved bins (0.10–0.16, 0.16–0.19, 0.19–0.22) all sit *below* our 0.175 GeV/ c threshold, in the region where range still tracks momentum; above 0.22 it uses a single wide bin and an overflow. The realistic options for this analysis are therefore to lower the pion threshold toward 0.10 GeV/ c and resolve the region where the estimator works, to adopt a two-bin-plus-overflow p_π result and say so, or to replace the range-based estimator altogether. The golden-pion BDT is a useful ingredient in any of those, but it is not by itself the fix.

Planned change of scope for p_π . On the strength of the numbers above, the current intention is to *withdraw the pion-momentum differential cross section from the inclusive selection* and to quote it instead on the golden-pion-enhanced selection, on the proton-tagged subsample, or on both. The inclusive p_π result is the one measurement in this analysis whose response cannot be made to satisfy the binning criterion the field applies, and reporting it at face value alongside observables that comfortably pass would misrepresent how well it is measured. The other inclusive observables are unaffected: the angular variables and p_μ retain their shape information, and the total cross section does not use this estimator at all.

This is stated as intent rather than as a completed study. The golden-pion route is quantified above. The proton-tagged route is not yet: requiring a tagged proton changes the phase space and the purity but does nothing directly to the pion reconstruction, so whether it improves the p_π response is an open question that needs the same migration-diagonal test applied to the CC1mu1pi1p sample before either option is adopted.

8.9 Model comparison: χ^2 against truth, tune and generators

Table 15 collects the bin-by-bin χ^2 of the unfolded fake data against each prediction, for every observable and every beam configuration. All predictions are A_C -smeared, so the comparison is made in the measurement space, and the full covariance is used. The *truth* column is the closure test: it is small throughout, confirming the chain, and is visibly larger for RHC than for FHC — consistent with the poorer RHC angular response noted in §9.2.

The $\cos\theta_\mu$ χ^2 is largely an artefact of the variable. One entry stands out: in $\cos\theta_\mu$ the NEUT and NuWro χ^2 reach 26.3/5 and 26.0/5 for FHC (and 18.4, 17.4 for RHC), an order of magnitude above every other observable. Taken at face value this would be strong evidence against those two generators. It is not. Repeating the identical comparison in θ_μ — the same events, the same data, the same bins mapped through arccos — gives 2.8/5 and 2.9/5, a factor

~ 9 smaller (Table 16). The mechanism is the A_C suppression discussed in §9.2: in $\cos \theta_\mu$ the additional smearing suppresses the sharply peaked external generators by a factor ~ 0.33 while leaving the smoother MicroBooNE tune much less affected, so NEUT and NuWro are pushed away from the data and the χ^2 inflates. In θ_μ , where the suppression is ~ 0.76 , the same models agree well. The pion angle, which is not sharply peaked, is essentially unchanged between the two parameterisations. This is a further argument for quoting the angular cross sections in θ rather than $\cos \theta$, and a caution against reading model discrimination directly off a $\cos \theta_\mu \chi^2$.

Table 15: χ^2/nfd of the unfolded fake data against each A_C -smeared prediction, for the inclusive selection. “truth” is the fake-data truth (the closure test).

Beam	Observable	truth	uB tune	GENIE	GiBUU	NEUT	NuWro
FHC	p_μ	1.66/7	3.09/7	5.99/7	3.87/7	3.77/7	2.90/7
FHC	p_π	1.96/5	3.77/5	5.05/5	3.78/5	3.33/5	3.46/5
FHC	$\cos \theta_\mu$	0.47/5	3.62/5	9.67/5	12.43/5	26.28/5	26.02/5
FHC	$\cos \theta_\pi$	0.49/4	1.20/4	3.82/4	0.99/4	1.34/4	1.19/4
FHC	$\theta_{\mu\pi}$	0.59/5	2.59/5	4.12/5	3.24/5	2.47/5	2.29/5
RHC	p_μ	5.27/7	8.57/7	10.88/7	8.81/7	8.15/7	8.48/7
RHC	p_π	6.09/5	7.04/5	9.36/5	7.28/5	5.96/5	6.13/5
RHC	$\cos \theta_\mu$	0.96/5	4.91/5	7.45/5	8.66/5	18.43/5	17.42/5
RHC	$\cos \theta_\pi$	1.47/4	4.00/4	6.64/4	3.92/4	4.21/4	4.30/4
RHC	$\theta_{\mu\pi}$	4.32/5	5.57/5	7.67/5	5.68/5	6.56/5	7.59/5
Comb	p_μ	3.41/7	3.70/7	5.82/7	3.99/7	3.83/7	3.30/7
Comb	p_π	1.32/5	2.32/5	3.85/5	2.65/5	1.82/5	2.05/5
Comb	$\cos \theta_\mu$	2.49/5	4.61/5	9.07/5	12.02/5	26.72/5	27.20/5
Comb	$\cos \theta_\pi$	0.42/4	1.83/4	4.01/4	1.69/4	1.89/4	1.95/4
Comb	$\theta_{\mu\pi}$	2.10/5	2.42/5	3.91/5	3.05/5	3.49/5	3.45/5

Table 16: The same comparison for the muon and pion angle in both parameterisations. The θ bin edges are exactly arccos of the $\cos \theta$ edges, so the phase space is identical; only the variable differs.

Beam	Variable	truth	uB tune	GENIE	GiBUU	NEUT	NuWro
FHC	$\cos \theta_\mu$	0.47/5	3.62/5	9.67/5	12.43/5	26.28/5	26.02/5
FHC	θ_μ	0.05/5	3.07/5	4.73/5	3.03/5	2.82/5	2.90/5
RHC	$\cos \theta_\mu$	0.96/5	4.91/5	7.45/5	8.66/5	18.43/5	17.42/5
RHC	θ_μ	0.48/5	4.41/5	6.97/5	5.54/5	8.18/5	8.46/5
FHC	$\cos \theta_\pi$	0.49/4	1.20/4	3.82/4	0.99/4	1.34/4	1.19/4
FHC	θ_π	0.44/4	1.15/4	3.77/4	0.93/4	1.30/4	1.18/4

8.10 Proton-tagged χ^2 and the A_C correlation

Table 17 repeats the comparison for the six proton-tagged observables (FHC), with the measured additional-smearing factor of §10.2 in the last column. Ordering the rows by $|A_C - 1|$ makes an unambiguous pattern visible: the generator χ^2 grows monotonically with the distance of A_C from unity, and does so in the same direction for every model.

Observable	A_C	$ A_C - 1 $	typical generator $\chi^2/6$
W_{had}	0.88	0.12	1.2
$W_{\pi p}$	0.78	0.22	1.0
$\delta\alpha_T$	0.63	0.37	3.3
$\delta\phi_T$	0.61	0.39	6.2
δp_T	0.43	0.57	9.4
p_n	1.71	0.71	17

The two observables the smearing barely touches, W_{had} and $W_{\pi p}$, are described well by *all four* generators ($\chi^2 \lesssim 1.8/6$). The two it distorts most — δp_T , suppressed to 0.43, and p_n , *enhanced* to 1.71 — carry the largest χ^2 , reaching 25.4/6 for NEUT in p_n . That a quantity as model-independent as $|A_C - 1|$ orders the χ^2 of four independent generators is the signature of a smearing artefact rather than of physics: it is the same effect identified in $\cos\theta_\mu$ (§8.9), where reparameterising to θ_μ reduced the NEUT and NuWro χ^2 by a factor ~ 9 on identical data. The closure column is excellent throughout ($\leq 0.46/6$), so the chain itself is sound; it is the *model comparison* in δp_T and p_n that should not be read as discrimination until the binning and response of those two observables are revisited.

Table 17: χ^2/ndf of the unfolded fake data against each A_C -smeared prediction for the proton-tagged observables (FHC), with the additional-smearing factor A_C .

Observable	truth	uB tune	GENIE	GiBUU	NEUT	NuWro	A_C
W_{had}	0.09/6	1.39/6	1.65/6	1.06/6	1.17/6	1.05/6	0.88
$W_{\pi p}$	0.09/6	0.95/6	1.78/6	0.62/6	0.56/6	0.87/6	0.78
$\delta\alpha_T$	0.26/6	2.74/6	4.73/6	2.95/6	2.78/6	3.53/6	0.63
$\delta\phi_T$	0.09/6	5.14/6	7.21/6	5.61/6	5.39/6	6.47/6	0.61
δp_T	0.46/6	5.34/6	9.89/6	8.27/6	9.02/6	10.38/6	0.43
p_n	0.09/6	4.24/6	7.72/6	18.74/6	25.40/6	16.45/6	1.71

8.11 Cross-check methods

Seven additional methods are run in parallel:

Table 18: Unfolding methods.

Label	Description
wiener_svd	Wiener-SVD (primary result)
bin_by_bin	Efficiency correction only; no migration unfolding
ibu	D’Agostini iterative Bayesian, 4 iterations
tikhonov	Tikhonov-regularised least-squares, $\tau = 1.0$
roo_bayes	RooUnfoldBayes, 4 iterations
roo_svd	RooUnfoldSvd
roo_invert	RooUnfoldInvert (direct matrix inversion)
roo_binbybin	RooUnfoldBinByBin

The D’Agostini iterative-Bayesian method provides the principal quantitative cross-check of the Wiener-SVD result. Because it does not apply the Wiener-SVD additional-smearing matrix A_C , its unfolded total is free of the (non-norm-preserving) A_C suppression and is stable against the iteration count. Table 19 compares the two integrated totals for all three configurations: the D’Agostini result sits above the Wiener-SVD total by $\sim 23\%$ (FHC) and $\sim 37\text{--}38\%$ (RHC, Combined) — equivalently A_C suppresses the quoted Wiener-SVD integral by $\sim 19\%$ (FHC) and $\sim 27\%$ (RHC/Combined), the larger suppression reflecting the coarser RHC/combined smearing

— and it varies by $< 1\%$ across 1, 2 and 4 iterations. For FHC the D’Agostini per-observable truth is flat to 1% (1.111–1.122), directly confirming that the observable-to-observable spread in the Wiener-SVD integrated cross sections (Table 23) is the A_C smearing rather than a physical or acceptance effect.

Table 19: Wiener-SVD vs. D’Agostini (iterative Bayesian, 4 iterations) integrated cross section, averaged over the five observables (fake-data truth, 10^{-38} cm²/Ar); both columns recomputed with the corrected MC POT normalisation. D’Agostini carries no A_C smearing, so it recovers a 25–29% higher total. Crucially the ratio is now *consistent* across the three configurations (1.25–1.29), whereas with the earlier, incorrect normalisation the RHC and combined ratios (1.37, 1.38) sat well above FHC (1.23) — the inconsistency was an artefact of the inflated RHC/combined Wiener-SVD values, not of the unfolding methods. D’Agostini is also nearly flat across observables (spread 1.2% FHC, 2.2% RHC, 1.5% combined), in contrast to the Wiener-SVD spread, confirming that the observable-to-observable variation of σ_{int} in Table 23 is the A_C smearing rather than a physical or normalisation effect.

Config	$\langle\sigma\rangle_{\text{WSVD}}$	$\langle\sigma\rangle_{\text{DAg}}$	DAg/WSVD	DAg 1→4 iter
FHC	0.923	1.170	1.27	$< 1\%$
RHC	0.836	1.041	1.25	$< 1\%$
Combined	0.852	1.102	1.29	$< 1\%$

9 Full-exposure results: FHC, RHC, and combined

This is the primary result of the analysis: the $\text{CC1}\mu 1\pi^\pm$ cross section extracted over the full available exposure in forward horn current (FHC), reverse horn current (RHC), and their combination, all on Poisson-thrown central-value fake data. The FHC measurement uses Runs 1, 2, 4 and 5 (data POT 8.857×10^{20}); RHC uses Runs 1, 2, 3, 4a, 4b, 4c (1.108×10^{21}); the combined measurement uses all fourteen overlay files at the joint POT 1.994×10^{21} . All three close cleanly against the A_C -smeared truth for every observable (Tables 20–22); the differential results are shown in Figs. 21–23.

Table 20: Full-exposure FHC ({Run 1,2,4,5}, data POT 8.857×10^{20} , native Run-4 FHC detector variations): reco-level and unfolded flux systematic, unfolding amplification, and fake-data closure χ^2/ndf (with p -value).

Observable	Flux (reco)	Flux (unfolded)	Amplification	Closure χ^2/ndf
p_μ	17.1%	23.9%	$1.40\times$	0.55/7 ($p = 1.00$)
p_π	16.9%	23.9%	$1.41\times$	0.29/5 ($p = 1.00$)
$\cos\theta_\mu$	16.7%	23.9%	$1.43\times$	0.93/5 ($p = 0.97$)
$\cos\theta_\pi$	16.6%	27.7%	$1.67\times$	0.10/4 ($p = 1.00$)
$\theta_{\mu\pi}$	16.8%	22.3%	$1.33\times$	1.84/5 ($p = 0.87$)

We extend the measurement to the reverse-horn-current beam as both a standalone measurement and a combined $\nu_\mu + \bar{\nu}_\mu$ result. The RHC *flux* is $\bar{\nu}_\mu$ -dominant (54.6% $\bar{\nu}_\mu$, versus FHC’s 35%), with authoritative > 60 MeV integrals $\Phi_{\nu_\mu}^{\text{RHC}} = 2.923 \times 10^{-10}$, $\Phi_{\bar{\nu}_\mu}^{\text{RHC}} = 3.523 \times 10^{-10}$, total 6.446×10^{-10} cm^{−2}POT^{−1} (the FHC-convention computation reproduces the authoritative FHC values exactly). The *measured event rate*, however, remains ν_μ -dominant (76% ν_μ , 20% $\bar{\nu}_\mu$) because the ν_μ cross section is $\sim 2\text{--}3\times$ larger than the $\bar{\nu}_\mu$ one. The RHC overlays (Runs 1, 2, 3, 4a, 4b, 4c) are all verified as genuine ν_μ overlays and processed — 4.05×10^6 events, combined MC POT 1.546×10^{22} , mutually consistent (2.62×10^{-16} events/POT) with the PPF weights — covering the full RHC exposure (Run 3, the largest period at 5.52×10^{21} MC POT / 1.44×10^6

events, was added last). Native Run-4 RHC detector-variation samples (the standard eight-knob set: light-yield down/Rayleigh, recombination, space charge, and the four wire-modification variations) are now used for the RHC detector systematic, verified as genuine ν_μ overlays with a uniform 3.80×10^{-16} events/POT across all knobs; they replace the Run-1 FHC stand-in used previously. Native RHC dirt is still taken from FHC. The standalone result below uses Poisson-thrown central-value fake data at the RHC data POT 1.108×10^{21} . The generator predictions are re-averaged to the RHC flux composition; the combined FHC+RHC predictions are the fluence-weighted combination of the two horn modes (0.458 FHC / 0.542 RHC, each mode weighted by $\Phi_{\text{tot}} \times \text{POT}$).

The combined FHC+RHC measurement uses the summed 14-file response over the full exposure (5.6×10^6 events, combined data POT $D_{\text{comb}} = 1.994 \times 10^{21} = D_{\text{FHC}} + D_{\text{RHC}}$), with each horn mode normalised to its *own* data exposure. This is a genuine subtlety: the detector-variation covariance is absolute (POT-scaled reco counts), so a single POT scale applied to all files would weight the modes by their Monte-Carlo POT (0.62:1 FHC:RHC) rather than by data exposure (0.80:1). Each mode’s numuMC and detVar samples therefore carry a per-mode `summed_pot` (FHC 2.158×10^{22} , RHC 2.782×10^{22} , $= S_{\text{mode}} \times D_{\text{comb}}/D_{\text{mode}}$) so the framework’s $D_{\text{comb}}/\text{summed_pot}$ scaling reduces to $D_{\text{FHC}}/S_{\text{FHC}}$ and $D_{\text{RHC}}/S_{\text{RHC}}$; the fake data is Poisson-thrown with the matching per-mode rate.

Contrary to an earlier concern, the joint *flux systematic* is correct as built. All fourteen numuMC files carry the *same* 600 PPFX multisim universes, which the universe maker accumulates index-matched across both horn modes in a single chain, so universe u is the same hadron-production draw in FHC and RHC. Per-universe summing therefore propagates one parameter draw through both beam configurations and lets each mode’s flux response set the result, capturing the physical (not unit) cross-horn correlation automatically — this is the rigorously correct multisim treatment, superior to any hand-built block covariance. The joint result was in fact blocked by a different, purely technical issue: a framework guard that forbade a second detector-variation file of the same type (needed to carry Run-4 FHC *and* Run-4 RHC as `detVarCV`). That guard now accumulates the per-mode detVar samples exactly as it already did for the alternate-CV samples, matching its own in-code TODO. With both in place the combined closure passes for every observable (Table 22), with an unfolded flux systematic of 22–27% that sits between the standalone FHC and RHC values, as expected for the fluence-weighted mix.

The standalone RHC extraction (native Run-4 RHC detector variations) closes cleanly for every observable (Table 21): the A_C -smeared truth is recovered with $\chi^2/\text{ndf} < 0.3$ and $p > 0.94$ throughout. Two features stand out. First, the reco-level flux systematic is $\sim 15\%$, genuinely below the FHC $\sim 17\%$ floor — the $\bar{\nu}_\mu$ -dominant flux carries a smaller PPFX hadron-production uncertainty. Second, the unfolding amplification ($1.29\text{--}1.99\times$) mirrors FHC, so the unfolded flux term (19.5–28%) is comparable, again data-statistics limited.

Table 21: Standalone RHC ($\bar{\nu}_\mu$ -dominant) extraction, full exposure 1.108×10^{21} with native Run-4 RHC detector variations: reco-level and unfolded flux systematic, unfolding amplification, and fake-data closure χ^2/ndf (with p -value).

Observable	Flux (reco)	Flux (unfolded)	Amplification	Closure χ^2/ndf
p_μ	15.4%	27.5%	$1.78\times$	1.37/7 ($p = 0.99$)
p_π	15.8%	23.6%	$1.49\times$	0.21/5 ($p = 1.00$)
$\cos \theta_\mu$	15.0%	26.5%	$1.76\times$	0.43/5 ($p = 0.99$)
$\cos \theta_\pi$	14.9%	28.5%	$1.91\times$	1.03/4 ($p = 0.91$)
$\theta_{\mu\pi}$	15.1%	20.4%	$1.35\times$	0.34/5 ($p = 1.00$)

Table 22: Combined FHC+RHC ($\nu_\mu + \bar{\nu}_\mu$) extraction over the full joint exposure $D_{\text{comb}} = 1.994 \times 10^{21}$, per-mode normalised, with native Run-4 FHC+RHC detector variations (accumulated per mode): reco-level and unfolded flux systematic, unfolding amplification, and fake-data closure χ^2/ndf (with p -value). The unfolded flux term (22–27%) lies between the standalone FHC and RHC values.

Observable	Flux (reco)	Flux (unfolded)	Amplification	Closure χ^2/ndf
p_μ	16.2%	22.5%	$1.39\times$	0.58/7 ($p = 1.00$)
p_π	16.3%	23.9%	$1.47\times$	1.27/5 ($p = 0.94$)
$\cos\theta_\mu$	15.8%	25.0%	$1.58\times$	0.28/5 ($p = 1.00$)
$\cos\theta_\pi$	15.7%	27.0%	$1.72\times$	1.19/4 ($p = 0.88$)
$\theta_{\mu\pi}$	15.9%	26.0%	$1.64\times$	0.43/5 ($p = 0.99$)




Figure 21: Differential cross sections $d\sigma/dx$ for all five observables on the current Poisson-thrown central-value fake data, FHC full exposure: unfolded fake data (black points), A_C -smeared fake-data truth (grey dashed), the MicroBooNE tune (black solid), and the four generator predictions (GENIE, GiBUU, NEUT, NuWro). The unfolded points recover the injected truth for every observable, the same closure summarised in Tables 21–22.




Figure 22: Standalone RHC ($\bar{\nu}_\mu$ -dominant) differential cross sections on the current Poisson-thrown fake data, with the RHC-flux-averaged generator predictions. Same layout as Fig. 21.




Figure 23: Combined FHC+RHC ($\nu_\mu + \bar{\nu}_\mu$) differential cross sections on the current per-mode-normalised fake data, with the fluence-weighted combined generator predictions.

Table 23: Integrated (unfolded) cross section σ_{int} over the measured phase space, for each observable and configuration (current fake data, coarsened $\cos\theta_\pi/\theta_{\mu\pi}$ binning, per-configuration flux; 10^{-38} cm²/Ar). The value differs between observables because the Wiener-SVD additional-smearing A_C is not norm-preserving; the D’Agostini cross-check (Sec. 8.11) recovers a flatter, $\sim 20\%$ higher total, confirming the spread is a smearing effect. All entries were recomputed in a single pass with the current code and a forced rebuild of the cached universes. The MC POT normalization scales each ntuple file by its own stored POT ($D_{\text{run}}/\text{POT}_{\text{file}}$): the processing writes the *total* run MC POT into every file of a run, so for the multi-file RHC runs (Run 3 split across 5 ntuple files, Run 4 across 3) the per-file scaling summed over the splits gives the correct total. An intermediate version of this analysis instead divided by the *sum* of the per-file POT values; because those values are already run totals, that over-counted the denominator by the split multiplicity and scaled the RHC/Combined simulation a factor of 5 (Run 3) and 3 (Run 4) too low, which under-subtracted the background and inflated the RHC/Combined cross sections. That change has been reverted and the RHC/Combined columns here are the corrected values, roughly 10–25% below the intermediate ones; FHC (one MC file per run) is unaffected by either convention. The last two rows give the alternative angular parameterisation θ_μ, θ_π (Sec. 9.2), now complete for all three configurations. (Values are fake-data closures; the signal is $\nu_\mu + \bar{\nu}_\mu$, so the RHC flux-average is not expected to lie far below FHC.)

Observable	FHC	RHC	Combined
p_μ	0.987	0.860	0.818
p_π	0.964	0.972	0.962
$\cos\theta_\mu$	0.819	0.662	0.702
$\cos\theta_\pi$	0.969	0.805	0.857
$\theta_{\mu\pi}$	0.878	0.880	0.922
θ_μ	0.996	0.828	0.783
θ_π	0.969	0.802	0.857

Uncertainties shown. Every differential cross section and every integrated value quoted in this note carries *both* statistical and systematic uncertainties. The error bars on the unfolded points are the square root of the diagonal of the *total* covariance matrix, propagated through the unfolding: flux (PPFX multisims), cross-section (GENIE multisims and unisims), detector variations, hadron re-interaction, POT counting and target number, together with the data, MC and EXT statistical terms. For p_μ in FHC the total is $\simeq 27\%$, dominated by flux (23%), with detector 7%, cross-section 8% and data statistics 8%. The cut-and-count total cross sections of §9.1 quote the same combination, written there as systematic $\oplus$ data-statistical (FHC: $25.5\% \oplus 4.3\%$).

9.1 Total flux-averaged cross section (cut-and-count)

As a check independent of the unfolding, the total flux-averaged cross section is also extracted directly from the selected event counts,

$$\sigma = \frac{N_{\text{sel}} - N_{\text{bkg}}}{\epsilon \Phi N_{\text{Ar}}}, \quad (11)$$

where N_{sel} is the number of selected events, N_{bkg} the predicted background (ν -background + EXT + dirt), $\epsilon = N_{\text{sel}}^{\text{sig}}/N_{\text{gen}}^{\text{sig}}$ the selection efficiency in the measured phase space ($\simeq 17.5\%$ for all configurations), Φ the integrated flux ($\Phi = \text{POT} \times \Phi_{\text{POT}}$, Table 3) and $N_{\text{Ar}} = 8.71 \times 10^{29}$ the number of argon nuclei in the fiducial volume (the dead z region removed), giving the same normalisation $N_{\text{Ar}} \Phi/10^{38}$ used for the differential result. This uses no matrix inversion; by avoiding the Wiener-SVD additional smearing it recovers the true (un-suppressed) total, consistent with the D’Agostini cross-check (Sec. 8.11) and lying above the smearing-suppressed σ_{int}

of Table 23, most visibly for RHC. The uncertainty is the quadrature of the systematic breakdown (Tables 28–30, flux-dominated) and the cut-and-count data-statistical term $\sqrt{N_{\text{sel}}}/(N_{\text{sel}} - N_{\text{bkg}})$, which is smaller (3–4%) than the unfolded data-statistical term because no regularisation amplifies it.

9.2 Angular variable: $\cos\theta_\mu$ versus θ_μ

The muon-angle distribution is measured in both $\cos\theta_\mu$ and θ_μ . The two carry identical information — the θ_μ bin edges are exactly arccos of the $\cos\theta_\mu$ edges, so the phase space is partitioned identically — but they behave very differently under the Wiener–SVD additional smearing A_C , and the choice matters for how the measurement compares to models.

In $\cos\theta_\mu$ the signal is sharply forward-peaked: the last two bins ($\cos\theta_\mu > 0.8$) hold most of the rate, and A_C suppresses the integral of any truth-level prediction by a factor ~ 0.33 , against ~ 0.6 for the momentum observables. The four generators are pushed $2\text{--}3\times$ below the tune and the data even though their truth-level total cross sections agree to within tens of percent.

An earlier version of this note attributed both effects to the sharpness of the peak — A_C smearing away an unresolvable feature, and thereby penalising the peaked generators more than the smoother tune. *Both halves of that explanation are wrong*, and are retracted here. Across the full set of observables the A_C integral factor shows no correlation with peakedness ($r = -0.17$): $W_{\pi p}$ is the sharpest distribution in the analysis (peak-to-minimum ratio 36) yet is suppressed only to 0.78, δp_T is comparatively flat (ratio 10) yet is suppressed to 0.43, and p_n is flat yet *enhanced* to 1.71 — an enhancement no peak argument can produce. Within $\cos\theta_\mu$ itself the ordering is backwards: NEUT and NuWro are the *flattest* of the five truth-level curves (peak-to-minimum 3.8 and 4.3, against 16.3 for the tune) and are suppressed the *most* (0.37 and 0.29, against 0.67). What actually drives their suppression is that A_C pushes their backward bins negative (-0.0005 to -0.0212 in normalised units), eating the integral. A_C is the Wiener filter acting on the SVD modes of the response: it removes whatever the data cannot statistically constrain, which is set by the conditioning of the response for that observable and binning, not by the shape of the truth distribution — and it is not norm-preserving, so it can suppress or enhance.

The FHC θ_μ result is withdrawn: A_C is rank-1 there

Reparameterising to θ_μ appeared to remove the problem entirely: the FHC suppression rose to ~ 0.76 and the NEUT/NuWro χ^2 fell from $\sim 26/5$ to $\sim 2.8/5$, with the generators sitting on top of the data. *That result is an artefact of over-regularisation and must not be used for model comparison.*

In the FHC θ_μ extraction the unfolded data, the A_C -smeared truth, the MicroBooNE tune and all four generators share the same normalised shape to four decimal places (Table 24). The cause is that A_C is *rank-1* for this config: writing the four generator predictions as the columns of a 5×4 input matrix and their A_C -smeared counterparts as the columns of the output matrix, the inputs are full rank (ratio of second to first singular value $s_2/s_1 = 5.0 \times 10^{-2}$) while the outputs have $s_2/s_1 = 2.8 \times 10^{-6}$. Every shape degree of freedom has been annihilated; each model is mapped onto a single common curve and only its normalisation survives. The agreement with the data is therefore true by construction, the effective number of degrees of freedom is ≈ 1 rather than 5, and the quoted $\chi^2/5$ is not a goodness-of-fit.

This is specific to FHC θ_μ , not a property of the θ parameterisation. Fed the identical generator inputs, FHC $\cos\theta_\mu$ returns $s_2/s_1 = 5.1 \times 10^{-2}$ and FHC p_μ returns 6.6×10^{-3} ; RHC θ_μ returns 2.3×10^{-2} and does retain its shape information, which is why its measured improvement is only $0.34 \rightarrow 0.49$ — previously explained as a poorer RHC angular response, when in fact RHC is the case that was *not* over-regularised. The generator FTE inputs themselves are sound: the θ_μ histograms are the exact bin-reversal of the $\cos\theta_\mu$ ones and differ properly between models.

Table 24: FHC θ_μ : normalised bin contents. Five independent inputs — the measurement and four generator predictions plus the tune — collapse onto one shape, the signature of a rank-1 A_C .

	bin 1	bin 2	bin 3	bin 4	bin 5
Unfolded data	0.1485	0.3361	0.2899	0.1934	0.0321
A_C -smeared truth	0.1484	0.3359	0.2901	0.1933	0.0323
MicroBooNE tune	0.1484	0.3359	0.2901	0.1934	0.0322
NEUT	0.1484	0.3359	0.2901	0.1934	0.0322

Consequently the muon-angle measurement is quoted in $\cos\theta_\mu$, which retains its shape modes, and the FHC θ_μ numbers are withdrawn pending a re-run with a regularisation strength that does not collapse the response. The strong $\cos\theta_\mu$ suppression and the negative backward bins in NEUT and NuWro remain real features of the measurement and are reported as such. The RHC generator predictions themselves are healthy, sitting a uniform 0.80–0.84 of their FHC counterparts as expected for the $\bar{\nu}_\mu$ -enriched beam. The same reparameterisation applied to the pion angle changes almost nothing (θ_π vs $\cos\theta_\pi$: 0.802 vs 0.805 for RHC); with the peak explanation retracted this is simply because the pion-angle response is comparably conditioned in the two parameterisations.

Is any other observable affected? A survey of all eighteen extractions

The rank-1 failure was found by inspection, so the unfolders was modified to write the slice-restricted A_C into the closure file (`h_A_C`; it had previously only been drawn to a canvas, and the attempt to save it failed silently with “the current directory is not associated with a file”). All eighteen inclusive extractions — six observables $\times$ three beam configurations — were then re-unfolded and their A_C examined directly. Table 25 reports, for each, how much of the model discrimination present in the truth-level predictions survives the smearing: the largest pairwise difference between normalised generator shapes before and after A_C , and their ratio.

The rank-1 collapse is *unique to FHC θ_μ* : it is the only entry with zero retention, and the only A_C whose rows are all the same vector (spread 9% across rows, effective rank 1.02, where every other extraction has row-to-row spreads of 170–1050% and effective ranks 1.24–1.91). RHC and combined θ_μ are not affected; the combined case is the weaker of the two (12% retention, effective rank 1.06) and its θ_μ model comparison should be read as a soft test, but it is not degenerate.

Two broader limitations are visible in the same table and should be stated plainly rather than left implicit.

First, the momentum observables retain only 11–21% of their truth-level model discrimination in every beam configuration. A_C does not destroy the shape information there, but it removes most of it, so a good χ^2 against p_μ or p_π is a weak statement about a model. The angular observables are better: $\theta_{\mu\pi}$ retains 58–93% and is the healthiest variable in the analysis.

Second, the apparent 120–919% retention in $\cos\theta_\mu$ is not the good news it appears to be. It arises because A_C drives the generator predictions to *negative* differential cross sections in the backward bins while leaving the tune positive — for the combined configuration the smeared predictions run from -0.43 (NuWro) to $+0.51$ (the tune) in the second bin. A negative differential cross section is unphysical, so the large spread reflects a regularisation artefact rather than genuine model separation, and any χ^2 computed against those curves inherits the problem. This is the same effect that produces the strong $\cos\theta_\mu$ integral suppression discussed above.

Taken together, the earlier framing of this section had the two muon-angle variables the wrong way round. $\cos\theta_\mu$ was described as the variable A_C ruins and θ_μ as the fix; in fact FHC θ_μ is the one extraction in which the measurement retains no shape information at all, while $\cos\theta_\mu$ retains its shape modes but at the cost of unphysical negative excursions. Neither is a

Table 25: Retention of model discrimination through A_C . “Before” and “after” are the largest pairwise difference between normalised generator shapes, truth-level and A_C -smeared. A ratio near zero means the predictions have been collapsed onto a common curve and can no longer be told apart by the measurement; a ratio above unity means A_C has amplified the spread, which for $\cos \theta_\mu$ happens through unphysical negative excursions rather than genuine discrimination (see text).

Beam	Observable	before	after	retention
FHC	p_μ	0.0248	0.0052	21%
FHC	p_π	0.0586	0.0080	14%
FHC	$\cos \theta_\mu$	0.0554	0.0724	131%
FHC	$\cos \theta_\pi$	0.0589	0.0211	36%
FHC	$\theta_{\mu\pi}$	0.0422	0.0391	93%
FHC	θ_μ	0.0554	2.5×10^{-6}	0%
RHC	p_μ	0.0232	0.0036	15%
RHC	p_π	0.0556	0.0068	12%
RHC	$\cos \theta_\mu$	0.0644	0.0771	120%
RHC	$\cos \theta_\pi$	0.0629	0.0165	26%
RHC	$\theta_{\mu\pi}$	0.0505	0.0295	58%
RHC	θ_μ	0.0644	0.0188	29%
COMB	p_μ	0.0240	0.0028	12%
COMB	p_π	0.0571	0.0064	11%
COMB	$\cos \theta_\mu$	0.0598	0.5500	919%
COMB	$\cos \theta_\pi$	0.0609	0.0166	27%
COMB	$\theta_{\mu\pi}$	0.0463	0.0360	78%
COMB	θ_μ	0.0598	0.0070	12%

clean muon-angle measurement, and both carry the caveats recorded here.

Table 26: Integrated cross section ($10^{-38} \text{ cm}^2/\text{Ar}$) for the unfolded fake data, the MicroBooNE tune and the four generators, in $\cos\theta_\mu$ and in θ_μ . All entries are A_C -smeared, i.e. live in the same space as the measurement. “ A_C supp.” is the mean ratio of the A_C -smeared to the truth-level generator integral; the momentum observable p_μ is shown for reference. **The FHC θ_μ row is retained for the record only and must not be used for model comparison:** A_C is rank-1 for that config, so every entry in the row is the same shape rescaled and the apparent agreement is by construction (Table 24).

Beam	Variable	data	tune	GENIE	GiBUU	NEUT	NuWro
FHC	$\cos\theta_\mu$	0.819	0.673	0.274	0.390	0.367	0.295
FHC	θ_μ	0.996	0.871	0.625	0.876	1.025	0.919
RHC	$\cos\theta_\mu$	0.662	0.633	0.226	0.338	0.320	0.256
RHC	θ_μ	0.828	0.824	0.328	0.487	0.505	0.423
A_C supp. FHC:		$\cos\theta_\mu$	0.23–0.34	θ_μ	0.73–0.78	$(p_\mu$ 0.59–0.67)	
A_C supp. RHC:		$\cos\theta_\mu$	0.25–0.36	θ_μ	0.41–0.52	$(p_\mu$ 0.69–0.75)	




Figure 24: Muon-angle differential cross section in θ_μ (left) and $\cos\theta_\mu$ (right), for FHC (top) and RHC (bottom). Unfolded fake data (black points), the MicroBooNE tune (black line) and the four generators are all A_C -smeared. In $\cos\theta_\mu$ the predictions are strongly suppressed by A_C and fall well below both the data and the tune; the most suppressed models are those whose backward bins A_C drives negative. **The FHC θ_μ panel is not a model comparison:** A_C is rank-1 there, so the data and every prediction are the same curve up to normalisation and their apparent agreement carries no shape information (Sec. 9.2).

Table 27 compares the result with the four generator predictions, each flux-averaged over the

corresponding beam. Because the analysis is blind the “measurement” is here the central-value overlay (the MicroBooNE tune) that the fake data is thrown from; at unblinding N_{sel} becomes the real beam-on count. It agrees with all four generators within $\sim 1\sigma$: standalone GENIE sits about 1σ low, while GiBUU, NEUT and NuWro lie within half a sigma. RHC and the combined total carry the larger uncertainty, driven by the RHC detector systematic.

Table 27: Total flux-averaged cross section from the cut-and-count method (Eq. 11; $10^{-38} \text{ cm}^2/\text{Ar}$), compared with the flux-averaged generator predictions. The measurement uncertainty is the systematic breakdown (flux/detector/cross-section/reinteraction, Tables 28–30) $\oplus$ cut-and-count data statistics. Values in parentheses give each generator’s deviation from the measurement in units of the measurement uncertainty.

Config	Measurement	GENIE	GiBUU	NEUT	NuWro
FHC	1.06 ± 0.27	$0.79 (1.0\sigma)$	$1.08 (0.1\sigma)$	$1.26 (0.8\sigma)$	$1.20 (0.5\sigma)$
RHC	0.98 ± 0.35	$0.63 (1.0\sigma)$	$0.89 (0.3\sigma)$	$1.04 (0.2\sigma)$	$0.98 (0.0\sigma)$
Combined	1.02 ± 0.30	$0.70 (1.1\sigma)$	$0.97 (0.1\sigma)$	$1.15 (0.4\sigma)$	$1.08 (0.2\sigma)$

The systematic breakdown by source (bin-averaged fractional uncertainty in cross-section space) is given for all three configurations in Tables 28–30. The flux (PPFX) term dominates throughout, followed by the detector variations; the data-statistical term reflects the fake-data POT of each configuration. The RHC detector term is the largest single-configuration entry, driven by the $\cos\theta$ observables at the current statistics.

Table 28: FHC ({Run 1,2,4,5}) systematic breakdown by source, bin-averaged fractional uncertainty (%).

Source	p_μ	p_π	$\cos\theta_\mu$	$\cos\theta_\pi$	$\theta_{\mu\pi}$
Total	27.3	31.9	33.3	30.0	29.2
Cross section (GENIE)	8.0	8.9	8.6	8.3	9.2
Flux (PPFX)	22.9	25.8	26.0	25.2	23.0
Detector	7.0	9.8	11.8	9.2	7.8
Reinteraction	3.4	4.0	3.5	3.9	3.7
MC stat	2.6	3.4	3.9	2.4	3.6
EXT stat	2.8	3.8	4.4	3.0	3.2
Data stat	8.3	10.6	12.3	7.6	11.5
POT + targets	2.7	3.3	3.3	3.2	2.9

Table 29: RHC ({Run 1,2,3,4a,4b,4c}) systematic breakdown by source, bin-averaged fractional uncertainty (%).

Source	p_μ	p_π	$\cos\theta_\mu$	$\cos\theta_\pi$	$\theta_{\mu\pi}$
Total	28.2	28.8	38.8	32.2	27.0
Cross section (GENIE)	8.4	8.2	10.8	9.2	7.9
Flux (PPFX)	21.4	22.4	28.5	25.3	20.8
Detector	10.4	12.0	16.0	13.7	9.5
Reinteraction	3.1	3.5	4.0	4.0	2.7
MC stat	2.8	2.1	4.1	2.2	2.8
EXT stat	3.4	3.4	5.9	3.2	3.1
Data stat	10.3	7.7	15.1	8.1	10.1
POT + targets	2.9	3.2	3.9	3.5	2.9

Table 30: Combined FHC+RHC systematic breakdown by source, bin-averaged fractional uncertainty (%).

Source	p_μ	p_π	$\cos \theta_\mu$	$\cos \theta_\pi$	$\theta_{\mu\pi}$
Total	30.1	33.7	41.0	34.8	29.3
Cross section (GENIE)	9.4	8.5	9.9	8.6	7.6
Flux (PPFX)	22.1	24.2	29.1	25.3	21.6
Detector	13.5	19.6	23.5	20.6	15.6
Reinteraction	3.8	3.5	3.9	3.8	3.0
MC stat	3.0	1.8	3.3	1.6	2.2
EXT stat	3.9	2.5	3.8	2.1	2.5
Data stat	10.0	6.1	11.0	5.4	7.6
POT + targets	2.7	3.2	3.8	3.3	2.8

9.3 Background-control sidebands

Because the analysis is blind, the background model cannot yet be confronted with beam-on data in the signal region. To exercise it before unblinding, four signal-depleted control regions (“sidebands”) are defined by inverting exactly one signal requirement, leaving the rest of the selection intact:

- **CC0 π** : the muon selection with no reconstructed charged pion (charged-current quasi-elastic / $2p2h$ -enriched);
- **multi- π** : two or more charged pions (deep-inelastic / resonant multi-pion);
- **π^0** : a reconstructed shower present but the π^0 shower veto failed (CC π^0 / NC π^0);
- **cosmic**: the full signal selection with the muon–pion opening angle inverted ($\theta_{\mu\pi} > 2.6$), enriched in out-of-time cosmic activity.

Each region is evaluated with the same per-run POT scaling and CV event weighting (§3) used for the measurement. Because the current “data” is a Poisson throw of the neutrino Monte Carlo central value only (it contains no EXT or dirt component), the meaningful closure quantity here is the ratio of the pseudo-data to the CV-weighted *neutrino* MC, $N_{\text{data}}/N_{\nu\text{MC}}$; the EXT and dirt yields are shown for scale but enter the background-model test only once real beam-on data are available. Table 31 summarises the four regions for the three beam configurations.

Table 31: Background-control sidebands. “Signal fraction” is the signal contamination of the (CV-weighted neutrino) MC in the region; all four are strongly signal-depleted. $N_{\text{data}}/N_{\nu\text{MC}}$ is the pseudo-data to neutrino-MC ratio: the high-statistics regions (CC0 π , π^0) close to unity across FHC/RHC/combined, validating the per-run weighting and POT-scaling machinery; the low-statistics regions (multi- π , cosmic) scatter within Poisson expectation ($\sqrt{N} \approx 6\text{--}10\%$). The background-model test against the full stack (with EXT and dirt) is performed at control-region unblinding with beam-on data.

Sideband	Signal frac. (of νMC)	$N_{\text{data}}/N_{\nu\text{MC}}$		
		FHC	RHC	Combined
CC0 π	$\sim 10\%$	1.00	1.00	1.00
multi- π	$\sim 23\%$	1.03	1.11	1.08
π^0	$\sim 10\%$	0.99	1.02	1.01
cosmic	$\sim 8\%$	0.94	1.24	1.10

Yields and composition. Table 32 gives the full breakdown for the three configurations. The regions differ enormously in size and in what they are sensitive to, which is what makes them complementary rather than redundant:

- **CC0 π** is by far the largest ($\sim 12\text{k}$ pseudo-data events per beam, $\sim 24\text{k}$ combined). Its neutrino component is dominated by the quasi-elastic and $2p2h$ modes that feed the signal region only through pion *absorption* in the final state. It is therefore the natural handle on the FSI model that converts $\text{CC } 1\pi^\pm$ into $\text{CC } 0\pi$ and vice versa — the largest single migration path into and out of the signal definition (§4).
- π^0 is the second largest (6.5–14.7k). It is enriched in $\text{CC}/\text{NC } \pi^0$ events, whose showers are precisely what the shower veto (Cut 8) is designed to remove; it therefore tests the veto efficiency and the π^0 production normalisation together. These are the events that leak into the signal region when a photon is missed.
- **multi- π** is small (127–303 events) but has the highest signal contamination of the four ($\sim 23\%$), because a true multi-pion event with one pion lost is exactly the wrong-topology background the selection fights. Its size is what limits its constraining power at present.
- **cosmic** is the only region where the beam-off (EXT) contribution *dominates* the prediction: 308 EXT against 68 neutrino-MC events for FHC, a ratio of 4.5, against 0.8 in $\text{CC } 0\pi$. It is consequently the region that constrains the cosmic normalisation, and the one whose interpretation changes most at unblinding, since the current pseudo-data contain no EXT component at all.

What the present numbers do and do not show. With Poisson-thrown neutrino-MC pseudo-data, the ratios of Table 31 test the *machinery* — per-run POT scaling, CV weighting, the selection logic in each inverted region — and not the background model. That test passes: the two high-statistics regions sit at $N_{\text{data}}/N_{\nu\text{MC}} = 1.00\text{--}1.02$ in all three configurations, which is a non-trivial check of the per-run normalisation that caused the difficulties described in §A. The apparent deviations in the small regions are statistical: with 69 cosmic-sideband events for FHC, a $\sqrt{N}$ fluctuation is 12%, so the observed 0.94 and the RHC 1.24 (on 100 events, 10%) are within about one and two standard deviations respectively.

Use at unblinding. Once the control regions are opened with beam-on data the same table becomes a genuine background test, and the statistical precision available is set by the yields above: the $\text{CC } 0\pi$ and π^0 regions would constrain their dominant background normalisations to roughly 1% and 1.5% statistically for the combined exposure, well below the $\sim 20\text{--}30\%$ cross-section systematic they are intended to test. The cosmic region would constrain the EXT normalisation to $\sim 8\%$, comparable to the 3–5% trigger-ratio uncertainty already assigned. The multi- π region, at ~ 300 combined events, gives $\sim 6\%$ — useful as a cross-check but not competitive with the others. The intended procedure is to unblind the control regions first, use them to constrain or reweight the dominant background components, and only then open the signal region.

Table 32: Sideband yields at full exposure. “Signal” is the signal contamination of the neutrino MC, “ ν bkg” the remaining neutrino MC, EXT the beam-off (cosmic) sample scaled by the trigger ratio, and dirt the out-of-cryostat overlay. The pseudo-data column is a Poisson throw of the neutrino MC only, so it should be compared with Signal + ν bkg; EXT and dirt are listed for scale and enter the comparison only with real beam-on data.

Beam	Sideband	Pseudo-data	Signal	ν bkg	EXT	Dirt	Sig. frac.
FHC	CC0 π	11 887	1 122	10 736	8 939	54	9.5%
FHC	multi- π	127	30	94	42	0.1	24.0%
FHC	π^0	6 555	661	5 931	5 148	12	10.0%
FHC	cosmic	69	6	68	308	2.2	8.1%
RHC	CC0 π	12 261	1 204	11 034	9 281	42	9.8%
RHC	multi- π	176	35	124	43	0.1	22.0%
RHC	π^0	8 139	778	7 234	5 346	9	9.7%
RHC	cosmic	100	6	74	320	1.7	8.0%
Comb	CC0 π	24 148	2 327	21 770	18 219	54	9.7%
Comb	multi- π	303	65	217	85	0.1	22.9%
Comb	π^0	14 694	1 439	13 165	10 494	12	9.9%
Comb	cosmic	169	12	142	629	2.2	8.0%

9.4 Control regions for the proton-tagged subsample

The proton-tagged selection inherits its cut logic from CC1mu1piXp, and with it the four sideband definitions of §9.3: the reprocessed proton-tagged files already carry CC1mu1pi1p_sb_cc0pi, _sb_multi1pi, _sb_pi0 and _sb_cosmic. No new definitions or reprocessing are required to exercise the background model of the W/TKI measurement; the same four regions can be evaluated either as inherited (the inclusive phase space) or with a reconstructed tagged proton additionally required, which places them in the same phase space as the measurement itself. Table 33 gives both.

Three things follow from the numbers. First, all four regions are *more* signal-depleted for the proton-tagged signal than they were for the inclusive one: 6.0–7.3% signal contamination in CC0 π and π^0 , against 9.5–10.0% for the same regions in the inclusive phase space and a 48% purity in the signal region, because requiring a proton does not make a pion-less or π^0 event look like signal. Second, the proton requirement costs about a third of the statistics (CC0 π : 11 887 $\rightarrow$ 8 268 for FHC) but suppresses the beam-off contribution by roughly a factor six (8 760 $\rightarrow$ 1 418): cosmic-induced events rarely contain a reconstructed proton above threshold, so the proton tag is itself a strong cosmic rejector. This is worth noting for the systematic budget — the EXT normalisation matters far less in the proton-tagged measurement than in the inclusive one. Third, the pseudo-data to neutrino-MC ratios remain consistent with unity in every region and configuration (0.956–1.158 with the tag), so the per-run normalisation machinery is validated in the proton-tagged phase space as well.

What these regions cannot do. All four *apply* the proton tag rather than invert it, so none of them constrains the proton *mis-tag* background — the $\sim 13\%$ of tagged protons that are really a pion or another track, which smears $W_{\pi p}$ and the TKI variables directly. Testing that requires a region defined by inverting the proton PID (selecting candidates that pass the momentum and containment requirements but fail the LLR cut), which does need a new branch in CC1mu1pi1p and a reprocess. Slicing a measured W/TKI observable is *not* an alternative: because the signal definition is inclusive in those variables, cutting on them splits the signal region rather than depleting it. Explicitly, on the Run-1 FHC overlay the signal region has 53.7% purity and the slices $W_{\pi p} > 1.5$ GeV, $W_{\pi p} < 1.15$ GeV, $p_n > 0.8$ GeV/c and $\delta p_T > 0.7$ GeV/c give 50.0%, 48.9%, 47.3% and 52.0% respectively — no depletion at all. Even the kinematically forbidden

region $W_{\text{had}} < m_\pi + m_p$ retains 53.4% signal, because the calorimetric mass is systematically under-reconstructed and true signal populates it through resolution.

Table 33: Sideband yields for the proton-tagged subsample at full exposure, evaluated with the inherited definitions both as-is and with a tagged proton additionally required. “Sig.” is the proton-tagged signal contamination of the neutrino MC. $N_{\text{data}}/N_{\nu\text{MC}}$ compares the Poisson-thrown pseudo-data with the CV-weighted neutrino MC; EXT and dirt are listed for scale.

Beam	Region	N_{data}	Signal	ν bkg	EXT	Sig. frac.	$N_{\text{data}}/N_{\nu\text{MC}}$
<i>inherited definitions (inclusive phase space)</i>							
FHC	CC0 π	11 887	1 122	10 736	8 760	9.5%	1.002
FHC	multi- π	127	30	94	41	24.0%	1.030
FHC	π^0	6 555	660	5 931	5 045	10.0%	0.994
FHC	cosmic	69	6.0	68	302	8.1%	0.937
RHC	CC0 π	12 261	1 204	11 034	9 095	9.8%	1.002
RHC	multi- π	176	35	124	42	22.0%	1.109
RHC	π^0	8 139	778	7 234	5 239	9.7%	1.016
RHC	cosmic	100	6.4	74	314	8.0%	1.239
Comb	CC0 π	24 148	2 326	21 770	17 855	9.7%	1.002
Comb	multi- π	303	64	217	83	22.9%	1.075
Comb	π^0	14 694	1 439	13 165	10 284	9.9%	1.006
Comb	cosmic	169	12	142	616	8.0%	1.095
<i>with a tagged proton required (measurement phase space)</i>							
FHC	CC0 π	8 268	484	7 636	1 418	6.0%	1.018
FHC	multi- π	54	9.8	46	17	17.7%	0.978
FHC	π^0	3 988	292	3 730	1 162	7.3%	0.991
FHC	cosmic	24	1.7	23	5.8	6.7%	0.956
RHC	CC0 π	8 287	499	7 812	1 473	6.0%	0.997
RHC	multi- π	83	14	58	18	19.8%	1.143
RHC	π^0	5 031	332	4 487	1 207	6.9%	1.044
RHC	cosmic	36	2.3	29	6.0	7.5%	1.158
Comb	CC0 π	16 555	983	15 448	2 891	6.0%	1.008
Comb	multi- π	137	24	104	36	18.9%	1.071
Comb	π^0	9 019	624	8 218	2 370	7.1%	1.020
Comb	cosmic	60	4.0	52	12	7.1%	1.067

10 Proton-tagged subsample: hadronic invariant mass and transverse kinematic imbalance

Requiring an identified proton in the final state, in addition to the muon and charged pion, reconstructs the full visible hadronic system and opens two classes of observable inaccessible to the inclusive CC $1\mu 1\pi Xp$ measurement: the *hadronic invariant mass*, which probes the $\Delta(1232)$ and higher resonances directly, and the *transverse kinematic imbalance* (TKI) between the muon and the hadronic system, which is sensitive to Fermi motion and final-state interactions. This subsample is implemented as a dedicated selection, **CC1 μ 1 π 1p**, that inherits the full CC $1\mu 1\pi Xp$ selection (§6) and additionally requires a leading proton candidate — the most proton-like track (log-likelihood-ratio PID score below the proton cut), contained in the fiducial volume and above the 0.3 GeV/ c tracking threshold. The signal is correspondingly narrowed to the inclusive signal plus at least one true proton above 0.3 GeV/ c . On the Run-1 FHC sample this yields a selection efficiency of 11.5% and a purity of 52% (versus 17.5% and $\sim 56\%$ for the inclusive selection); the proton tag trades signal statistics for a fully reconstructed hadronic final state.

10.1 Selection performance and cut-flow

The full-exposure cut-flow of the proton-tagged subsample is the inherited $\text{CC } 1\mu 1\pi Xp$ sequence followed by the leading-proton identification (“IdentProton”), which defines the subsample. Table 34 gives the per-category, per-cut breakdown for the three beam configurations and Figure 25 the corresponding stacks; both use the same per-run POT scaling as the inclusive cut-flow (Table 9), and “signal” here is the proton-tagged signal (inclusive signal plus a true proton above 0.3 GeV/ c). The proton-identification cut is the key discriminator: at FHC it rejects $\sim 64\%$ of the neutrino background while retaining $\sim 71\%$ of the signal, lifting the purity from 33% at the inclusive final cut (Nonprotons) to 48% in the proton-tagged sample.




Figure 25: Proton-tagged ($\text{CC1}\mu 1\pi 1p$) cut-flow yields at full exposure (log scale; MC-prediction stack, Okabe-Ito: signal blue, background orange, EXT grey, dirt green). The last stage, IdentProton, is the leading-proton identification that defines the subsample.

Table 34: Proton-tagged (CC1 μ 1 π 1 p) cut-flow yields at full exposure for FHC, RHC and combined (MC prediction; data blind). Stages None–Nonprotons are the inherited inclusive selection with the proton-tagged signal definition; the final IdentProton stage adds the leading-proton identification. Signal/Bkg from the ν_μ overlays POT-scaled per run, EXT by the data/EXT trigger ratio, dirt from the GENIE NuMI dirt overlay. Pred. = Signal + Bkg + EXT + Dirt. Eff. is the cumulative signal efficiency relative to the generated signal (the None stage, where all true signal is counted); Pur. = Signal/Pred.

Cut	Signal	Bkg MC	EXT	Dirt	Pred.	Eff. [%]	Pur. [%]
<i>FHC (Runs 1,2,4,5; 8.857×10^{20} POT)</i>							
None	3 387	281 372	24 410 022	36 142	24 730 924	100.0	0.01
InFiducialVol	2 834	63 944	1 041 637	1 890	1 110 306	83.7	0.26
Topological	2 238	29 303	35 184	204	66 929	66.1	3.3
MuonCandidate	2 031	20 092	15 516	137	37 776	60.0	5.4
ContainedPion	852	2 637	1 376	25	4 890	25.2	17.4
MuonIn3Planes	812	2 499	1 145	17	4 473	24.0	18.2
PionIn3Planes	754	2 255	872	7.0	3 888	22.3	19.4
ShowerCut	632	1 313	569	6.2	2 521	18.7	25.1
OpeningAngle	628	1 226	190	2.0	2 045	18.5	30.7
Nonprotons	568	1 049	130	1.3	1 749	16.8	32.5
IdentProton	401	373	53	0.4	828	11.8	48.4
<i>RHC (Runs 1,2,3,4a,4b,4c; 1.108×10^{21} POT)</i>							
None	3 638	306 840	25 344 643	28 032	25 683 152	100.0	0.01
InFiducialVol	3 064	71 162	1 081 520	1 466	1 157 212	84.2	0.26
Topological	2 455	33 134	36 532	158	72 278	67.5	3.4
MuonCandidate	2 234	22 769	16 110	106	41 219	61.4	5.4
ContainedPion	935	3 291	1 429	19	5 674	25.7	16.5
MuonIn3Planes	892	3 119	1 189	13	5 213	24.5	17.1
PionIn3Planes	831	2 835	905	5.5	4 577	22.9	18.2
ShowerCut	687	1 544	591	4.8	2 827	18.9	24.3
OpeningAngle	683	1 444	197	1.5	2 325	18.8	29.4
Nonprotons	611	1 217	136	1.0	1 964	16.8	31.1
IdentProton	436	443	55	0.3	934	12.0	46.6
<i>Combined (1.994×10^{21} POT)</i>							
None	7 025	588 212	49 755 076	36 142	50 386 456	100.0	0.01
InFiducialVol	5 898	135 106	2 123 174	1 890	2 266 069	84.0	0.26
Topological	4 693	62 436	71 717	204	139 050	66.8	3.4
MuonCandidate	4 265	42 860	31 627	137	78 889	60.7	5.4
ContainedPion	1 787	5 928	2 805	25	10 545	25.4	16.9
MuonIn3Planes	1 704	5 618	2 333	17	9 672	24.3	17.6
PionIn3Planes	1 585	5 091	1 777	7.0	8 460	22.6	18.7
ShowerCut	1 320	2 857	1 161	6.2	5 344	18.8	24.7
OpeningAngle	1 310	2 670	387	2.0	4 369	18.7	30.0
Nonprotons	1 179	2 266	266	1.3	3 712	16.8	31.8
IdentProton	837	816	109	0.4	1 762	11.9	47.5

Proton-tag performance. Beyond the yields, the leading-proton identification (log-likelihood-ratio PID score below the proton cut, $p_p \geq 0.3$ GeV/ c , contained) is a reliable tag. Truth-matched on the FHC ν_μ overlay, 87% of the tagged events contain an energetic true proton ($p_p > 0.3$ GeV/ c) — a $\sim 13\%$ mis-tag rate, dominated by a charged pion or other track being taken for the proton — and 52% are the full $\mu\pi p$ signal topology. The reconstructed proton momentum, from the proton-hypothesis range kinetic energy, carries a -9% bias (underestimate) and a 21% resolution. In cross-section terms the tag rejects $\sim 64\%$ of the neutrino background for a $\sim 30\%$ signal loss, improving the signal-to-background figure of merit $S/\sqrt{B}$ by $\sim 17\%$; the residual mis-tag and

momentum smearing are folded into the response matrix and corrected by the unfolding.

Eleven differential observables are measured in the proton-tagged subsample: the six hadronic-mass and transverse-kinematic-imbalance (W/TKI) variables specific to this selection, plus the five muon/pion kinematic observables of the inclusive analysis (p_μ , p_π , $\cos\theta_\mu$, $\cos\theta_\pi$, $\theta_{\mu\pi}$) re-measured on the proton-tagged phase space so that the two selections share a common set of kinematic distributions. The two hadronic-mass variables are the *directly measured* πp invariant mass, $W_{\pi p} = [(E_\pi + E_p)^2 - (\vec{p}_\pi + \vec{p}_p)^2]^{1/2}$, which makes no assumption on the neutrino direction, and the *calorimetric* hadronic mass $W_{\text{had}} = [m_p^2 + 2m_p(E_{\text{cal}} - E_\mu) - Q^2]^{1/2}$ built from the reconstructed calorimetric energy. The four TKI observables — the transverse-momentum imbalance δp_T , its two angles $\delta\alpha_T$ and $\delta\phi_T$, and the inferred struck-nucleon momentum p_n — are computed from the muon versus the summed ($\pi + p$) hadronic system using the framework STV tool. Proton momenta are obtained from the proton-hypothesis range-based kinetic energy. The six W/TKI observables are each measured in six equal-population bins — matching the granularity of the inclusive analysis, now that the per-run POT normalisation (§A) removes the background-subtraction bias that had previously forced a coarser three-bin binning for the lower-statistics reverse-horn-current sample — while the five inherited kinematic observables use the inclusive-analysis binning (four to seven bins). All bin edges are set from the validated true-signal distributions (Table 35).

Reconstruction from input quantities

All six observables are built from the three candidate tracks (muon, charged pion, proton) using only reconstructed quantities stored in the input ntuples: the track direction cosines $\hat{n} = (\text{track_dirx}, \text{track_diry}, \text{track_dirz})$, the muon-hypothesis range and multiple-scattering momenta ($\text{track_range_mom_mu}$, track_mcs_mom_mu) and the proton-hypothesis calorimetric kinetic energy ($\text{track_kinetic_energy_p}$).

Candidate three-momenta. The muon momentum magnitude is range-based when the muon track ends inside the fiducial volume, $p_\mu = \text{track_range_mom_mu}[\mu]$, and multiple-Coulomb-scattering based otherwise, $p_\mu = \text{track_mcs_mom_mu}[\mu]$. The charged pion, being minimum-ionising, is assigned the muon-hypothesis range momentum $p_\pi = \text{track_range_mom_mu}[\pi]$. The proton momentum follows from its calorimetric kinetic energy under the proton hypothesis, $p_p = \sqrt{T_p^2 + 2m_p T_p}$ with $T_p = \text{track_kinetic_energy_p}[p]$. Each three-momentum is $\vec{p}_i = p_i \hat{n}_i$, and the energies are $E_\mu = \sqrt{p_\mu^2 + m_\mu^2}$, $E_\pi = \sqrt{p_\pi^2 + m_\pi^2}$, $E_p = T_p + m_p$ (masses $m_\mu = 105.7$, $m_\pi = 139.6$, $m_p = 938.3$ MeV).

Directly measured mass $W_{\pi p}$. The invariant mass of the pion-proton four-vector sum, using no neutrino-direction assumption:

$$W_{\pi p} = \sqrt{(E_\pi + E_p)^2 - |\vec{p}_\pi + \vec{p}_p|^2}.$$

Transverse kinematic imbalance. The TKI observables treat the muon against the visible hadronic system $\vec{p}_{\text{had}} = \vec{p}_\pi + \vec{p}_p$, $E_{\text{had}} = E_\pi + E_p$. With the components transverse to the beam ($\hat{z}$) written $\vec{p}_T^\mu$ and $\vec{p}_T^{\text{had}}$,

$$\delta p_T = |\vec{p}_T^\mu + \vec{p}_T^{\text{had}}|, \quad \delta\alpha_T = \arccos \frac{-\vec{p}_T^\mu \cdot (\vec{p}_T^\mu + \vec{p}_T^{\text{had}})}{|\vec{p}_T^\mu| \delta p_T}, \quad \delta\phi_T = \arccos \frac{-\vec{p}_T^\mu \cdot \vec{p}_T^{\text{had}}}{|\vec{p}_T^\mu| |\vec{p}_T^{\text{had}}|},$$

both angles in degrees folded onto $[0, 180]$. For a stationary nucleon with no final-state interactions $\delta p_T \rightarrow 0$; Fermi motion broadens it and FSI drive $\delta\alpha_T \rightarrow 180^\circ$.

Calorimetric mass W_{had} and struck-nucleon momentum p_n . The calorimetric available energy is $E_{\text{cal}} = E_\mu + (E_{\text{had}} - m_p) + E_b$ with the argon removal energy $E_b = 24.78$ MeV, giving the

momentum transfer $Q^2 = -[(0, 0, E_{\text{cal}}, E_{\text{cal}}) - p^\mu]^2$ and

$$W_{\text{had}} = \sqrt{m_p^2 + 2m_p(E_{\text{cal}} - E_\mu) - Q^2}$$

(returned as 0 where reconstruction fluctuations make the argument negative). The inferred initial-state nucleon momentum combines the transverse imbalance with a longitudinal component p_L from energy-momentum balance, $p_n = \sqrt{\delta p_T^2 + p_L^2}$. The transverse-imbalance and calorimetric quantities are evaluated with the framework STV tool (`STVTools::CalculateSTVs`, option 1); it was formulated for a single-proton hadronic system and is applied here to the summed (πp) system with the proton rest mass, an approximation adequate for the leading-order kinematics.

Table 35: Binning of the six W/TKI proton-tagged observables (six equal-population bins each), chosen from the validated true-signal distributions. The first and last bins are open so the full signal phase space is partitioned. $W_{\pi p}$ brackets the $\Delta(1232)$ peak; the $\delta\alpha_T$ and $\delta\phi_T$ binnings are finer toward 180° , where final-state interactions pile up. The five inherited kinematic observables (p_μ , p_π , $\cos\theta_\mu$, $\cos\theta_\pi$, $\theta_{\mu\pi}$) use the inclusive-analysis binning of §C unchanged.

Observable	Definition	Bin edges
$W_{\pi p}$	$M(\pi p)$, direct	1.08, 1.19, 1.23, 1.27, 1.34, 1.47, 2.00 GeV/ c^2
W_{had}	calorimetric	0, 0.78, 1.08, 1.21, 1.31, 1.44, 1.90 GeV/ c^2
δp_T	$ \vec{p}_T^\mu + \vec{p}_T^{\text{had}} $	0, 0.43, 0.59, 0.72, 0.88, 1.17, 1.60 GeV/ c
$\delta\alpha_T$	TKI angle	0, 80, 120, 142, 158, 170, 180°
$\delta\phi_T$	TKI angle	0, 37, 63, 84, 107, 139, 180°
p_n	inferred nucleon momentum	0, 0.52, 0.67, 0.81, 0.98, 1.25, 1.70 GeV/ c

The pipeline is validated end-to-end on FHC with a fake-data closure: the Poisson-thrown central-value pseudo-data are unfolded with the same Wiener-SVD procedure (§8) and compared to the A_C -smeared truth. Table 36 shows the result. The bin-by-bin χ^2/ndf is well below unity for every observable, confirming the unfolded pseudo-data are statistically consistent with truth and that the response, background subtraction and normalisation are self-consistent. The integrated-cross-section ratios (0.94–1.19) reflect the same non-norm-preserving Wiener-SVD additional smearing A_C seen in the inclusive measurement (§8.11), largest for $\delta\alpha_T$ and $\theta_{\mu\pi}$. The systematic treatment here includes the flux, cross-section and reinteraction terms from the MC weight universes; the detector-variation term is deferred pending a dedicated reprocessing of the detector-variation samples through this selection.

The reverse-horn-current measurement is complete at the same finer binning (Table 37). All eleven RHC observables close positive, with σ_{int} clustered in $0.41\text{--}0.58 \times 10^{-38} \text{ cm}^2/\text{Ar}$ — just below the FHC band (0.50–0.67), as expected for the $\bar{\nu}_\mu$ -enriched RHC flux — and $\chi^2/\text{ndf} \lesssim 1$ for almost every observable.

Reaching this point required resolving a normalisation subtlety that is worth recording, because it produced spurious *negative* cross sections for this subsample at an intermediate stage. The processed ntuples use two different POT conventions: the files behind the inclusive selection store the *total* run MC POT duplicated in every file of a multi-file run, whereas the proton-tagged files store each file’s *own* POT. Normalising by the individual file POT is therefore correct for the first set but scales every split of a multi-file run to the full run exposure in the second, inflating the simulated background by the split multiplicity ($\sim 3\times$ for RHC) and driving the background-subtracted spectrum negative. Normalising by the summed per-file POT has the mirror problem. The analysis now normalises by the sum of the *distinct* per-file POT values, which reduces to the correct run total under either convention; FHC, with one MC file per run, is unaffected by the choice.

The combined-sample measurement is in production and will be added here.

Table 36: FHC fake-data closure for the eleven proton-tagged observables at the finer binning. σ_{int} is the integrated (unfolded) cross section over the measured phase space ($10^{-38} \text{ cm}^2/\text{Ar}$); “unf./truth” is the ratio of integrals; χ^2 is the bin-by-bin comparison of unfolded pseudo-data to A_C -smeared truth (ndf = number of bins). The upper block is the six W/TKI observables (six bins each); the lower block is the five kinematic observables inherited from the inclusive analysis. The small χ^2/ndf validates the chain; the integral ratios are the A_C smearing effect, not a bias.

Observable	σ_{int}	unf./truth	χ^2/ndf
$W_{\pi p}$	0.62	1.11	0.09/6
W_{had}	0.60	1.09	0.09/6
δp_T	0.61	1.12	0.46/6
$\delta \alpha_T$	0.67	1.19	0.26/6
$\delta \phi_T$	0.63	1.11	0.09/6
p_n	0.52	0.94	0.09/6
p_μ	0.56	1.16	0.39/7
p_π	0.60	1.10	2.58/5
$\cos \theta_\mu$	0.52	1.11	0.56/5
$\cos \theta_\pi$	0.54	1.13	0.16/4
$\theta_{\mu\pi}$	0.50	1.19	0.84/5

Figure 26 shows the six FHC W/TKI differential cross sections at the finer six-bin binning, overlaid with all four generator predictions. The generator W/TKI predictions were regenerated at this binning (GENIE and GiBUU natively, NuWro and NEUT from their event vectors inside the SL7 container); the rebinning conserves each generator’s integral to the previous three-bin value (GENIE 0.501, NuWro 0.646, GiBUU 0.735, NEUT $0.783 \times 10^{-38} \text{ cm}^2/\text{Ar}$). All curves — data, tune and generators — are A_C -smeared, so the comparison is made consistently in the measurement space (§10.2). The predictions are built at truth level from each generator’s event-level final state — GENIE `gst` events, GiBUU `FinalEvents`, and the NuWro and NEUT event vectors reprocessed in their respective SL7-container environments — by applying the inclusive signal plus a leading true proton above 0.3 GeV/ c and evaluating the six observables with the same formulas as the selection (§10.1), then combining the ν_μ and $\bar{\nu}_\mu$ samples over the FHC flux to a per-argon cross section. The four generators span a factor of ~ 1.6 in the integrated proton-tagged cross section — GENIE lowest (0.50), then NuWro (0.65), GiBUU (0.74) and NEUT highest ($0.78 \times 10^{-38} \text{ cm}^2/\text{Ar}$) — bracketing the data (~ 0.58), and they differ most in the low- δp_T / high- p_n region sensitive to the proton kinematics and final-state interactions. As throughout, the unfolded fake data sit above the A_C -smeared truth by the Wiener-SVD normalisation factor.

Table 37: RHC fake-data closure for the eleven proton-tagged observables at the finer binning, columns as in Table 36. All observables close positive with the convention-agnostic MC POT normalisation described above; σ_{int} sits just below the FHC band, consistent with the $\bar{\nu}_\mu$ -enriched RHC flux.

Observable	σ_{int}	unf./truth	χ^2/ndf
$W_{\pi p}$	0.578	1.18	0.44/6
W_{had}	0.500	1.10	2.19/6
δp_T	0.493	1.20	1.00/6
$\delta\alpha_T$	0.567	1.20	0.44/6
$\delta\phi_T$	0.568	1.17	1.40/6
p_n	0.548	1.30	3.07/6
p_μ	0.531	1.23	0.94/7
p_π	0.572	1.23	0.71/5
$\cos\theta_\mu$	0.512	1.22	2.23/5
$\cos\theta_\pi$	0.444	1.06	3.74/4
$\theta_{\mu\pi}$	0.410	1.15	0.86/5

Table 38: Combined FHC+RHC fake-data closure for the eleven proton-tagged observables at the finer binning, columns as in Table 36. All eleven close positive and every χ^2 is comfortably below its number of bins. This completes the proton-tagged measurement in all three beam configurations. Note that the combined result is an independent extraction, not a fluence-weighted average of the FHC and RHC ones, so it is not required to lie between them; see the discussion below.

Observable	σ_{int}	unf./truth	χ^2/ndf
$W_{\pi p}$	0.564	1.08	0.07/6
W_{had}	0.504	1.02	1.30/6
δp_T	0.497	1.10	0.71/6
$\delta\alpha_T$	0.667	1.28	0.86/6
$\delta\phi_T$	0.573	1.07	0.05/6
p_n	0.556	1.18	1.00/6
p_μ	0.523	1.15	0.65/7
p_π	0.582	1.18	2.49/5
$\cos\theta_\mu$	0.529	1.17	0.52/5
$\cos\theta_\pi$	0.454	1.03	2.17/4
$\theta_{\mu\pi}$	0.447	1.15	0.21/5

With all three configurations in hand, the combined σ_{int} lies between the FHC and RHC values for seven of the eleven observables. The four exceptions are small and worth stating explicitly rather than glossing: p_n sits 6.9% above the FHC value, $\cos\theta_\mu$ 1.7% above it, and $W_{\pi p}$ and p_μ fall 2.4% and 1.5% below the RHC value. The combined measurement is *not* required to bracket the two single-mode ones: it is an independent extraction with its own response matrix, its own regularisation and its own A_C , not a fluence-weighted average of the FHC and RHC results. Excursions at the few-per-cent level are therefore expected, and all four are far inside the $\sim 30\%$ systematic uncertainty.

The spread across observables within a configuration — 0.45 to 0.67 for the combined set — is likewise not a physics spread but the residual of the additional smearing discussed in the next section: each observable’s A_C rescales the integral by a different factor, so the eleven numbers

are not eleven independent measurements of the same quantity.

10.2 Additional smearing in the proton-tagged observables

Because the Wiener–SVD additional-smearing matrix A_C is not norm-preserving, each observable’s predictions are rescaled by a different factor when they are brought into the measurement space. Table 39 gives that factor, measured as the ratio of the A_C -smeared to the truth-level generator integral, for the six proton-tagged observables. It is a property of the response and the binning, not of the model: GENIE and GiBUU agree to a few per cent throughout.

The two hadronic-mass variables are the *least* affected in the whole analysis (0.77–0.88). This is worth stating explicitly because $W_{\pi p}$ is the sharpest distribution measured here — the ratio of its largest to smallest bin is ~ 36 , against ~ 19 for $\cos\theta_\mu$, whose forward peak is suppressed by a factor ~ 0.33 (§9.2). Sharpness alone therefore does not drive the suppression: A_C removes structure that the *response* cannot resolve, and $W_{\pi p}$ is built from directly measured pion and proton momenta with good resolution and equal-population bins that track the $\Delta(1232)$ peak, so the peak survives. No reparameterisation of W is required.

Two observables do carry a caveat. δp_T is strongly suppressed (0.43–0.45), closer to the $\cos\theta_\mu$ regime than to W ; and p_n is *enhanced* by a factor ~ 1.7 . Both effects are applied identically to data, tune and generators, so the closure is unaffected, but any model comparison read off the δp_T and especially the p_n panels is distorted by these factors and should not be interpreted as a direct statement about the truth-level predictions until the binning and response for those two observables have been revisited.

Table 39: Additional-smearing factor A_C for the proton-tagged observables (FHC), measured as the ratio of the A_C -smeared to the truth-level generator integral. The inclusive observables are given for scale.

Observable	GENIE	GiBUU	Comment
W_{had}	0.88	0.86	least smeared in the analysis
$W_{\pi p}$	0.78	0.77	sharpest distribution, yet barely affected
$\delta\alpha_T$	0.63	0.62	
$\delta\phi_T$	0.61	0.60	
δp_T	0.43	0.45	strongly suppressed
p_n	1.71	1.65	<i>enhanced</i>
<i>Inclusive, for scale: p_μ 0.64; θ_μ 0.76; $\cos\theta_\mu$ 0.33</i>			




Figure 26: Proton-tagged (CC1mu1pi1p) differential cross sections for FHC: the πp invariant mass $W_{\pi p}$, the calorimetric hadronic mass W_{had} , and the four transverse-kinematic-imbalance observables δp_T , $\delta\alpha_T$, $\delta\phi_T$ and p_n . Unfolded fake data (black points) are compared to the A_C -smeared fake-data truth (grey dashed), the MicroBooNE tune (black), and the GENIE (blue), GiBUU (green), NEUT (purple) and NuWro (orange) predictions built for the same proton-tagged phase space. $W_{\pi p}$ falls from the $\Delta(1232)$ threshold region; $\delta\alpha_T$ rises toward 180° , the signature of final-state interactions.

11 Systematic Uncertainties

Full systematic uncertainty propagation has not yet been performed. Table 40 lists the sources and their estimated impacts.

Table 40: Systematic uncertainty sources. “Not quantified” means the branch is readable but propagation code has not been written.

Source	Branch / method	Impact
Flux (PPFX multisims)	<code>weightsFlux</code>	$\sim 10\text{--}15\%$ norm.
GENIE cross sections	<code>weightsGenie</code>	$\sim 20\text{--}30\%$ (model)
Hadron re-interaction	<code>weightsReint</code>	$\sim 5\text{--}10\%$
Detector response	Dedicated samples	$\sim 5\text{--}15\%$
EXT normalisation	Trigger-count ratio	$\sim 3\text{--}5\%$
POT counting	Beam toroids	$\sim 2\%$
Flux integral (absolute)	PPFX integral	Blocks absolute σ
Statistical (427 data events)	—	$\sim 5\%$ on σ_{int}

12 Known Issues and Required Follow-up

Table 41: Outstanding issues and their impact on the final result.

Issue	Status	Impact
FLUX_PER_POT confirmed	Resolved	$\Phi = 6.81159 \times 10^{-10} \text{ cm}^{-2} \text{ POT}^{-1}$ from NuMI flux histograms
Unfolding covariance (stat-only $\rightarrow$ full)	Resolved	Framework <code>PredTotal</code> +data Poisson imported; removed $\sim 26\%$ bias
EXT subtraction on NuWro fake data	Resolved	Cancelled via <code>fake_data_ext_cancel</code> ; re-enable for real data
$\theta_{\mu\pi}$ binning (10 vs 9, C4)	Resolved	Custom aligned to framework 9-bin scheme
Multisim systematics not propagated (custom)	Not started	Imported from framework for the covariance; full custom propagation pending
Negative tail bins	Resolved	All bins now positive after EXT cancellation; no longer present
Runs 2, 4, 5 beam-on files not available	Data files needed	$\times 5.4$ exposure increase for the real-data measurement
Generator comparisons (GENIE, NuWro)	Blocked by missing AVX on compute nodes	NEUT comparison runnable locally
CRT veto branches (<code>crthitpe</code> , <code>crtveto</code>)	Dead (all-zero in NuMI ntuples)	Replaced by <code>bdt_cosmic</code> (disabled at default 0.0)
Official XSecAnalyzer NuWro closure	Done (all 5 obs)	Cross-validates the unfolding; NuWro $\approx 2\times$ GENIE (Sec. B)
Native RHC detector variations	Resolved	Run-4 RHC detVar (8 knobs) now used; replaces the Run-1 FHC stand-in (Sec. 9)
Run-5 FHC detVar Recomb2/SCE	Resolved	Were intrinsic- ν_e overlays (0% ν_μ); Run-4 FHC used as the stand-in for those two knobs
Combined FHC+RHC joint flux systematic	Resolved	Per-universe summing over the shared 600 PPFx universes is correct; the block was a detVar duplicate-file guard, now fixed to accumulate (Sec. 9)
Combined per-mode normalisation	Resolved	Per-mode <code>summed_pot</code> so each horn scales to its own data exposure, not MC POT (Sec. 9)

13 Analysis Pipeline

The primary measurement is produced by a single driver, `run_analysis.sh`, which chains the selection, the per-observable cross-section extraction, and the comparison/diagnostic stages:

```

selection/ccpi_selection.C -> histograms.root + unfolding_inputs.root
                             | (5 observables x 8 methods)
xsec/ccpi_xsec.C(obs,method) -> xsec_{obs}_{method}.root (40 files)
xsec/compare_unfolding.C -> compare_plots/
xsec/compare_generators.C -> gen_compare/

```

selection/event_display.C -> event_displays/

Each output file is validated for freshness and non-zero content after the step that produces it. The official-framework cross-check of Section B is driven separately (xsec_analyzer/run_nurow_observables.sh).

14 Binning Optimisation, Systematic Breakdown, and Threshold Validation

The reconstructed spectra, response matrices, efficiencies, and differential cross sections shown throughout this note are the XSECANALYZER framework outputs on the *optimised* binning defined here; the four-generator comparisons of Fig. 28 onward likewise use it.

14.1 Momentum-threshold validation

The two signal-definition momentum thresholds sit at the selection efficiency turn-on, measured against a threshold-free true signal (the CC ν_μ $1\mu 1\pi^\pm$ topology with the momentum/angle cuts removed):

Threshold	below	above	turn-on
$p_\mu > 0.15$ GeV/c	$\sim 0\text{--}1.5\%$	$\sim 9\text{--}12\%$	0.145 $\rightarrow$ 0.155
$p_\pi > 0.175$ GeV/c	$\sim 0\text{--}3\%$	$\sim 12\text{--}16\%$	0.165 $\rightarrow$ 0.185

Both are correctly placed — lowering them would add near-zero-efficiency phase space, raising them would discard measurable signal. The reconstructed pion momentum (muon-hypothesis range) is biased low and increasingly so with momentum (mean reco–true ≈ 0 at threshold, -0.15 GeV/c at 0.40) from the μ -vs- π mass hypothesis plus pion re-interaction; a pion-hypothesis mass correction is applied to the p_π reco bins.

14.2 Binning optimisation

Bins were designed to be at least the local RMS detector resolution wide (so the response stays well conditioned and the Wiener-SVD does not amplify the correlated flux/cross-section uncertainty) and to hold ≥ 45 selected-signal events (statistical floor). This lowers the total fractional uncertainty in every observable, chiefly through the statistical and unfolding-amplified terms; the flux/beamline term is correlated across bins and sets a floor:

Observable	bins	Total before	Total after
p_μ	22 $\rightarrow$ 7	33.2%	25.0%
p_π	5 $\rightarrow$ 5	27.7%	27.0%
$\cos\theta_\mu$	12 $\rightarrow$ 5	31.3%	26.7%
$\cos\theta_\pi$	12 $\rightarrow$ 5	48.5%	33.9%
$\theta_{\mu\pi}$	9 $\rightarrow$ 7	28.4%	23.5%

14.3 Systematic breakdown by source

The framework builds a covariance matrix for every systematic; the source-group decomposition method is illustrated below on a single configuration (the definitive per-configuration breakdowns for the full FHC, RHC, and combined exposures are Tables 28–30). The bin-averaged fractional uncertainty from each source group (on the optimised binning) is:

Table 42: Bin-averaged fractional systematic uncertainty (%) by source group on the optimised binning.

Source	p_μ	p_π	$\cos \theta_\mu$	$\cos \theta_\pi$	$\theta_{\mu\pi}$
Total	27.2	28.3	28.8	33.8	25.1
Cross section (GENIE)	8.6	8.5	9.1	9.5	8.1
Flux — PPFX	21.8	23.6	20.7	25.4	18.7
Flux — beamline geom.	2.3	2.2	2.3	3.3	1.6
Detector	9.6	10.2	14.0	14.6	10.5
Reinteraction	3.2	2.6	3.4	3.4	2.5
MC + data stats	8.6	6.5	9.4	11.9	9.3
POT + targets	2.7	2.9	2.7	3.0	2.4

The NuMI flux uncertainty has two parts, listed separately above: the PPFX hadron-production term (the dominant systematic, 19–26%) and the beamline-geometry term. The latter is built from the 12 geometry-variation weights (`Horn_2kA`, `Horn1/2.x/y_3mm`, `Beam_spot_*`, `Beam_shift_*`, `Horns*_water`, `Target_z_7mm`) injected into the numuMC overlay via `AddBeamlineGeometryWeights`, summed in quadrature as two-sided unisim variations. It is sub-dominant, 1.6–3.3% across the observables. (Building it required a framework fix: the covariance code dereferenced a null histogram when a systematic — here the beamline weights, present in the numuMC but not the dirt overlay — was absent from some MC sample; it now falls back to that sample’s central value, i.e. unit weight.) The per-bin breakdowns are shown in Fig. 27.




Figure 27: Bin-averaged fractional systematic uncertainty by source vs. observable, for FHC (left), RHC (centre) and combined (right) — the figure form of Tables 28–30. Colours use the Okabe-Ito colourblind-safe palette with distinct markers per source: total (black), flux (blue), detector (green), cross section (vermillion), reinteraction (purple), data stat (orange).

14.4 Anatomy of the flux term: genuine flux vs. unfolding amplification

The flux (PPFX) term dominates the budget, and at 17–26% it is larger than the $\sim 17\%$ normalisation uncertainty usually quoted for the NuMI flux. It is therefore worth asking how much of it is the flux model itself and how much is statistical amplification introduced by the unfolding. To separate the two we evaluate the flux covariance at two stages — on the reconstructed spectrum (before unfolding) and on the unfolded result — and, at each stage, both integrated (the fully correlated, normalisation-like number, from the sum of all covariance elements) and bin-averaged (the mean of the diagonal fractional errors, i.e. the number quoted in Table 42 and Fig. 27).

Table 43: Flux (PPFX) fractional uncertainty before and after unfolding. The reconstructed-level value is $\sim 17\%$ and flat across observables — the genuine flux-normalisation uncertainty. The unfolded value is that same $\sim 17\%$ amplified by the Wiener–SVD propagation through the finite-statistics response; the amplification tracks the data+MC statistical quality of each observable.

	p_μ	p_π	$\cos \theta_\mu$	$\cos \theta_\pi$	$\theta_{\mu\pi}$
Flux, reco (integrated / norm)	16.9	17.0	16.7	16.7	16.7
Flux, reco (bin-averaged)	17.7	16.4	16.6	17.0	17.0
Flux, unfolded (bin-averaged)	19.8	22.4	19.8	25.5	17.5
Amplification factor	1.12	1.37	1.20	1.49	1.03
Excess from unfolding (abs.)	+2.1	+6.0	+3.2	+8.5	+0.5
Data+MC stats (unfolded)	7.9	6.3	8.4	12.2	8.8

Three points follow. First, the *genuine* flux uncertainty is $\sim 17\%$ and identical for all five observables, as expected for a normalisation term. Second, it is not itself contaminated by Monte Carlo statistics: at reco level the integrated (correlated) and bin-averaged (diagonal) numbers agree to better than a percent, whereas statistical noise leaking into the covariance would inflate the diagonal well above the correlated norm. This is expected — the ~ 600 PPFX universes reweight the *same* events, so there is no independent per-universe statistical scatter. Third, the elevated unfolded values are the $\sim 17\%$ flux *amplified* by the Wiener–SVD propagation $A_C M V_{\text{flux}} M^\top A_C^\top$ through the finite-statistics response. It is at this step, not in the flux multisim, that the response statistics transfer into the flux term: the amplification is largest exactly where the statistics are thinnest ($\cos \theta_\pi$: stat. 12.2%, amplification $1.49\times$, $+8.5\%$) and negligible where the response is well conditioned ($\theta_{\mu\pi}$: $1.03\times$, $+0.5\%$). The lever for reducing the flux term is therefore the response conditioning (more overlay statistics, or the observable’s binning and regularisation strength), not the flux model; $\cos \theta_\pi$ — the noisiest observable across the entire budget — is where it would matter most.

14.5 Multi-run overlay and a resonance-model closure

Two facts frame what more statistics can and cannot do for the flux term. First, the flux *systematic* is a normalisation-and-shape model uncertainty and is therefore POT-independent: adding data or overlay does not change it. Second, Section 14.4 showed that the elevated *unfolded* flux is that $\sim 17\%$ model uncertainty amplified by the Wiener–SVD propagation through a statistically-limited response, and that the limitation is the *data* statistics, not the overlay. Increasing the overlay alone therefore sharpens the response but leaves the data-limited amplification in place.

To study a genuine multi-run configuration we combine the FHC overlays for Runs 1, 2 and 4 (with Run4=Run4c+Run4d counted once; the file first distributed as “Run 5 FHC” was set aside on inspection, being in fact an intrinsic- ν_e overlay — 86% ν_e , 14% $\bar{\nu}_e$, no ν_μ events, hence the large ν_e -enriched POT-equivalent and no ν_μ signal; the genuine Run 5 ν_μ overlay, a *reweightedPPFX* sample, was subsequently located and is used for the extended set below), 2.32×10^6 unique events, normalised to the FHC data POT 6.626×10^{20} (Section 14.6). Because $\text{CC}\pi^\pm$ is resonance-dominated, robustness to the resonant axial mass is tested with Poisson-thrown fake data reweighted to $M_A^{\text{RES}} = 1.12 \pm 0.22$ GeV ($\pm 1\sigma$; axial dipole ratio in the true interaction Q^2 , a $+8.1\% / -7.2\%$ rate distortion). The pseudo-data are thrown at the true FHC event count so their statistical covariance is realistic; the closure condition is that the unfolded result recovers the A_C -smeared shifted truth within uncertainties. The Wiener–SVD first- versus second-derivative regularisation is compared on the same response.

The closure succeeds for every observable, both signs of the shift, and both regularisations: the unfolded result recovers the A_C -smeared shifted truth with $\chi^2/\text{ndf} < 1$ throughout (Table 45), all p -values above 0.63 and most above 0.9. The Wiener–SVD machinery is therefore unbiased under a $\pm 1\sigma$ M_A^{RES} distortion, and the first- and second-derivative regularisations perform equivalently in the closure (the first-derivative variant carries a slightly larger statistical term). Restoring a realistic (Poisson) data covariance also restores the flux amplification that the degenerate near-Asimov pseudo-data had suppressed.

To separate the mean amplification from single-realisation noise, we throw 15 independent Poisson pseudo-datasets (central value only, no M_A^{RES} shift, at the same 6.626×10^{20} FHC exposure) and unfold each (Table 44). Averaged over throws, the unfolded flux stabilises at 24.7–27.1%, a mean amplification of $1.44\text{--}1.63\times$ over the $\sim 17\%$ floor, with a data-statistical term of 8–13%. This confirms that the amplification is driven by the data statistics, not the overlay, and is consistent with the mechanism of Section 14.4. The throw-to-throw scatter is small for the momenta (p_μ, p_π : $\pm 0.06\text{--}0.08\times$, flux $\pm 1.1\text{--}1.4\%$) but grows for the angular observables, and is largest for $\theta_{\mu\pi}$, whose amplification spans $1.27\text{--}2.02\times$ (flux 21.4–34.2%) across the 15 throws. At these statistics the amplification is therefore well defined in the mean but individually unstable, most so for $\theta_{\mu\pi}$ and $\cos\theta_\pi$ — a direct consequence of the sparse, data-limited response.

Table 44: Flux amplification averaged over 15 independent Poisson-thrown pseudo-datasets at the FHC data POT: mean $\pm$ standard deviation and min–max range over throws, for the unfolded flux (bin-averaged) and the amplification relative to the $\sim 17\%$ reco-level floor.

Observable	Amplification (mean $\pm$ sd)	Amp. range	Unfolded flux	Flux range
p_μ	$1.44 \pm 0.08\times$	[1.27, 1.60]	$24.7 \pm 1.4\%$	[21.9, 27.5]
p_π	$1.52 \pm 0.06\times$	[1.40, 1.65]	$26.2 \pm 1.1\%$	[24.0, 28.4]
$\cos\theta_\mu$	$1.60 \pm 0.14\times$	[1.41, 1.85]	$26.8 \pm 2.3\%$	[23.7, 31.1]
$\cos\theta_\pi$	$1.63 \pm 0.11\times$	[1.45, 1.89]	$27.1 \pm 1.9\%$	[24.1, 31.6]
$\theta_{\mu\pi}$	$1.58 \pm 0.21\times$	[1.27, 2.02]	$26.7 \pm 3.6\%$	[21.4, 34.2]

Table 45: Fake-data closure: χ^2/ndf of the unfolded result against the A_C -smeared M_A^{RES} -shifted truth, for the two shift signs and the two Wiener–SVD regularisations (number of bins in parentheses). All values are below unity per degree of freedom ($p > 0.63$).

Observable	+1 σ (2nd)	+1 σ (1st)	−1 σ (2nd)	−1 σ (1st)
p_μ (7)	1.26	2.15	0.95	2.14
p_π (5)	1.24	3.44	0.25	0.39
$\cos\theta_\mu$ (5)	0.47	0.63	1.90	1.24
$\cos\theta_\pi$ (5)	2.99	1.07	0.24	0.33
$\theta_{\mu\pi}$ (7)	1.71	3.32	1.96	4.33

Adding the genuine Run 5 ν_μ overlay extends the set to {Run 1, 2, 4, 5} — 2.87×10^6 events (+24%), combined MC POT 8.549×10^{21} , at the full FHC data POT 8.857×10^{20} . The extraction closes for every observable ($\chi^2/\text{ndf} < 1$, $p > 0.92$) and the flux term (22–29% unfolded, amplification $1.29\text{--}1.72\times$) is statistically consistent with the {Run 1, 2, 4} throws above: the extra 24% of overlay does not move the data-limited amplification, as expected, but the measurement now spans the complete FHC exposure.

14.6 Consistency with the sibling ν_e CC 1π measurement

This analysis is one of a pair of MicroBooNE NuMI CC single-charged-pion measurements; the companion ν_e CC 1π analysis shares the same cross-section-extraction framework, beam and reconstruction. A cross-check against its internal note confirms consistency on the essential conventions and identifies the differences that a combined NuMI CC π programme should reconcile.

Consistent by construction. Both use the Wiener–SVD extraction framework; the updated NuMI flux with the fixed geometry bug and Geant v4.10.4; the same signal topology (exactly one $\pi^\pm$, zero π^0 or heavier mesons, any number of protons, charged current); and the same systematic suite — PPFx flux hadronic universes plus the *same 12 beamline-geometry unisims*, the 600-universe GENIE cross-section multisim with unisims, the Geant4 reinteraction reweighting, the standard detector-variation set, and 2% / 1% / 100% POT / target / dirt normalisation

uncertainties. The FHC beam-on data POT per run agrees exactly (Run 2 1.268 , Run 4 2.075×10^{20}).

Differences reconciled or flagged. (i) **POT accounting:** the good-POT is the same beam data for the ν_e and ν_μ analyses, so the exposures are identical; Run 1 FHC = 3.283×10^{20} (standard trigger; the open-trigger era is treated separately), giving a total FHC exposure of 8.857×10^{20} over Runs 1–5. (ii) **Run 4:** the Run4 overlay is identical to Run4c+Run4d combined (same 859,368 events and 2.834×10^{21} POT); only one may be used to avoid double counting. (iii) **Pion threshold:** ν_e uses $\text{KE}_\pi > 40$ MeV ($p_\pi > 0.113$ GeV/ c) versus our $p_\pi > 0.175$ GeV/ c (Table 5). (iv) **Opening-angle cut:** ν_e excludes $\theta_{e\pi} > 170^\circ$ (a shower-specific back-to-back misreconstruction), whereas our muon-track topology uses $\theta_{\mu\pi} < 2.6$ rad (149°). (v) **Regularisation:** the ν_e analysis smooths on the first derivative, this analysis on the second (Section 14.5). (vi) **Scope:** both analyses cover Runs 1–5 in FHC and RHC; the present measurement reports FHC, RHC, and their combination.

15 Summary

We have presented a simultaneous measurement of five $\text{CC}\pi^\pm$ differential cross sections on argon using the full MicroBooNE NuMI exposure in three configurations — FHC (8.857×10^{20} POT), RHC (1.108×10^{21} POT), and their combination (1.994×10^{21} POT) — the first such multi-observable measurement in the MicroBooNE NuMI dataset.

Selection. The event selection achieves a signal purity of $\approx 55\%$ and a flat efficiency of $\sim 13\text{--}20\%$ in the measured phase space (reference phase space: 57.2% purity, $12.1\text{--}12.6\%$ efficiency; Table 10), identical across the three configurations since the selection is the same. The blind data/prediction ratio of ≈ 0.75 is consistent with the known GENIE v3 over-prediction of $\text{CC}\pi^\pm$ production documented in MiniBooNE [21], MINERvA [22], T2K [25], and the MicroBooNE BNB $\text{CC}\pi^+$ measurement [28].

Cross-section results, framework validation, and phase space. Differential cross sections are extracted via a framework-consistent Wiener-SVD unfold using each configuration’s own integrated $\nu_\mu + \bar{\nu}_\mu$ flux (FHC $\Phi = 6.81159 \times 10^{-10}$ cm $^{-2}$ POT $^{-1}$, $\bar{\nu}_\mu$ fraction 34.9%; RHC and combined in Sec. 2). At a common *reference* phase space ($p_\mu, p_\pi > 0.1$ GeV/ c , reco $\theta_{\mu\pi} < 2.6$ rad), the custom pipeline was validated against the fully independent official XSecAnalyzer framework to $\leq 2\%$ on every observable ($d\sigma/dp_\mu$: 1.016, $d\sigma/dp_\pi$: 0.979, $d\sigma/d\cos\theta_\mu$: 0.987, $d\sigma/d\cos\theta_\pi$: 1.004, $d\sigma/d\theta_{\mu\pi}$: 1.001) after three corrections (Section A): importing the full systematic covariance into the Wiener filter (a stat-only covariance over-corrects by $\sim 26\%$), cancelling the EXT subtraction for the pure-MC fake data ($\sim 3.7\%$), and matched binning; the two pipelines were also shown to select identical events (1914 vs 1913.6 NuWro events at the final cut). The final results (Table 23) then refine the *measured phase space* to $p_\mu > 0.15$, $p_\pi > 0.175$ GeV/ c , $\theta_{\mu\pi} < 2.6$ rad, chosen from the efficiency turn-on curves so the selection efficiency is flat at $\sim 13\text{--}20\%$ (Section 4): the momentum thresholds remove the $\sim 5\%$ low-momentum bins and the 2.6 rad cut rejects the cosmic-dominated large-angle region. No lower opening-angle cut is applied, since the small-angle region is high-purity. This raises the efficiency from 12.1% to 15.6% at $\sim 55\%$ purity. The integrated values are correspondingly smaller and, being a different phase space, are not directly comparable to the framework’s full-0.1 GeV/ c numbers.

Context in the experimental landscape. The measured σ_{int} corresponds to $\approx 0.65\text{--}0.70$ times the GENIE v3 prediction, consistent with generator over-predictions of $20\text{--}35\%$ seen by all prior experiments. This measurement provides:

- The first $\text{CC}\pi^\pm$ differential cross sections from the MicroBooNE NuMI (off-axis) dataset, complementing the existing BNB-beam $\text{CC}\pi^+$ result [28] from the same detector.

- Constraints at lower neutrino energies ($E_\nu \sim 0.1\text{--}0.5$ GeV) than the previous NuMI $\text{CC}\pi^+$ results (ArgoNEUT [26] at $\langle E_\nu \rangle \approx 9.6$ GeV, MINERvA [22] at $\langle E_\nu \rangle \approx 3.6$ GeV).
- A measurement on an argon target at the kinematic regime most relevant to the planned DUNE [3] oscillation analysis, which operates at a similar beam energy with a LArTPC far detector.
- Sensitivity to the $\bar{\nu}_\mu$ contribution ($\approx 38\%$) in $\text{CC}\pi^\pm$ production on argon, previously unmeasured.

Required for publication. Native systematic propagation in the custom pipeline (flux, cross-section, and detector-response multisims — currently imported from the framework for the unfolding covariance) and inclusion of Runs 2, 4, and 5 beam-on data ($\approx 5.4\times$ additional exposure) are needed before a real-data result can be published. With the EXT cancellation and full covariance in place the unfolded spectra are everywhere positive and stable; the remaining $\leq 2\%$ residual against the framework is dominated by the covariance-import approximation (Section A). A comparison with NEUT and GiBUU predictions is also planned.

A Data/MC Normalisation

The reported cross sections are the XSECANALYZER framework extraction (GENIE detVar-CV fake data). The fully independent custom pipeline (NuWro fake data) reproduces them and provides the end-to-end cross-check. Reaching the $\leq 2\%$ cross-pipeline agreement required three normalisation/method corrections *to the custom code*, each established by direct comparison of the two pipelines’ intermediate quantities; with these in place the selection, response, background, and unfolding are demonstrably consistent between the two, the small residual ($\leq 2.1\%$) dominated by the covariance-import approximation (below).

A.1 Unfolding covariance ($\sim 26\%$ effect)

Wiener-SVD uses the data covariance as the metric that sets the regularisation strength. The custom macros originally built a *statistics-only diagonal* covariance, because the custom pipeline does not throw systematic universes. With the $\text{CC}\pi^\pm$ response strongly off-diagonal (50–99% bin migration), a stat-only covariance under-states the noise, under-regularises, and inflates the unfolded integral by $\sim 26\%$. This was proven by a swap test: feeding the framework’s own inputs through the same Wiener-SVD unfold with a stat-only covariance reproduced the inflated result, while its full systematic covariance recovered the NuWro truth.

The fix imports the framework’s covariance into the custom unfold. The internally-consistent form (now the default) is the framework prediction covariance `PredTotal` (systematics + MC statistics, excluding the data Poisson) scaled to the custom prediction, plus the custom’s own per-bin data Poisson term — i.e. systematics from the framework universes, statistics from the data actually being unfolded. This removes the 26% bias and brings all five observables to $\leq 2\%$ of the framework.

A.2 EXT (cosmic) subtraction for fake data

Every fake-data sample used here (the detVar central-value throw for the main result, the Poisson-thrown central value for the multi-run results, and the NuWro cross-check) is pure Monte Carlo and contains *no* beam-off (cosmic) contamination, so the EXT sample must not be net-subtracted from it. The framework handles this by *adding* the EXT histogram to the fake-data spectrum, so the EXT term cancels in $\text{data}_{\text{signal}} = (\text{data} + \text{EXT}) - (\text{MC bkg} + \text{EXT} + \text{dirt})$. The custom macros originally subtracted EXT (≈ 41 events at 3.283×10^{20} POT) from a sample that never contained it, a flat $\sim 3.7\%$ deficit on every observable. Replicating the framework’s

EXT cancellation (`fake_data_ext_cancel` in `ccpi_xsec.C`) removes it. For *real* beam-on data this term must be re-enabled, since real data does contain cosmics.

A.3 Opening-angle binning (review item C4)

The custom $\theta_{\mu\pi}$ binning used 10 reco/true bins; the framework uses 9 (the last two custom bins merged into $[2.513, 3.142]$). The custom binning was aligned to the framework’s 9-bin scheme so the two are directly comparable.

A.4 Generator normalisation (NuWro vs GENIE)

In the closure test the NuWro fake data carries roughly twice the selected $\text{CC}\pi^\pm$ signal of the GENIE central-value prediction used to build the response, consistent with the well-documented generator spread in $\text{CC}\pi^\pm$ production on argon [28]. The unfolding correctly recovers the injected NuWro truth (Section B), confirming that the $\sim 2\times$ generator difference is propagated faithfully rather than absorbed into an efficiency or normalisation artefact.

A.5 Software trigger (verified to be a no-op)

A hypothesis that the absence of a `swtrig == 1` filter on MC events was causing a $\sim 13\%$ over-normalisation was investigated. The cut was implemented and the full pipeline rerun; the result was byte-identical. All MC events passing the topological cut already carry `swtrig == 1`: the $\sim 13\%$ of raw MC events failing `swtrig` also fail the topological cut. The cut has been retained for formal correctness but contributes nothing to the current normalisation.

Detector variations. The detector systematic in Table 42 (9.6–14.6%) was built from the Run-1 FHC detector variations applied globally. The multi-run FHC, RHC, and combined results instead use *native* Run-4 detector variations per horn mode (the standard eight-knob set), with the combined result accumulating the Run-4 FHC *and* Run-4 RHC detVar samples — each scaled to its own mode’s data exposure — which the framework now supports via multi-file detVar summing. (The Run-5 FHC Recomb2 and SCE detVar productions turned out to be intrinsic- ν_e overlays with zero ν_μ content; the Run-4 FHC knobs stand in for those two.) Likewise, adding the GENIE spline weight to the central-value weight was a verified no-op: `weightSpline` $\equiv 1$ in the NuMI ntuples and the NuWro overlay’s CV weights are sentinel (-1) values clamped to unity.

B NuWro Fake-Data Closure and Cross-Pipeline Validation

As a generator-model closure and an independent cross-check of the extraction, the XSEC-ANALYZER framework used throughout this note was additionally run (`ProcessNTuples` $\rightarrow$ `UniverseMaker` $\rightarrow$ `Unfolder`) with a NuWro overlay as *fake data* in place of the GENIE detVar-CV sample of the main results. The NuWro reconstructed spectrum is unfolded with the GENIE detector response, efficiency, and GENIE-predicted background subtraction, then compared to the GENIE MicroBooNE-tune prediction. If the unfolding is unbiased, the ratio reflects the genuine NuWro-versus-GENIE model difference rather than a procedural artefact.

Extending the framework beyond p_μ required six additional output branches in the `CC1mu1piXp` selection (reconstructed and true pion momentum, muon $\cos\theta$, and pion $\cos\theta$); the reconstructed pion momentum uses the proton-hypothesis range kinetic energy (`trk_energy_proton_v`), matching the custom pipeline. The true $\theta_{\mu\pi}$ opening angle, previously written as all-zeros, was also corrected, and all fourteen samples were re-processed.

Table 46: Configuration of the NuWro fake-data closure.

Item	Value
Fake data	NuMI FHC NuWro overlay, 6.6504×10^{20} POT native, high statistics (onBNB)
Exposure	all MC/EXT/dirt/detVar scaled to NuWro POT (high-statistics closure)
Response/eff.	GENIE ν_μ MC overlay (central value)
Unfolding	Wiener-SVD, second-derivative regularisation
Systematics	full multisim + detector-variation universes
Observables	all five (p_μ , p_π , $\cos\theta_\mu$, $\cos\theta_\pi$, $\theta_{\mu\pi}$)

The NuWro fake data unfolds to approximately twice the GENIE MicroBooNE tune in every observable, uniformly in shape (Fig. 28, Table 47); the two independent unfolding chains recover the same ratio, cross-validating both. The flat $\sim 2\times$ offset here is a genuine NuWro-versus-GENIE-tune $\text{CC}\pi^\pm$ model difference in this NuMI FHC phase space: NuWro is used at its raw cross section (its `weightTune` and `ppfx_cv` branches are unset and default to unity, so no GENIE tune is applied to the fake data). This must *not* be conflated with the apparent $\sim 2\times$ seen in the primary GENIE-detVar-CV result: the two coincide in magnitude but not in cause — the primary-result excess is a POT/flux *bookkeeping* offset from using the detVar-CV sample as fake data ($1.29\times$ more signal events/POT, times the A_C smearing; Sec. A), whereas the NuWro-closure excess is a real generator difference. A caveat on the latter is that part of the excess may be NuWro’s larger *background* being attributed to signal, since the procedure subtracts GENIE-predicted rather than NuWro backgrounds.




Figure 28: NuWro fake-data closure in the official XSECANALYZER framework. Points: unfolded NuWro fake data; shaded line: GENIE 3.0.6 G18 μB tune. The unfolded NuWro result sits a uniform $\approx 2\times$ above the GENIE prediction in all five observables.

Table 47: NuWro fake-data closure: ratio of integrated cross sections and the χ^2 of the unfolded NuWro result relative to the GENIE tune. The p_μ χ^2 is undefined because its 22-bin Wiener-SVD covariance is numerically singular (condition number $\sim 4 \times 10^{13}$).

Observable	$\sigma_{\text{NuWro}}/\sigma_{\text{Genie}}$	χ^2/bins
$d\sigma/dp_\mu$	1.137	—
$d\sigma/dp_\pi$	1.125	20.3/5
$d\sigma/d\cos\theta_\mu$	1.159	35.0/12
$d\sigma/d\cos\theta_\pi$	0.931	34.0/12
$d\sigma/d\theta_{\mu\pi}$	1.262	25.8/9

Technical notes. Two framework-level issues were resolved. First, the selection’s opening-angle cut ($\theta_{\mu\pi} < 2.6$ rad) leaves the reconstructed bin $[2.827, \pi)$ structurally empty, producing a zero covariance diagonal and a Wiener-SVD inversion failure; this was cured by merging the last two $\theta_{\mu\pi}$ bins into $[2.513, \pi)$, giving nine bins. Second, the χ^2 computation in the diagnostic plotter was made robust to the singular p_μ covariance (it now falls back to a sentinel rather than aborting), so the closure plots render for all observables.

B.1 Extraction-bias investigation and algorithm equivalence

The fake-data closures show a consistent $\sim 1.3\times$ ratio of the unfolded total to the raw generator truth, and the detector-variation central-value fake data sits $\sim 1.29\times$ above the GENIE tune. Because a persistent multiplicative offset in a closure test can indicate an unfolding bug, this was investigated end-to-end; the conclusion is that **there is no framework Wiener-SVD bias**.

Algorithm equivalence. The framework unfolders (`WienerSVDUnfolder.cxx`) was compared line-by-line with the independent `RunWienerSVD.FW` implementation used for the custom pipeline. The two are algorithmically identical: covariance whitening $Q = \text{chol}(V^{-1})$, response pre-scaling $R = QA$, the SVD of RC^{-1} , the per-mode Wiener filter $W_t = \eta_t/(\eta_t + 1)$ with $\eta_t = (D_C V_C^\top C x_{\text{prior}})_t^2$, and the additional smearing matrix $A_C = C^{-1}V_C W_C V_C^\top C$. Feeding the framework’s *exact* extracted matrices (smearacceptance, prior, data, and full covariance) into a standalone replica reproduces the framework’s unfolded result bin-for-bin ($\sum \hat{x} = 5836.5$ for $\cos\theta_\mu$), and the framework smearacceptance column sums equal the raw efficiency.

Self-closure. The decisive test rebuilds the universes with the GENIE ν_μ overlay itself in the `onBNB` slot, so data and response come from the same sample and the detector-variation POT offset is absent by construction. The unfolded total then recovers the truth for all five observables and all three regularisations (Table 48): $\hat{x}/x_{\text{true}} = 0.94\text{--}1.06$. The closure holds even where the A_C diagonal is as low as 0.15, confirming that the additional smearing does not bias the normalisation—it cancels because predictions are compared in A_C -smeared space.

Table 48: GENIE-overlay self-closure: $\hat{x}/x_{\text{true}}$ for the integrated cross section, by observable and regularisation. Data and response are the same sample, removing the fake-data POT offset.

Observable	Wiener-SVD (id.)	Wiener-SVD (2nd-deriv)	D’Agostini (4 it.)
$d\sigma/dp_\mu$	1.137	0.937	1.040
$d\sigma/d\cos\theta_\mu$	1.159	0.973	0.989
$d\sigma/d\cos\theta_\pi$	0.931	0.996	0.992
$d\sigma/dp_\pi$	1.125	1.063	1.059
$d\sigma/d\theta_{\mu\pi}$	1.262	0.996	0.991

Decomposition. The $\sim 1.29\times$ detector-variation offset is a pure POT/flux bookkeeping factor from using the Run-3b detVar central-value sample as fake data (measured flat across all event types, so not a physics or generator difference); the A_C smearing cancels in the closure; and the $\sim 2\times$ NuWro/GENIE ratio (Table 47) is the genuine model difference. None of the three is an unfolding artefact. The closure χ^2 (small, high p -value) is not the most stringent test given the systematics-dominated covariance—the central $\hat{x}/x_{\text{true}}$ ratio is the more discriminating metric and is what establishes the result above. Since the extraction machinery (`ProcessNTuples`→`UniverseMaker`→`Unfolder`) is identical across horn modes, this closure validates the unbiasedness of the FHC, RHC, and combined full-exposure extractions equally.

C Observable Binning

Table 49: Bin edges for all five differential cross-section observables.

Observable	Bin edges
p_μ (GeV/ c)	0.15, 0.20, 0.30, 0.40, 0.50, 0.60, 0.70, 0.80, 0.90, 1.00, 1.10, 1.20, 1.30, 1.40, 1.50, 1.60, 1.70, 1.80, 1.90, 2.00, 2.30, 2.60, 3.00 (first edge 0.15 = signal threshold)
p_π (GeV/ c)	0.175, 0.205, ∞ (first edge 0.175 = signal threshold; second bin open above). <i>Changed 2026-08-17</i> — see §8.8. The previous five-bin scheme (0.175, 0.250, 0.320, 0.420, 0.550, open) had migration diagonals of 99/24/13/6/6%, so only its first bin met the $> 68\%$ usable-bin criterion; the two-bin scheme gives 97/62%. (An earlier draft of this table listed 0.175, 0.20, 0.30, 0.40, 0.50, 1.00, which never matched the configuration files.)
$\cos \theta_\mu$	−1.0, −0.5, 0.0, 0.2, 0.4, 0.55, 0.65, 0.75, 0.82, 0.88, 0.93, 0.97, 1.0
$\cos \theta_\pi$	−1.0, −0.7, −0.4, −0.2, 0.0, 0.2, 0.35, 0.5, 0.65, 0.78, 0.88, 0.95, 1.0
$\theta_{\mu\pi}$ (rad)	0.0, 0.314, 0.628, 0.942, 1.257, 1.571, 1.885, 2.199, 2.513, 2.6 (9 bins; measured range $\theta_{\mu\pi} < 2.6$ rad, no lower cut)

D Selection Cut Parameter Summary

Table 50: All selection threshold constants from `selection/ccpi_selection.C`.

Parameter	Value	Ntuple variable
Topological score	> 0.67	<code>topological_score</code>
Muon track score	> 0.80	<code>trk_score_v</code>
Muon vertex distance	≤ 4.0 cm	
Muon track length	> 10.0 cm	<code>trk_len_v</code>
Muon LLR PID	> 0.20	<code>trk_llr_pid_score_v</code>
Muon MIP BDT	≥ -0.10	TMVA BDT
Pion track length	> 20.0 cm	<code>trk_len_v</code>
Pion vertex distance	< 4.0 cm	
Pion LLR PID	> 0.10	<code>trk_llr_pid_score_v</code>
Pion MIP BDT	> -0.10	TMVA BDT
Pion pion-BDT	> -0.10	TMVA BDT
Pion Bragg score	≥ 0.08	<code>trk_bragg_pion_v</code>
Max. opening angle	< 2.6 rad	
Shower Y-plane hits	< 50 (if 1 shower)	
Shower vertex distance	> 10 cm (if 1 shower)	
Non-proton multiplicity	< 4	LLR PID > 0.1
Primary track multiplicity	< 5	
Cosmic BDT	> 0.0 (disabled)	<code>bdt_cosmic</code>
Software trigger	$= 1$	<code>swtrig</code>

References

- [1] NOvA Collaboration, M. A. Acero *et al.*, “New constraints on oscillation parameters from ν_e appearance and ν_μ disappearance in the NOvA experiment,” *Phys. Rev. D* **98**, 032012 (2018).
- [2] T2K Collaboration, K. Abe *et al.*, “Constraint on the matter–antimatter symmetry-violating phase in neutrino oscillations,” *Nature* **580**, 339–344 (2020).
- [3] DUNE Collaboration, B. Abi *et al.*, “Prospects for Beyond the Standard Model Physics Searches at the Deep Underground Neutrino Experiment,” *Eur. Phys. J. ST* **231**, 3159–3260 (2022).
- [4] J. S. Marshall and M. A. Thomson, “The Pandora Software Development Kit for Pattern Recognition,” *Eur. Phys. J. C* **75**, 439 (2015).
- [5] W. Tang, X. Li, X. Qian, H. Wei, and C. Zhang, “Data Unfolding with Wiener-SVD Method,” *JINST* **12**, P10002 (2017), arXiv:1705.03568 [physics.data-an].
- [6] L. Aliaga *et al.* (MINERvA Collaboration), “Neutrino Flux Predictions for the NuMI Beam,” *Phys. Rev. D* **94**, 092005 (2016).
- [7] C. Andreopoulos *et al.*, “The GENIE Neutrino Monte Carlo Generator,” *Nucl. Instrum. Methods A* **614**, 87–104 (2010).
- [8] MicroBooNE Collaboration, P. Abratenko *et al.*, “Calorimetric reconstruction of interactions with the MicroBooNE LArTPC,” *JINST* **15**, P12037 (2020).

- [9] MicroBooNE Collaboration, P. Abratenko *et al.*, “First Measurement of Energy-Dependent Inclusive Muon-Neutrino Charged-Current Cross Sections on Argon with the MicroBooNE Detector,” *Phys. Rev. Lett.* **128**, 151801 (2022), arXiv:2110.14023 [hep-ex].
- [10] MicroBooNE Collaboration, R. Acciarri *et al.*, “Design and Construction of the MicroBooNE Detector,” *JINST* **12**, P02017 (2017).
- [11] MicroBooNE Collaboration, C. Adams *et al.*, “A Method to Determine the Electric Field of Liquid Argon Time Projection Chambers Using a UV Laser System and Its Application in MicroBooNE,” *JINST* **15**, P07010 (2020).
- [12] D. Rein and L. M. Sehgal, “Neutrino Excitation of Baryon Resonances and Single Pion Production,” *Ann. Phys.* **133**, 79–153 (1981).
- [13] K. S. Kuzmin, V. V. Lyubushkin, and V. A. Naumov, “Lepton polarization in neutrino nucleon interactions,” *Mod. Phys. Lett. A* **19**, 2815–2829 (2004).
- [14] K. M. Graczyk and J. T. Sobczyk, “Form factors in the quark resonance model,” *Phys. Rev. D* **77**, 053001 (2008) [Erratum: *Phys. Rev. D* **79**, 079903 (2009)].
- [15] C. Berger and L. Sehgal, “Lepton mass effects in single pion production by neutrinos,” *Phys. Rev. D* **76**, 113004 (2007).
- [16] J. A. Nowak (MiniBooNE Collaboration), “Four-momentum transfer discrepancy in the charged current π^+ production in the MiniBooNE: Data vs. theory,” *AIP Conf. Proc.* **1189**, 243–248 (2009).
- [17] Y. Hayato, “A neutrino interaction simulation program library NEUT,” *Acta Phys. Polon. B* **40**, 2477–2489 (2009).
- [18] G. M. Radecky *et al.* (ANL Collaboration), “Study of single pion production by weak charged currents in low-energy νd interactions,” *Phys. Rev. D* **25**, 1161–1173 (1982).
- [19] T. Kitagaki *et al.* (BNL Collaboration), “Charged-current exclusive pion production in neutrino–deuterium interactions,” *Phys. Rev. D* **34**, 2554–2565 (1986).
- [20] A. Rodriguez *et al.* (K2K Collaboration), “Measurement of single charged pion production in the charged-current interactions of neutrinos in a 1.3 GeV wide band beam,” *Phys. Rev. D* **78**, 032003 (2008).
- [21] A. A. Aguilar-Arevalo *et al.* (MiniBooNE Collaboration), “Measurement of ν_μ -induced charged-current single positive pion production cross sections on mineral oil at $E_\nu \in 0.5$ –2.0 GeV,” *Phys. Rev. D* **83**, 052007 (2011).
- [22] B. Eberly *et al.* (MINERvA Collaboration), “Charged pion production in ν_μ interactions on hydrocarbon at $\langle E_\nu \rangle = 4.0$ GeV,” *Phys. Rev. D* **92**, 092008 (2015).
- [23] C. L. McGivern *et al.* (MINERvA Collaboration), “Cross sections for ν_μ and $\bar{\nu}_\mu$ induced pion production on hydrocarbon in the few-GeV region using MINERvA,” *Phys. Rev. D* **94**, 052005 (2016).
- [24] MINERvA Collaboration, J. Wolcott *et al.*, “Evidence for Quasielastic Charge Current Antineutrino Scattering Off Protons in Hydrocarbon at $\langle E_{\bar{\nu}} \rangle = 3.5$ GeV,” *Phys. Rev. Lett.* **116**, 081802 (2016).
- [25] K. Abe *et al.* (T2K Collaboration), “Measurement of ν_μ charged-current cross sections on ^{12}C by the T2K ND280 detector,” *Phys. Rev. D* **95**, 012010 (2017).

- [26] R. Acciarri *et al.* (ArgoNEUT Collaboration), “First measurement of the cross section for ν_μ and $\bar{\nu}_\mu$ induced single charged pion production on argon using ArgoNEUT,” *Phys. Rev. D* **98**, 052002 (2018), arXiv:1804.10294 [hep-ex].
- [27] MicroBooNE Collaboration, P. Abratenko *et al.*, “First measurement of ν_e and $\bar{\nu}_e$ charged-current single charged-pion production differential cross sections on argon using the MicroBooNE detector,” *Phys. Rev. Lett.* **135**, 061802 (2025), arXiv:2503.23384 [hep-ex].
- [28] MicroBooNE Collaboration, P. Abratenko *et al.*, “Measurement of single charged pion production in charged-current ν_μ -Ar interactions with the MicroBooNE detector,” *Phys. Rev. D* **113**, 032007 (2026), arXiv:2509.03628 [hep-ex].