

Measurement of $\nu_\mu + \bar{\nu}_\mu$ Charged-Current Single Pion Production Differential Cross Sections on Argon with the MicroBooNE Detector in the NuMI Beam

Analysis note for the full FHC, RHC and combined exposures

Ishanee Pophale and Jarek Nowak

Lancaster University

Abstract

We measure flux-averaged $\nu_\mu + \bar{\nu}_\mu$ charged-current single-charged-pion ($\text{CC}\pi^\pm$) cross sections on argon with the full MicroBooNE NuMI exposure, in three configurations: forward horn current (FHC, 7.766×10^{20} POT), reverse horn current (RHC, 11.082×10^{20} POT) and their combination (18.848×10^{20} POT).

The selection requires a muon candidate and exactly one contained charged-pion candidate attached to a vertex in the fiducial volume, vetoes electromagnetic showers and limits the muon-pion opening angle to 2.6 rad. It reaches 17.5% efficiency at 58–59% purity. Differential cross sections in the muon momentum, the cosine of the muon polar angle, the pion polar angle, the muon-pion opening angle and the pion momentum (three regions) are extracted by Wiener-SVD unfolding. The total flux-averaged cross section is extracted with the same machinery as a single bin. Every result comes with its additional-smearing matrix A_C , through which a prediction must be passed before comparison, and with a covariance that contains the flux, interaction-model, hadron re-interaction, detector, normalisation and statistical terms. The total cross section carries a 38% uncertainty, dominated by the flux (28–30%) and detector (20–22%) terms.

The signal region is blind, and all results are extracted from Poisson fluctuations of the central-value simulation. The four background control regions have been compared with beam-on data: they require no correction of the background model, and the data lie 12–20% above the prediction, within its flux and cross-section uncertainties, in every run period except FHC Run 1, whose 13% offset from the other periods is not yet explained. The extraction is validated by closure, by ensembles of 100 independent pseudo-experiments in six extractions (pull widths 0.92–1.07) and by pseudo-data from reweighted interaction models (largest bias 1.2σ , in one sparse bin). A closure test with an independent generator remains to be done.

Contents

1	Introduction	1
1.1	Measurement overview	1

2	Physics context	2
2.1	Production mechanisms and nuclear effects	2
2.2	Generators	2
2.3	Previous measurements	3
3	Beam, flux and exposure	4
3.1	NuMI beam and flux composition	4
3.2	Beam direction in detector coordinates	5
3.3	Data exposure	6
3.4	Integrated flux and target count	7
4	Simulated and data samples	8
5	Signal definition	10
6	Fiducial volume	11
6.1	Use of the volumes	12
7	Event selection	12
7.1	BDT-based particle identification	12
7.2	Cut definitions	13
7.3	Kinematic reconstruction	14
7.4	Proton-tagged subsample	14
8	Selection performance and backgrounds	15
8.1	N−1 distributions	17
8.2	Proton rejection for the pion candidate	17
8.3	Distributions at the final cut	18
8.4	Background composition	19
8.5	Event displays	19
9	Background control regions	20
9.1	Region definitions and expected yields	20
9.2	Comparison with beam-on data	22
9.3	Validation of the background fit	24
10	Cross-section extraction	26
10.1	Extraction chain	26
10.2	Response matrix	28
10.3	Regularised unfolding	29
10.4	Wiener-SVD implementation and normalisation	30
10.5	Reconstructed spectra, response and efficiency	30
10.6	Reconstruction resolution	34
11	Systematic uncertainties	35
11.1	Breakdown by source	36
11.2	Flux uncertainty and background subtraction	37
11.3	Muon momentum scale	37
11.4	Effects not included in the covariance	38

12 Validation	38
12.1 Closure under the nominal model	39
12.2 Ensemble coverage	39
12.3 Robustness to a different interaction model	39
13 Results	41
13.1 Result set	41
13.2 Differential cross sections	42
13.3 Pion momentum in three regions	45
13.4 Comparing a prediction with the results	46
13.5 Model comparisons	47
13.6 Total cross section	47
13.7 Proton-tagged subsample	49
14 Limitations and requirements for unblinding	49
14.1 Requirements before the signal region is opened	50
15 Summary	50
Glossary	51
A Observable binning	53
B Alternative binnings	54
C Distributions not reported	54
C.1 The differential $W_{\pi p}$ shape	55
C.2 The five-bin $d\sigma/dp_\pi$	55
C.3 The muon polar angle θ_μ	55
D Selection cut parameters	56
E The MicroBooNE detector	56
E.1 Operating Principle	56
E.2 Active Volume and Electric Field	57
E.3 Wire Readout and Signal Formation	57
E.4 Photon Detection System	57
E.5 Calibration and Corrections	57
E.6 Cosmic Ray Environment and Reconstruction	58

1 Introduction

Charged-current single charged-pion ($\text{CC}\pi^\pm$) production,

$$\nu_\mu(\bar{\nu}_\mu) + \text{Ar} \longrightarrow \mu^\mp + \pi^\pm + X, \quad (1)$$

is among the most frequent neutrino interactions at energies of 1–5 GeV and a significant background to oscillation measurements at the same energies, in particular to ν_e appearance at NOvA [3] and T2K [4] and at the future DUNE [5]. The cross section on nuclear targets is poorly constrained: the GENIE v3, NuWro and NEUT generators differ by 20–30% in absolute rate and show shape differences in the individual kinematic variables.

MicroBooNE is a liquid argon time-projection chamber (LArTPC) [33] at Fermilab that recorded data from 2015 to 2021, on axis in the Booster Neutrino Beam and off axis in the NuMI beam. The NuMI exposure contains neutrino- and antineutrino-enhanced samples with a softer and broader spectrum than the Booster beam. We measure $\text{CC}\pi^\pm$ cross sections on argon with the full NuMI exposure, using Pandora pattern recognition [6, 12] and Wiener-SVD unfolding [7].

1.1 Measurement overview

The two horn polarities are analysed as independent measurements, FHC (7.766×10^{20} POT) and RHC (11.082×10^{20} POT), each with its own exposure and integrated flux, and as their combination (18.848×10^{20} POT). The signal is defined by the particles that leave the nucleus: one muon, one charged pion, no neutral pion and any number of nucleons, with the vertex in the fiducial volume and $p_\mu > 0.15 \text{ GeV}/c$, $p_\pi > 0.175 \text{ GeV}/c$ and $\theta_{\mu\pi} < 2.6 \text{ rad}$ (§5). Coherent pion production satisfies this definition and is part of the signal. All polar angles are measured about the neutrino direction, which at the MicroBooNE position is 28.6° from the detector axis (§3.2).

Figure 1 shows the extraction. The same selection is applied to the beam data, to the beam-off data that measure the cosmic-ray background, and to the simulation. The simulated neutrino background, the beam-off data and the simulated interactions outside the cryostat are subtracted, and the remaining spectrum is unfolded with the simulated response. The result of each extraction is a cross section, the additional-smearing matrix A_C that describes the bias of the regularised unfolding, and the covariance matrix of the uncertainties.




Figure 1: The extraction chain. The same selection is applied to data, beam-off data and simulation; the simulation also supplies the response and the systematic universes. The output is a cross section together with the additional-smearing matrix A_C and the covariance, all three of which are needed to compare the result with a prediction.

Five differential distributions and the total cross section are measured in each configuration (Table 19): $d\sigma/dp_\mu$, $d\sigma/d\cos\theta_\mu$, $d\sigma/d\cos\theta_\pi$, $d\sigma/d\theta_{\mu\pi}$, and the cross section in three regions of pion momentum. The coarse three-region form of p_π follows from the saturation of the range-based pion momentum near $0.34 \text{ GeV}/c$ (§13.3). The differential results estimate A_C times the true spectrum, so a prediction is compared after multiplication by A_C (§13.4); A_C does not

preserve the normalisation, and the total cross section is taken from the one-bin extraction, for which $A_C = 1$ (§13.6).

The signal region is blind. All results are extracted from fake data, Poisson fluctuations of the central-value simulation at the recorded exposure, and every figure of an unfolded result carries the label “BLINDED: CV PSEUDO-DATA”. The four background control regions are disjoint from the signal region and have been compared with beam-on data (§9.2).

The inclusive analysis, that is the signal definition, event selection, background treatment, unfolding, uncertainty model and the result set of Table 19, is presented for approval. Opening the signal region further requires the items listed in §14.1, among them the confirmation of the absolute flux normalisation and an understanding of the FHC Run-1 offset seen in the control regions.

The proton-tagged subsample, which adds hadronic-mass and transverse-kinematic-imbalance observables, is analysed in a companion note [53]. The optimisation studies and cross-checks behind the choices described here are collected in the technical supplement, and the comparison of the control regions with beam-on data in the control-sample note. The additional-smearing and covariance matrices of every extraction are provided as text files in `report/data_release/`.

2 Physics context

2.1 Production mechanisms and nuclear effects

Between 0.5 and 5 GeV, charged-current single charged-pion production proceeds mainly through baryon resonances, dominated by the $\Delta(1232)$ and modelled with Rein–Sehgal-type form factors [35, 36, 37, 38, 39]. A non-resonant continuum and deep-inelastic scattering together contribute about 20% of the rate in the present phase space and are added incoherently by every generator used here. Coherent production off the whole nucleus contributes at most 10%; it satisfies the topological signal definition of §5 and is part of the signal. In argon, Fermi motion, Pauli blocking and binding modify the kinematics, and final-state interactions (pion absorption at 25–30%, charge exchange and quasi-elastic scattering) move events between the $CC\pi^\pm$, $CC\pi^0$ and $CC0\pi$ topologies at the 30% level. For this reason the signal is defined by the final-state particle content rather than by the production process. The same mixing makes $CC0\pi$ the largest background: pion absorption, and a proton reconstructed as the pion, bring such events into the selection (Sec. 12 of the technical supplement).

2.2 Generators

Four generators are compared with the results; the model choices relevant to this measurement are listed in Table 1. The response and the fake data are built from the MicroBooNE tune of GENIE v3 [16]. In the measured phase space the standalone generators span a factor of 1.6 in the total cross section, from GENIE (lowest) to NEUT (highest) (Table 19 of the technical supplement), and bracket the tune. Since the present results are extracted from fake data, the model comparisons test the procedure; the spread indicates the size of the model differences the measurement is expected to resolve.

Table 1: Model content of the generators compared, restricted to what matters for interpreting this measurement.

Generator	Resonance model	Non-resonant / DIS	Final-state interactions
GENIE v3 [9, 20]	KLN-BS, 18 resonances to $W \lesssim 2$ GeV	Bodek–Yang, incoherent; AGKY transition	INTRANUKE hA (one effective interaction)
NuWro [21]	$\Delta(1232)$ only (Paschos–Yu)	separate amplitude; smooth transition to Bodek–Yang	Salcedo–Oset (in-medium Δ , two-body absorption)
NEUT [40, 23]	modified Rein–Sehgal, 18 resonances, Graczyk–Sobczyk Δ form factors	GRV98 with Bodek–Yang above $W \approx 2$ GeV	step-by-step intranuclear cascade
GiBUU [22]	Δ and higher resonances with in-medium spectral functions	PYTHIA fragmentation above $W \approx 1.6$ GeV	transport equations with pion optical potential

2.3 Previous measurements

Table 2 lists published $\text{CC}\pi^+$ measurements on nuclear targets. The two most relevant are from MicroBooNE. The Booster-beam measurement [52], on the same detector with an on-axis spectrum peaked at 0.7 GeV and a negligible $\bar{\nu}_\mu$ component, gives $\sigma = (3.75 \pm 0.07_{\text{stat}} \pm 0.80_{\text{syst}}) \times 10^{-38} \text{ cm}^2/\text{Ar}$ and finds all tested generators above the data at very forward muon angles. The NuMI $\nu_e + \bar{\nu}_e$ $\text{CC}\pi^+$ measurement [51] uses the same beam, reconstruction and extraction framework and conventions as this analysis (Sec. 11.1 of the technical supplement). The NuMI flux at MicroBooNE is softer and broader than the Booster flux and has a 35% $\bar{\nu}_\mu$ component in FHC, and the flux-averaged cross section divides by the whole flux above 60 MeV while pion production starts near 0.3 GeV; the value of about $1.0 \times 10^{-38} \text{ cm}^2/\text{Ar}$ found here is therefore not directly comparable with the Booster result (§3.4). Among other experiments, the T2K off-axis measurement on carbon [49] has the most similar spectrum. A longer survey is given in App. 16 of the technical supplement.

Table 2: Selected $\text{CC}\pi^+$ cross-section measurements on nuclear targets. $\langle E_\nu \rangle$ is the approximate mean or peak neutrino energy. “Type” indicates whether a total integrated (T) or differential (D) cross section was reported. “Target” gives the nucleus primarily used for cross-section extraction.

Experiment	Target	$\langle E_\nu \rangle$ (GeV)	Type	Key result / reference
ANL 12-ft [41]	D	1.0	T	Bubble chamber; foundational σ vs E_ν
BNL 7-ft [42]	D	1.5	T	Higher statistics; 15% normalisation tension with ANL
K2K [43]	C	1.3	T	First accelerator $\text{CC}\pi^+$ on carbon near detector
MiniBooNE [44]	CH_2	0.7	D	Q^2 , W , p_μ , θ_μ ; 20% excess over NUANCE
MINERvA [46]	CH	3.6	D	p_π , θ_π , Adler angles; NuMI medium-energy
T2K ND280 [49]	CH	0.85	D	p_μ , $\cos\theta_\mu$; off-axis beam similar to this work
ArgoNEUT [50]	Ar	9.6	T	First $\text{CC}\pi^+$ on Ar; 115 ν + 158 $\bar{\nu}$ events (1.25×10^{20} POT)
MicroBooNE BNB [52]	Ar	0.7	D	Full BNB dataset (1.11×10^{21} POT); $\sigma =$ (3.75 ± 0.80) $\times 10^{-38}$ cm^2/Ar
MicroBooNE NuMI ν_e [51]	Ar	0.73	D	First $\nu_e + \bar{\nu}_e$ $\text{CC}\pi^+$ on Ar; $\sigma = (0.93 \pm 0.28) \times$ 10^{-39} $\text{cm}^2/\text{nucleon}$

3 Beam, flux and exposure

3.1 NuMI beam and flux composition

The NuMI beam is produced by 120 GeV protons from the Fermilab Main Injector on a graphite target. Two magnetic horns focus charged secondary hadrons of one sign. In forward horn current (FHC) the horns focus positive pions and kaons and the beam is predominantly ν_μ ; in reverse horn current (RHC) they focus negative mesons and the beam is predominantly $\bar{\nu}_\mu$. MicroBooNE is located 134.6 mrad (7.71°) off the NuMI beam axis¹, which places the flux peak between 0.5 and 2.5 GeV and narrows the spectrum. At this off-axis position and energy the sign selection of the horns is incomplete, and both modes contain a large wrong-sign component, largest in RHC.

The flux composition at the detector is taken from the new-G4 central-value flux histograms, integrated above $E_\nu = 60$ MeV. The threshold removes an unphysical accumulation near $E_\nu = 0$, below the reaction threshold, which the flux group removes in the same way. The muon-flavour fractions of these histograms reproduce the authoritative integrated fluxes of §3.4 exactly.

¹From the detector position in beam coordinates, (5502, 7259, 67270) cm: 91.1 m transverse at 672.7 m downstream (§3.2).

Table 3: NuMI flux composition at the MicroBooNE site by neutrino type, for the FHC and RHC horn modes (new-G4 central-value flux, integrated over $E_\nu > 60$ MeV, the same threshold as the authoritative integrals of Table 5; above 0.1 GeV the RHC split would read 42.0/55.2%). FHC is ν_μ -dominant; RHC is $\bar{\nu}_\mu$ -dominant with a large wrong-sign ν_μ fraction. The combined row weights each mode by its fluence, $\Phi \times \text{POT}$ (FHC 5.29, RHC 7.14×10^{11} cm $^{-2}$; Tables 4 and 5), giving a near-balanced $\nu_\mu/\bar{\nu}_\mu$ mix. The signal is charge-blind ($|\nu_{\text{pdg}}| = 14$), so both ν_μ and $\bar{\nu}_\mu$ contribute in every mode.

Horn mode	ν_μ	$\bar{\nu}_\mu$	ν_e	$\bar{\nu}_e$
FHC	63.4%	34.0%	1.7%	0.9%
RHC	44.2%	53.3%	1.4%	1.2%
Combined	52.4%	45.1%	1.5%	1.1%

The Package to Predict the FluX (PPFX) [8] corrections are applied event by event through the central-value weight stored in the analysis ntuples. The absolute normalisation of §3.4 uses the same new-G4 histograms.

3.2 Beam direction in detector coordinates

Two different angles describe the geometry of MicroBooNE in the NuMI beam (Fig. 2). The off-axis angle is a property of the beamline: the detector lies 91.1 m from the NuMI axis at 672.7 m downstream of the target, a baseline of 678.8 m and an off-axis angle of 7.71° (134.6 mrad).² This angle does not describe the orientation of the TPC.

The angle that enters the measurement is the direction of the neutrinos relative to the detector z axis. Averaged over interactions in the active volume, the true neutrino direction in detector coordinates is (0.4624, 0.0489, 0.8853), 28.6° from z . The detector z axis is therefore neither the direction from which the neutrinos arrive nor the NuMI beamline axis, which is itself 23.2° from z .




Figure 2: The NuMI geometry seen by MicroBooNE. (a) Beamline geometry, to scale: the detector sits 91.1 m off the NuMI axis at a baseline of 678.8 m, an off-axis angle of 7.7° . (b) The same neutrinos in TPC coordinates, to scale (the x - z projection; the quoted 28.6° is the full three-dimensional angle): the beam arrives 28.6° away from the detector z axis, so one and the same muon is assigned $\theta_\mu = 12^\circ$ about z but $\theta_\mu = 40.6^\circ$ about the neutrino. The generators mean the second. (c) The resulting shift in the true $\cos \theta_\mu$ distribution; dotted lines mark the analysis bin edges. Under the two conventions 64.6% of signal events fall in a different bin (44.6% for $\cos \theta_\pi$).

²Computed from the detector position in beam coordinates, (5502, 7259, 67270) cm, with the rotation matrix used for the NuMI beamline geometry weights.

Every polar angle and every transverse-kinematic-imbalance variable is defined with respect to the neutrino direction. In the generator predictions the neutrino defines the $+z$ axis (the GENIE event records used here have $p_x^\nu = p_y^\nu = 0$ and $p_z^\nu = E_\nu$ exactly), so angles measured about the detector z axis would not correspond to the predictions they are compared with. The difference is large: 64.6% of signal events change $\cos\theta_\mu$ bin between the two conventions, and 44.6% change $\cos\theta_\pi$ bin. For the transverse-kinematic-imbalance variables the error exceeds the size of the observable, with a mean δp_T of 0.742 GeV/ c about z against 0.264 GeV/ c about the beam.

All angular and transverse observables are therefore measured about the neutrino direction, following the β convention of the NuMI ν_e CC1 π analysis [1] and of the earlier NuMI ν_e measurement it builds on [2], which keeps the two NuMI measurements directly comparable. At truth level the exact neutrino direction of each event is used, which reproduces the generator convention. At reconstruction level the neutrino direction is approximated by the vector from the NuMI target, at $(-31388, -3316, -60100)$ cm in detector coordinates (684 m from the detector centre and 678.8 m from the TPC origin), to the reconstructed interaction vertex.

The result is insensitive to this choice of reconstruction-level axis. A single fixed axis and the per-event β differ by a median of 2.4 mrad (at most 7.3 mrad), which moves 0.42% of muons across a $\cos\theta_\mu$ bin edge, compared with 64.6% for the detector- z convention. Against the true direction the two perform almost identically, with mean opening angles of 1.95° (β) and 2.00° (fixed axis). Neither recovers the divergence of the beam: the neutrinos come from decays along the whole decay pipe rather than from the target, and over a 684 m baseline a 2.5 m detector subtends about 4 mrad, so the vertex position carries little information on the direction. This residual is a resolution effect and is absorbed by the response matrix.

The muon and pion momenta, the muon–pion opening angle $\theta_{\mu\pi}$, the invariant mass $W_{\pi p}$ and all integrated cross sections do not depend on the choice of axis. Neither does the event selection, since no selection cut and no BDT input uses a polar angle.

3.3 Data exposure

The FHC measurement uses Runs 1, 2, 4 and 5 and the RHC measurement Runs 1, 2, 3, 4a, 4b and 4c; the combined measurement uses all fourteen overlay files. The exposure of each run period is listed in Table 4. For FHC Run 1 it is that of the data recorded without the open trigger, 2.192×10^{20} POT and 5 748 692 triggers (runs 4952–6748); the earlier open-trigger running (runs 3490–4950) is not used. Every run is normalised to its own exposure, so the prediction contains the run periods in the proportion of the recorded data.

The analysis is blind and the beam-on data of the signal region are not used. Each configuration is extracted from fake data thrown run by run at that run’s exposure, so every fake-data sample has the same run structure as the beam-on sample it stands in for. At unblinding the beam-on files of the same runs, with the same POT and trigger counts, replace the fake data.

Table 4: Exposure by run period. “MC POT” is the simulated exposure of each period and “Data POT” the recorded exposure to which the fake data of that period are thrown. Each run is normalised independently by $\text{Scale} = (\text{Data POT})/(\text{MC POT})$, so the prediction is the data-weighted mixture of run periods, the normalisation a per-run beam-on dataset produces. The per-run factors range from 0.045 to 0.116; a single factor per horn mode would weight the runs by their simulated instead of their recorded exposure. The RHC Run-4 row sums the 4a, 4b and 4c files and the Run-3 row the aa–ae files; files of one run are scaled by the total simulated exposure of that run. Run 5 uses the ν_μ overlay produced with reweighted PPFx flux weights, the only Run-5 overlay available.

Run period	MC POT ($\times 10^{20}$)	Data POT ($\times 10^{20}$)	Scale
Run 1 (FHC)	23.282	2.192	0.0942
Run 2 (FHC)	24.934	1.268	0.0509
Run 4 (FHC)	28.335	2.075	0.0732
Run 5 (FHC)	19.300	2.231	0.1156
FHC total	95.85	7.766	<i>per run</i>
Run 1 (RHC)	8.997	0.605	0.0673
Run 2 (RHC)	57.865	2.591	0.0448
Run 3 (RHC, aa–ae)	55.185	5.003	0.0907
Run 4 (RHC, 4a–4c)	32.586	2.883	0.0885
RHC total	154.63	11.082	<i>per run</i>
Combined (FHC+RHC)	250.49	18.848	<i>per run</i>

With the full exposure the data-statistical uncertainty is smaller than the flux and detector terms in every observable (Table 20 of the technical supplement–Table 22 of the technical supplement).

3.4 Integrated flux and target count

The absolute normalisation uses the new-G4 NuMI flux integrated above $E_\nu = 60$ MeV. The signal does not distinguish the charge of the lepton ($|\nu_{\text{pdg}}| = 14$ covers both $\nu_\mu \rightarrow \mu^- \pi^+$ and $\bar{\nu}_\mu \rightarrow \mu^+ \pi^-$), so the normalisation is the summed flux $\Phi = \Phi_{\nu_\mu} + \Phi_{\bar{\nu}_\mu}$. The RHC flux is $\bar{\nu}_\mu$ -dominated and differs in size from the FHC flux, so each configuration is normalised with its own integral (Table 5).

Table 5: Integrated ν_μ and $\bar{\nu}_\mu$ fluxes per POT used by the analysis (new-G4 flux, $E_\nu > 60$ MeV), and the combined $\nu_\mu + \bar{\nu}_\mu$ value used as the per-configuration normalisation. The combined row is the data-POT-weighted average of the two modes (weights 7.766 : 11.082; Table 4). Each configuration is normalised with its own flux: the RHC flux is 5.4% below the FHC one, so using a single value for both would bias the RHC cross section by that amount.

Config	Φ_{ν_μ} ($10^{-10} \text{ cm}^{-2} \text{ POT}^{-1}$)	$\Phi_{\bar{\nu}_\mu}$	$\Phi_{\nu_\mu + \bar{\nu}_\mu}$	$\bar{\nu}_\mu$ frac.
FHC	4.43515	2.37644	6.81159	34.9%
RHC	2.92348	3.52298	6.44646	54.6%
Combined	3.54634	3.05057	6.59691	46.2%

The mean energy of the $\nu_\mu + \bar{\nu}_\mu$ flux above 60 MeV is 0.479 GeV (FHC) and 0.475 GeV (RHC). For ν_μ alone it is 0.468 and 0.544 GeV, for $\bar{\nu}_\mu$ alone 0.500 and 0.417 GeV; in RHC the antineutrinos are the focused component and are softer at this off-axis angle. Only 40% of the

flux lies above 0.3 GeV, 26% above 0.5 GeV and 13% above 1 GeV. Since the flux-averaged cross section divides by the whole flux above 60 MeV while single-pion production starts near 0.3 GeV, the total is about a factor 2.5 smaller than it would be for the Booster spectrum peaked at 0.7 GeV. This, together with the measured phase space ($p_\pi > 0.175$ GeV/ c , $\theta_{\mu\pi} < 2.6$ rad), is why the value of about 1.0×10^{-38} cm²/Ar found here differs from the 3.75×10^{-38} cm²/Ar of the Booster measurement (§2.3).

Every simulated event carries the PPFX central-value flux weight (§4; mean 0.944 on selected signal), so the integrated flux that divides the event rate must be the PPFX-corrected central value. The central-value histograms of the new-G4 flux file are taken to be that central value. An integral of the uncorrected GEANT4 flux would lower every absolute cross section by 5–6%. The integrated values were validated in four independent ways: standalone NuWro and GENIE event generation on the same flux, the numbers of the flux author, and a GENIE closure. A written confirmation from the NuMI flux contact that the constants of Table 5 are PPFX-corrected central values, consistent with the per-event weighting, is required before the signal region is opened.

The number of argon nuclei in the fiducial volume is

$$N_{\text{Ar}} = \frac{\rho_{\text{LAr}} \cdot V_{\text{FV}} \cdot N_A}{A_{\text{Ar}}} = 8.710 \times 10^{29}, \quad (2)$$

with $\rho_{\text{LAr}} = 1.3836$ g/cm³, $A_{\text{Ar}} = 39.948$ g/mol and $N_A = 6.02214076 \times 10^{23}$ mol⁻¹. The fiducial box, $10 < x < 246$, $-101 < y < 101$, $10 < z < 986$ cm (4.6528×10^7 cm³), minus the unresponsive region $675.1 < z < 775.1$ cm gives $V_{\text{FV}} = 4.1761 \times 10^7$ cm³; without the subtraction the target count would be 9.705×10^{29} . The extraction computes N_{Ar} from the same box and subtraction, and the value enters the conversion factor of §10 and the total of §13.6. The normalisation of each configuration is $N_{\text{norm}} = \Phi_{\text{mode}} \cdot N_{\text{POT,mode}} \cdot N_{\text{Ar}}$, with the flux and exposure of that mode (Tables 4 and 5).

4 Simulated and data samples

Neutrino interactions are simulated with the MicroBooNE tune [16] of GENIE v3 [9, 20] on argon-40, propagated through the detector with Geant4 [24], and overlaid on cosmic-ray data recorded without beam, so that simulated events contain the detector noise and cosmic activity of real events. Table 6 lists the samples.

Table 6: Input samples (final, full FHC+RHC). “ n ” is the number of overlay files; “MC POT” is the summed native simulated exposure ($\times 10^{20}$). The blind Poisson-thrown fake data stands in for the (unused) beam-on data; the detector variations use the native Run-4 FHC and RHC samples. The per-file lists and POT and trigger values are in the configuration files of the analysis.

Sample	Composition	n	MC POT
ν_μ overlay (FHC)	Runs 1, 2, 4 (pandora) + Run 5 (reweighted-PPFX)	4	95.85
ν_μ overlay (RHC)	Runs 1, 2, 3, 4a, 4b, 4c (pandora)	10	154.63
Dirt MC (interactions outside the cryostat)	GENIE NuMI dirt overlays, per run (Run-1 file standing in for Run 2)	7	n/a
EXT (beam-off)	cosmic data, per run, horn-mode-independent	7	n/a
Detector variations	Run-4 FHC + Run-4 RHC, CV + 8 knobs each	18	n/a
Fake data (blind)	Poisson-thrown central value, per configuration	n/a	n/a
Beam-on data	blinded, not used until unblinding	n/a	n/a

Event weights. Each simulated event is weighted by

$$w = \text{Scale}[\text{run}] \times w_{\text{PPFX}} \times w_{\text{tune}}, \quad (3)$$

where $\text{Scale}[\text{run}] = (\text{run Data POT})/(\text{run MC POT})$ is the per-run factor of Table 4, w_{PPFX} the PPFX central-value flux correction and w_{tune} the GENIE tune weight. In the framework each overlay file receives the exposure of its run from the matching beam-on (here fake-data) entry, and the simulated exposure of a run is the sum of the distinct exposure values stored in its files, which is the run total under both file conventions in use. One rule is applied to every event weight,

$$w_{\text{safe}} = \begin{cases} 1, & w \text{ non-finite, } w < 0, \text{ or } w > 30, \\ w, & \text{otherwise (a weight of exactly zero is kept),} \end{cases} \quad (4)$$

to the product $w = \text{Scale} \times w_{\text{PPFX}} \times w_{\text{tune}}$, including the 0.65 normalisation weight of the dirt overlay that scales its generated exposure to the rate of interactions outside the cryostat expected at the NuMI site. The universes, the covariance, the unfolding inputs, the cut-flow histograms, the count-level tables, the control-region yields and transfer factors, and the beam-on comparison at unblinding all use Eq. (4), so that count-level and covariance-level numbers agree at the per-cent level.

Beam-off normalisation. Beam-off data contain cosmic rays only, so a beam-off sample is matched to a beam-on run by run period, independently of the horn polarity in which it was recorded. Each run period uses the beam-off samples of that period, scaled by the ratio of beam-on to beam-off gates times the NuMI beam-occupancy factor 0.98. Several beam-off files of one run are pooled:

$$\text{EXT}_{\text{run}} = 0.98 \frac{\text{bnb_gates}[\text{run}]}{\sum_{\text{files} \in \text{run}} \text{ext_gates}} \sum_{\text{files} \in \text{run}} N_{\text{EXT}}^{\text{sel}}. \quad (5)$$

No beam-off sample exists for Run 2; the Run-1 sample, which like Run 2 predates the cosmic-ray tagger, is used in its place and scaled by its own gate count. The same beam-off files serve the FHC and RHC runs of a period, and the combined configuration lists each (file, run) pair once. The selected beam-off rate per 10^6 gates is 1.75 (Run 1), 0.43 (Run 3b), 1.15 (Run 4) and 0.52 (Run 5), from 8, 14, 40 and 10 selected beam-off events, so the rates carry Poisson uncertainties of 16–35%. The Run-1 rate is 1.5 to 4 times the rates of the later periods, which is why the samples are matched by period instead of pooled. At the final cut the beam-off background is 23 events in FHC, 23 in RHC and 46 in the combination, 1.6%, 1.3% and 1.5% of the selected sample. Figure 3 shows the rates, scale factors and yields per period; the numbers are in `data_release/ext_gates.tsv`, and the normalisation manifest (Sec. 15 of the technical supplement) checks that every (beam-off file, run) pair enters each configuration once and that the gate counts are consistent.




Figure 3: Beam-off normalisation by run period (`data_release/ext_gates.tsv`): selected beam-off rate per 10^6 gates at the final inclusive cut, with its Poisson uncertainty (left), the scale factor $0.98 \times (\text{beam-on gates}) / (\text{beam-off gates})$ (centre), and the resulting beam-off events at the final cut (right), for each beam-on run period of the FHC and RHC configurations. Run 2 of both modes uses the Run-1 sample. The Run-1 rate is 1.5 to 4 times the rates of the later periods.

Cosmic-ray rejection. No cosmic-ray rejection beyond the selection of §7 is applied. The beam-off contribution, 1.6% of the selected FHC prediction, is removed by a data-driven subtraction, so an additional cut could only reduce the statistical uncertainty of that subtraction, about five events in a sample of 1418, and every variant studied removes more signal than that. The directional and calorimetric variants are described in Sec. 14.2 of the technical supplement.

Software trigger. Beam-on and beam-off files are filtered on the software trigger. The beam-off files of Runs 4 and 5 carry a trigger flag that is zero throughout, and for them the `swtrig_pre` flag is used instead. Simulated events that fail the software trigger are removed at the NeutrinoSlice cut; every simulated event that passes the topological cut also passes the trigger, so the requirement has no effect on the selected sample.

5 Signal definition

The signal is defined at generator level and is used to build the efficiency and the response; it is never applied as a reconstruction cut. It does not distinguish the charge of the lepton: $\nu_\mu \rightarrow \mu^- \pi^+$ and $\bar{\nu}_\mu \rightarrow \mu^+ \pi^-$ are both signal.

Neutrino flavour

$$|\nu \text{ PDG}| = 14 \text{ } (\nu_\mu \text{ or } \bar{\nu}_\mu).$$

Interaction type

Charged current.

True vertex

Inside the fiducial volume (§6).

Primary multiplicity

Among the primary final-state particles, exactly one muon ($|\text{PDG}| = 13$), exactly one charged pion ($|\text{PDG}| = 211$) of any momentum, no neutral pion ($|\text{PDG}| = 111$), no kaon ($K^\pm, K^0, K_S, K_L$) and no other meson. A second charged pion below the momentum threshold makes the event background. Pions from the decay of primary hyperons are not primary particles and are not counted. Any number of protons and neutrons is allowed.

Kinematic thresholds

$$p_\mu > 0.15 \text{ GeV}/c, p_\pi > 0.175 \text{ GeV}/c \text{ and } \theta_{\mu\pi} < 2.6 \text{ rad.}$$

The thresholds define the measured phase space. The momentum efficiency rises sharply between 0.15 and 0.20 GeV/c, set by the 10 cm muon and 20 cm pion track-length requirements; below it the efficiency is 1–3% and the correction strongly model dependent. The upper bound on the opening angle removes the region above 2.6 rad, which is dominated by cosmic-ray and dirt background (about 7% purity in the $\theta_{\mu\pi} > 2.513$ rad bin). No lower bound is applied: the small-angle region has a purity of about 58%, so a cut there would remove signal without improving the purity. The same opening-angle requirement is applied at reconstruction level so that the signal and selection acceptances match. The first p_μ and p_π bins start, and the last $\theta_{\mu\pi}$ bin ends, at these thresholds.

Pion threshold. The NuMI ν_e CC1 π analysis defines its pion signal by kinetic energy, $\text{KE}_\pi > 40$ MeV, equivalent to $p_\pi > 0.113$ GeV/c. The threshold used here, $p_\pi > 0.175$ GeV/c ($\text{KE}_\pi > 84$ MeV), is higher because the pion momentum, reconstructed from the track range under the muon hypothesis and converted to the pion mass, degrades rapidly below it (Table 7). The pion-momentum bias is the signed quantity $b(p_\pi) = \langle (p_\pi^{\text{reco}} - p_\pi^{\text{true}})/p_\pi^{\text{true}} \rangle$, which has two regimes. Below the threshold the few pions that are reconstructed at all are biased upwards (+30 to +131%): they are the tracks long enough to pass the selection. At the threshold the mass conversion removes most of the hypothesis bias (−8%, with 31% resolution). Above about 0.4 GeV/c the bias becomes increasingly negative, because pions interact before they stop and the visible range is truncated; the mean reconstructed momentum saturates near 0.34 GeV/c, and the bias reaches −65% above 0.8 GeV/c. Between 0.113 and 0.175 GeV/c the resolution is 51–56% and the selection efficiency 2.2%. A lower threshold would therefore add a bin with large, strongly model-dependent unfolding corrections, even as a single coarse bin, and it is not adopted; no cross section is reported below $p_\pi = 0.175$ GeV/c.

Table 7: Pion momentum reconstruction versus true p_π for correctly reconstructed signal pions. Below the 0.175 GeV/c threshold the bias and resolution grow and the reconstructed fraction halves.

True p_π [GeV/c]	Bias $\langle (p_\pi^{\text{reco}} - p_\pi^{\text{true}})/p_\pi^{\text{true}} \rangle$	RMS resolution	Reco fraction
0.080–0.113	+131%	73%	3.4%
0.113–0.150	+66%	56%	3.4%
0.150–0.175	+30%	51%	4.3%
0.175–0.250 (threshold)	−8%	31%	10.0%
0.250–0.400	−30%	21%	10.4%

6 Fiducial volume

The fiducial volume (FV) is

$$x \in [10.0, 246.0] \text{ cm}, \quad y \in [-101.0, +101.0] \text{ cm}, \quad z \in [10.0, 986.0] \text{ cm}, \quad (6)$$

with the region of unresponsive wires, $z \in [675.1, 775.1]$ cm, excluded. Its volume is

$$\begin{aligned} V_{\text{FV}} &= (246 - 10) \times 202 \times (976 - 100) \text{ cm}^3 \\ &= 236 \times 202 \times 876 \text{ cm}^3 \approx 4.17 \times 10^7 \text{ cm}^3. \end{aligned} \quad (7)$$

A larger containment volume, inset by 2 cm from the faces of the active volume and by 10 cm from the top face, is used to require that the pion-candidate track ends inside the detector:

$$x \in [2.0, 254.35] \text{ cm}, \quad y \in [-113.53, +107.47] \text{ cm}, \quad z \in [2.1, 1034.9] \text{ cm}. \quad (8)$$

6.1 Use of the volumes

volume	extent	used for
fiducial (FV)	$[10, 246] \times [-101, 101] \times [10, 986]$ cm minus $z \in [675.1, 775.1]$ cm	vertex (true and reco)
containment	$[2, 254.35] \times [-113.53, 107.47] \times [2.1, 1034.9]$ cm	track end points
active TPC	$[0, 256.35] \times [-115.5, 117.5] \times [0, 1036.9]$ cm	reference only

The true and reconstructed fiducial volumes are the same box. The signal definition tests the true neutrino vertex and the selection the reconstructed vertex, through the same function, so both apply the exclusion of the unresponsive region. Vertex resolution and space-charge distortions can move a reconstructed vertex across a boundary; such events are part of the simulated response, true-in and reconstructed-out as inefficiency and true-out and reconstructed-in as the out-of-fiducial-volume background of Table 10.

The exclusion of the unresponsive region applies to the vertex only. An event with its vertex outside the region is kept even if one of its tracks crosses it; such tracks have degraded calorimetry and range, and the resulting inefficiency is part of the response. The region is also removed from the target count (Eq. (2)); keeping it would raise the target number by $100/876 \approx 11\%$ and lower every cross section by the same factor.

The fiducial volume is inset from the active TPC by about 10 cm in x , 15 cm in y , 10 cm at the upstream z face and 51 cm at the downstream face. The containment volume is inset by only 2 cm (10 cm at the top), since the range-based momentum and the pion-containment cut need the track to end inside the detector, not inside the fiducial region.

7 Event selection

The selection is implemented as a selection class of the extraction framework, which produces the response matrices and spectra of every result; an independent standalone implementation of the same cut sequence is kept as a cross-check. It consists of ten sequential cuts, with Cut 0 the unselected reference sample.

7.1 BDT-based particle identification

Three gradient-boosted decision trees (TMVA [26]), trained on the MicroBooNE overlay simulation, provide the calorimetric particle identification of the muon and pion candidates. All three use the same per-track inputs: the Bragg-peak dE/dx likelihoods under the proton, muon and minimum-ionising hypotheses, the log-likelihood-ratio (LLR) PID score, the Pandora track-shower score, and the space-charge-corrected track end points.

- **MIP BDT:** separates minimum-ionising tracks (muons, pions) from stopping protons. Both candidates must have a MIP score above -0.1 (Cuts 4 and 5).
- **Pion BDT:** a pion-versus-proton discriminant for the pion candidate, with a required score above -0.1 (Cut 5).
- **Muon BDT:** the same inputs plus the track length. It applies no threshold. The tracks that qualify as muon candidates are defined by the cuts of Table 8, and the qualifying track with the highest muon-BDT score is taken as the muon.

The two candidate-PID BDTs do not use the track length, so that their scores do not bias the momentum spectra. A fourth, cosmic-rejection BDT is available but not used: it brought no improvement over the Pandora topological score, which provides the cosmic rejection at Cut 3.

7.2 Cut definitions

- Cut 0. EventsInTrueFV.** All events; the reference for truth-level studies.
- Cut 1. NeutrinoSlice.** At least one reconstructed track and the software trigger. The cosmic-BDT requirement at this stage is disabled.
- Cut 2. InFiducialVol.** The space-charge-corrected reconstructed neutrino vertex lies inside the FV.
- Cut 3. Topological.** Topological score > 0.67 . This Pandora multivariate score separates cosmic-ray from neutrino topologies and provides the main cosmic rejection: the beam-off yield falls by a factor of 28–29 between Cut 2 and Cut 3 (Table 10).
- Cut 4. MuonCandidate.** At least one primary track satisfies the muon criteria of Table 8. Of the qualifying tracks, the one with the highest muon-BDT score (§7.1) is the muon candidate; the track length enters only through the 10 cm requirement and as a BDT input.
- Cut 5. ContainedPion.** Exactly one primary track other than the muon candidate satisfies the pion criteria of Table 9 and ends inside the containment volume. The track with the highest LLR PID score is the pion candidate.
- Cut 6. MuonIn3Planes.** The muon candidate has at least one hit on each of the U, V and Y planes.
- Cut 7. PionIn3Planes.** The pion candidate has at least one hit on each of the U, V and Y planes.
- Cut 8. ShowerCut.** No primary electromagnetic shower, unless there is exactly one shower with fewer than 50 Y-plane hits more than 10 cm from the vertex, as expected for a delta ray or a low-energy nuclear photon.
- Cut 9. OpeningAngle.** The reconstructed muon–pion opening angle satisfies $\theta_{\mu\pi} < 2.6$ rad ($\approx 149^\circ$), which removes the large-angle cosmic-ray and dirt background. The same requirement is part of the signal definition (§5).
- Cut 10.**
Nonprotons (final cut). Fewer than four primary tracks with LLR PID > 0.1 (non-proton-like), and fewer than five primary tracks in total.

Table 8: Muon candidate PID thresholds (applied at Cut 4).

Variable	Threshold	Description
Track score	> 0.80	Track-like PFP
Vertex distance	≤ 4.0 cm	Attached to ν vertex
Track length	> 10.0 cm	Minimum track length
LLR PID	> 0.20	Muon-like
MIP BDT	≥ -0.10	Minimum-ionising-particle

Table 9: Pion candidate PID thresholds (applied at Cut 5). Endpoint containment requires the track endpoint inside the containment volume. The Bragg-pion score is the pion-hypothesis Bragg-peak likelihood score (rejects stopping protons).

Variable	Threshold	Description
Track score	≥ 0.50	Track-like PFP
LLR PID	> 0.10	Pion / non-proton-like
Vertex distance	≤ 4.0 cm	Attached to ν vertex
Track length	> 20.0 cm	Minimum track length
MIP BDT	> -0.10	MIP-like
Pion BDT	> -0.10	Pion-like (vs. proton)
Bragg-pion score	≥ 0.08	Proton rejection
Endpoint contained	yes	Inside containment volume

7.3 Kinematic reconstruction

Muon momentum. The muon momentum uses the range when the muon-candidate track ends inside the detector and multiple Coulomb scattering (MCS) [14] when it exits, since the range of an exiting track measures only the distance to the detector boundary. In the selected FHC sample 37% of the muons are contained and 63% exit, so the muon momentum is mostly measured by MCS: NuMI muons are energetic and the TPC is small. This matters for the resolution and for the uncertainty, since the two estimators have different resolutions and systematic effects. The MCS estimator has a -10% bias and 30% resolution (Fig. 24), which enter through the response matrix. The detector variations do not change the MCS momentum scale directly; a dedicated momentum-scale term is evaluated in §11.3.

Pion momentum. The pion momentum starts from the range momentum under the muon hypothesis. Before a reconstructed p_π bin is assigned, it is converted to the pion hypothesis at fixed kinetic energy, $T = \sqrt{p_{\mu\text{-hyp}}^2 + m_\mu^2} - m_\mu$ and $p_\pi = \sqrt{(T + m_\pi)^2 - m_\pi^2}$. The stored branch is unchanged; the conversion is part of the reconstructed-bin definition in the bin configuration, and it raises the reconstructed-to-true ratio of the lowest bin from 0.888 to 0.976 (Sec. 2.1 of the technical supplement). Because the muon mass (105.7 MeV) is close to the pion mass (139.6 MeV), the range-to-momentum conversion is a good estimate for a pion that stops, and it yields a momentum that is binned against the true momentum on the same axis. For a pion that interacts before stopping, the visible range is truncated and the estimate is biased low; this truncation, more than the mass hypothesis, limits the observable (§5). The kinetic energy from the proton-hypothesis range is not used: binned with the momentum edges it would shift the migration matrix systematically (a true 0.2 GeV/c pion has $KE \approx 0.104$ GeV, two bins lower, and a true 0.175 GeV/c pion would fall below the first reconstructed bin). The range momentum is only meaningful for contained tracks, which the containment requirement ensures.

Angles. Polar angles are measured with respect to the neutrino direction $\hat{u}_\nu$: $\cos\theta_\mu = \hat{u}_\mu \cdot \hat{u}_\nu$ and $\cos\theta_\pi = \hat{u}_\pi \cdot \hat{u}_\nu$. For reconstructed quantities $\hat{u}_\nu$ is the β direction of Refs. [1, 2], from the NuMI target to the reconstructed vertex; at truth level it is the neutrino direction of the event (§3.2). The opening angle $\theta_{\mu\pi} = \arccos(\hat{u}_\mu \cdot \hat{u}_\pi)$ is computed from the reconstructed track directions and does not depend on $\hat{u}_\nu$.

7.4 Proton-tagged subsample

Requiring in addition a contained proton candidate with $p_p > 0.3$ GeV/c selects a proton-tagged subsample with one further cut. It gives access to hadronic-mass and transverse-kinematic-

imbalance observables but carries its own proton-assignment and model-dependence requirements; it is described in the companion note (Sec. 2 of the proton-tagged note) and is not part of the approval request.

8 Selection performance and backgrounds

Figure 4 shows the cut-flow yields at full exposure for the three configurations, and Table 10 the yields by category after each cut. The final selection keeps 17.5% of the signal at 58.8% purity in FHC (17.6% and 57.7% in RHC, 17.6% and 58.2% combined).




Figure 4: Cut-flow yields at *full exposure* for FHC, RHC and combined (log scale; MC-prediction stack: signal blue, background orange, EXT grey, dirt green; colourblind-safe Okabe-Ito palette), with signal and background scaled per run to the recorded exposure. The detailed breakdown by category is in Table 10.

Table 10: Cut-flow yields at full exposure for FHC, RHC and combined, from the simulation (the expectation of the fake data). The yields carry the full central-value weight, $\text{tune} \times \text{PPFX} \times \text{normalisation}$ (§4, with Eq. (4)), as in the fake-data throw and the unfolding, so the final-cut yields agree with Table 28 to 1%. Signal and “Bkg” (all other neutrino interactions: out of FV, NC, $\text{CC}1\pi^0$, $\text{CC}0\pi$, $\text{CC}N\pi$, CCK and other CC) come from the ν_μ overlays scaled per run to the recorded exposure (Table 4); EXT is the beam-off data scaled by the per-run ratio of beam-on to beam-off gates times 0.98, each run period using its own beam-off sample (§4); dirt is the GENIE NuMI dirt overlay. Pred. = Signal + Bkg + EXT + Dirt. Eff. is the cumulative signal efficiency relative to all generated signal (the None stage); Pur. = Signal/Pred.

Cut	Signal	Bkg MC	EXT	Dirt	Pred.	Eff. [%]	Pur. [%]
<i>FHC (Runs 1,2,4,5; 7.766×10^{20} POT)</i>							
None	4 753	235 281	3 256 689	29 494	3 526 216	100.0	0.1
InFiducialVol	3 955	52 835	122 860	1 546	181 195	83.2	2.2
Topological	2 995	24 262	4 299	168	31 724	63.0	9.4
MuonCandidate	2 689	16 484	1 870	113	21 156	56.6	12.7
ContainedPion	1 195	1 823	173	21	3 212	25.1	37.2
MuonIn3Planes	1 143	1 722	144	14	3 023	24.1	37.8
PionIn3Planes	1 062	1 549	113	6.0	2 730	22.3	38.9
ShowerCut	896	783	83	5.1	1 768	18.9	50.7
OpeningAngle	890	710	32	1.6	1 634	18.7	54.5
Nonprotons	833	560	23	1.1	1 418	17.5	58.8
<i>RHC (Runs 1,2,3,4a,4b,4c; 1.108×10^{21} POT)</i>							
None	5 778	286 540	3 560 975	26 089	3 879 381	100.0	0.1
InFiducialVol	4 828	65 474	139 122	1 367	210 791	83.6	2.3
Topological	3 743	30 446	4 767	149	39 104	64.8	9.6
MuonCandidate	3 363	20 709	2 090	100	26 262	58.2	12.8
ContainedPion	1 488	2 576	191	19	4 274	25.7	34.8
MuonIn3Planes	1 426	2 430	158	13	4 028	24.7	35.4
PionIn3Planes	1 329	2 194	121	5.3	3 650	23.0	36.4
ShowerCut	1 101	1 026	86	4.5	2 218	19.1	49.7
OpeningAngle	1 093	933	33	1.5	2 060	18.9	53.1
Nonprotons	1 015	720	23	1.0	1 759	17.6	57.7
<i>Combined FHC+RHC (1.885×10^{21} POT)</i>							
None	10 530	521 821	6 817 663	55 583	7 405 598	100.0	0.1
InFiducialVol	8 783	118 309	261 981	2 913	391 986	83.4	2.2
Topological	6 738	54 707	9 066	317	70 828	64.0	9.5
MuonCandidate	6 052	37 193	3 960	213	47 418	57.5	12.8
ContainedPion	2 682	4 399	364	40	7 485	25.5	35.8
MuonIn3Planes	2 570	4 152	302	27	7 051	24.4	36.4
PionIn3Planes	2 391	3 744	234	11	6 380	22.7	37.5
ShowerCut	1 997	1 809	169	9.5	3 985	19.0	50.1
OpeningAngle	1 984	1 643	65	3.1	3 694	18.8	53.7
Nonprotons	1 848	1 280	46	2.1	3 177	17.6	58.2

Phase space and opening-angle requirement. The thresholds were chosen from the efficiency turn-on and purity studies of the Run-1 FHC sample (Sec. 14.1 of the technical supplement); the measured phase space excludes the momentum regions of near-zero efficiency and the low-purity region of large opening angle. The events removed by the requirement $\theta_{\mu\pi} < 2.6$ rad form the cosmic control region (Table 12): in FHC 5.1 signal, 60.1 neutrino-background, 40.3 beam-off and 1.9 dirt events, a signal purity of 4.7%. Without the requirement the purity of the selected sample would fall from 58.8% to 55.0% and its beam-off content would rise from 23 to 64 events. No lower bound on the opening angle is applied, since the small-angle region has a purity of 58% and a cut there would remove signal without improving the purity.

8.1 N–1 distributions

Figure 5 shows the N–1 distributions (all cuts applied except the one shown) of the topological score (Cut 3) and the muon–pion opening angle (Cut 9), the two cuts instrumented event by event, for the three configurations. The shapes agree between the horn modes, as expected for an identical selection; the RHC and combined samples are larger. The remaining cut variables are shown at the final cut in §8.3.




Figure 5: N–1 distributions for the topological score (Cut 3, top) and μ – π opening angle (Cut 9, bottom), for FHC, RHC and combined. Signal (blue) and background (orange) stacked; dashed line = applied threshold. Colourblind-safe Okabe-Ito palette.

8.2 Proton rejection for the pion candidate

Figure 6 shows the two main proton-rejection variables of the pion candidate for simulated pions and protons at the final cut, the distributions used to set the working point.






Figure 6: Pion-candidate PID distributions for true $\pi^\pm$ (red) and true protons (blue) at the final cut stage. The Bragg-pion score provides the proton rejection; χ_p^2 is shown as a cross-check.

8.3 Distributions at the final cut


Figures 7–9 show the particle-identification and track-length variables of the two candidates and two topological variables at the final cut, for FHC and RHC.




(a) Muon LLR PID, FHC.



(b) Muon LLR PID, RHC.




(c) Muon length, FHC.




(d) Muon length, RHC.


Figure 7: Muon-candidate LLR PID (top) and track length (bottom) at the final cut, for FHC and RHC (the combined configuration is consistent with both and is not shown). Signal (blue) / background (orange) stacked.




(a) Pion LLR PID, FHC.



(b) Pion LLR PID, RHC.



(c) Pion length, FHC.



(d) Pion length, RHC.

Figure 8: Pion-candidate LLR PID (top) and track length (bottom) at the final cut, for FHC and RHC (the combined configuration is consistent with both and is not shown). Signal (blue) / background (orange) stacked.




Figure 9: Topological variables at the final cut, fake data and simulation.

8.4 Background composition

Table 11 gives the composition of the background at the final cut.


Table 11: Background composition at the final cut, FHC full exposure: truth classes of the selection's counters, weighted with the central-value weight and scaled per run to the recorded exposure (the signal-region column of Table 25 of the technical supplement). Fractions are of the 560 neutrino-background events; the beam-off and dirt yields are listed for scale.

Category	Events	Fraction
$\nu_\mu CC$ 0π (proton or other track as pion)	215	38%
$\nu_\mu CC \geq 2\pi^\pm$, no π^0	128	23%
NC and $\nu_e CC$	79.0	14%
$\nu_\mu CC$ with a π^0	51.4	9%
$\nu_\mu CC$ $1\pi^\pm$ outside the measured phase space	46.5	8%
Out of fiducial volume	41.3	7%
EXT (cosmic)	23	n/a
Dirt	1.1	n/a

The largest background, 38% of the neutrino background, is $CC0\pi$ interactions in which a proton is reconstructed as the pion candidate. Event displays labelled with the simulated particle type confirm this: the pion-candidate track is a stopping proton with a Bragg peak, while the muon is identified correctly. At the working point (Bragg-pion score ≥ 0.08) the pion candidate is a proton in about 18% of the selected events. Tighter Bragg-score thresholds (0.20–0.30) do not improve the figure of merit: the pion-to-proton ratio of the pion candidates in signal events is about 4 : 1, and a tighter cut removes pions at nearly the same rate as protons.

8.5 Event displays

Figure 10 shows two-dimensional event displays, labelled with the simulated particle type, of the two largest backgrounds and of a beam-off event. Tracks are coloured by the matched particle: μ (blue), $\pi^\pm$ (red), p (green), cosmic or unmatched (grey). The muon candidate is drawn thick and solid, the pion candidate thick and dashed, and the reconstructed vertex as a star.



(a) CC0 π background: a proton (green, dashed) is reconstructed as the pion candidate. (b) CCN π background: two genuine charged pions present; one is selected as the pion candidate. (c) EXT cosmic event: all tracks are grey (no truth match), faking a neutrino vertex topology.

Figure 10: Truth-labelled 2D event displays (ZX projection) for representative background events at the final cut. Three views (ZX top, ZY side, XY face) are generated per event.

9 Background control regions

The background prediction is tested in four control regions, each defined by inverting one requirement of the signal selection (Fig. 11). No event can belong to a control region and to the signal region, so the control regions can be compared with beam-on data while the signal region remains blind.




Figure 11: The four control regions and the signal region. Each region inverts one requirement of the signal selection, so no event belongs to a control region and to the signal region; this allows the control regions to be compared with beam-on data while the signal region remains blind.

9.1 Region definitions and expected yields

Each region inverts one requirement of the signal selection; requirements that the inversion leaves undefined are dropped explicitly (branch-level definitions in Table 24 of the technical supplement):

- **CC0 π** : the muon selection without a reconstructed charged pion, enriched in quasi-elastic and $2p2h$ interactions;
- **multi- π** : two or more charged-pion candidates, enriched in resonant and deep-inelastic multi-pion production;
- **π^0** : a reconstructed shower that fails the shower veto, enriched in charged- and neutral-current π^0 production;
- **cosmic**: the full signal selection with the opening-angle requirement inverted ($\theta_{\mu\pi} > 2.6$ rad), enriched in cosmic-ray activity.

The regions use the per-run normalisation and the central-value weights of the measurement (§4). Table 12 gives the expected yields.

Table 12: Control-region yields at full exposure. “Signal” is the signal content of the neutrino simulation, “ ν bkg” the remaining neutrino simulation, EXT the beam-off samples scaled by the gate ratios of their run periods (§4) and dirt the overlay of interactions outside the cryostat; all yields and the pseudo-data apply Eq. (4). The pseudo-data column is a Poisson throw of the neutrino simulation only and is compared with Signal + ν bkg.

Beam	Sideband	Pseudo-data	Signal	ν bkg	EXT	Dirt	Sig. frac.
FHC	CC0 π	10 613	987	9 703	1 080	49	9.2%
FHC	multi- π	102	27	83	5.4	0.1	24.4%
FHC	π^0	5 789	582	5 265	613	10	10.0%
FHC	cosmic	77	5.1	60	40	1.9	7.8%
RHC	CC0 π	12 590	1 206	11 452	1 214	43	9.5%
RHC	multi- π	175	35	124	5.9	0.1	22.0%
RHC	π^0	8 176	780	7 300	678	9.1	9.6%
RHC	cosmic	81	6.4	75	42	1.7	7.9%
Comb	CC0 π	23 203	2 193	21 155	2 294	92	9.4%
Comb	multi- π	277	62	207	11	0.2	23.0%
Comb	π^0	13 965	1 362	12 565	1 291	19	9.8%
Comb	cosmic	158	11.5	136	82	3.6	7.8%

The regions differ in size and in the background they probe:

- **CC0 π** is the largest region, with about 10.6k events in FHC, 12.6k in RHC and 23k combined. Its neutrino component is dominated by quasi-elastic and $2p2h$ interactions, which reach the signal region through pion absorption or a proton reconstructed as the pion. It probes the final-state interactions that move events between $CC\pi^\pm$ and CC0 π , the largest migration path into and out of the signal definition (§5).
- π^0 is the second largest, 5.8–8.2k events per mode and 14k combined. It is enriched in charged- and neutral-current π^0 events, whose photon showers the veto (Cut 8) removes, and it probes the veto efficiency together with the π^0 production rate: these are the events that enter the signal region when a photon is missed.
- **multi- π** is small, 102–175 events per mode (277 combined), and has the largest signal content, 22–24%, because a multi-pion event with one pion lost is the background the selection is most exposed to. Its size limits its constraining power.
- **cosmic** has the largest beam-off share: in FHC 40 beam-off against 65 neutrino events, a ratio of 0.6, against 0.1 in the CC0 π region. It probes the beam-off normalisation and shape. Its beam-off purity is 38% in FHC and 34% in RHC, below the target purity of the frozen procedure (Sec. 13.5 of the technical supplement), so it can add an uncertainty to the beam-off normalisation but not correct it.

The combined CC0 π and π^0 regions have raw statistical precisions of 0.7% and 0.8%. Their constraint on the signal-region background is weaker, because it passes through transfer factors that depend on the model (Sec. 13.3 of the technical supplement): the factors are small (0.03 and 0.02), but the transfer from a CC0 π event to a selected CC1 π topology carries its own cross-section and detector uncertainties, which limit the constraint more than the sample size does. The cosmic region would constrain the beam-off normalisation only to about 60%: the combined region contains 82 beam-off events among about 233, with 15 events of statistical uncertainty and a 32% systematic uncertainty on its 147 neutrino events (± 47 events); the statistics alone would give 19%. With the beam-off background at 1.6% of the signal region this is sufficient. The multi- π region, 277 combined events, gives about 6%, useful as a cross-check.

Pseudo-data test. Before the comparison with beam-on data, the region definitions and the per-run normalisation were tested on pseudo-data, a Poisson throw of the neutrino simulation alone, compared with the neutrino simulation (Table 13). The two large regions agree with unity to 1% in all three configurations. The small regions scatter within their Poisson uncertainties: FHC multi- π at 0.93 on 102 events (10%), FHC cosmic at 1.18 on 77 events (11%, 1.6σ) and RHC multi- π at 1.10 on 175 events (8%). A variant with the beam-off and dirt components added to the pseudo-data tests the complete bookkeeping (Sec. 13.2 of the technical supplement).

Table 13: Pseudo-data test of the control regions. “Signal fraction” is the signal content of the neutrino simulation in each region. $N_{\text{data}}/N_{\nu\text{MC}}$ is the ratio of a Poisson throw of the neutrino simulation to the simulation itself: the two large regions (CC0 π , π^0) are within 1% of unity in the three configurations, and the two small regions (multi- π , cosmic) scatter within their Poisson uncertainties (8–11%). The test checks the region definitions and the per-run normalisation; the background model is tested with beam-on data (§9.2).

Sideband	Signal frac. (of νMC)	$N_{\text{data}}/N_{\nu\text{MC}}$		
		FHC	RHC	Combined
CC0 π	$\sim 9\%$	0.99	0.99	0.99
multi- π	$\sim 23\%$	0.93	1.10	1.03
π^0	$\sim 10\%$	0.99	1.01	1.00
cosmic	$\sim 8\%$	1.18	0.99	1.07

9.2 Comparison with beam-on data

The four regions have been compared with beam-on data at reconstruction level (Fig. 12), under a procedure fixed beforehand: frozen region definitions, a pre-agreed test statistic and the conditions under which a correction may be applied (Sec. 12 of the technical supplement). No beam-on sample of Run 2 exists, and Run 2 is removed from both data and prediction in this comparison.

- **Global test.** The frozen global test passes, $\chi^2 = 8.1/8$ ($p = 0.42$).
- **No class-specific correction.** A joint FHC+RHC fit of the background-class normalisations gives the dominant CC0 π class at 0.949 ± 0.190 over all run periods (0.984 ± 0.191 without FHC Run 1), consistent with unity (Sec. 6 of the control-sample note). The nominal background prediction is kept. The fit uses the CC0 π , π^0 and cosmic regions; the multi- π region, which shares 47 FHC and 68 RHC events with the π^0 region, is diagnostic and not in the fit, so the overlap does not enter the fit covariance.
- **Common excess.** The neutrino-dominated regions lie 12–20% above the prediction, within the correlated flux and cross-section uncertainties. The excess does not follow the signal content of a region: tightening each region to reduce its signal (the CC0 π region from 8.6% to 2.2% signal while keeping 58% of its events) leaves the data-to-prediction ratio unchanged, 1.12 to 1.13 (Sec. 7 of the control-sample note), and the cosmic region approaches unity as its beam-off share increases. The excess therefore lies in the neutrino simulation and is not signal in the control regions.
- **FHC Run 1.** Period by period, the data lie 12–25% above the prediction in every run period except FHC Run 1, which is consistent with unity (Table 13 of the control-sample note). A free normalisation factor on FHC Run 1 fits 0.869 ± 0.022 and improves the per-period fit by $\Delta\chi^2 = 37$, while a free factor on any other period gains at most 15. The common excess of the other periods is covered by the correlated uncertainties; a 13% offset

of one period relative to the others is not. The cosmic region, dominated by beam-off events, is at 0.95 ± 0.17 in Run 1, so the Run-1 trigger count is not obviously wrong; the POT bookkeeping, the completeness of the Run-1 beam-on sample with respect to its run list and the normalisation of the Run-1 overlay remain to be checked (§14.1). The same comparison fixes the Run-1 exposure assignment: normalised to the open-trigger bookkeeping (3.283×10^{20} POT and 9 846 635 triggers), FHC Run 1 lies at 0.60–0.72 in the CC0 π , π^0 and cosmic regions, with a fitted exposure factor of 0.520 ± 0.020 , whereas the non-open-trigger values used here (2.192×10^{20} POT and 5 748 692 triggers) give 0.95–1.10 in the same regions.

- **Muon-momentum shape.** In two dimensions ($\cos \theta_\mu \times p_\mu$) the residual is organised in p_μ : the data over the prediction are 1.16–1.82 in the 0.35–0.75 GeV/ c bins, 1.05–1.40 below and mostly 0.8–1.1 above, in every angular bin, in both modes and in regions of different composition. Normalisation parameters do not remove it. It points to the muon momentum rather than to a background class, remains open, and two-dimensional normalisation fits are not quoted while it stands (Sec. 6 of the control-sample note).

A normalisation fit propagated under the frozen rules cannot reduce the uncertainty of the measurement. Conditioning the prediction on the control-region totals could, through the flux and detector correlations: on an Asimov dataset it reduces the uncertainty of the one-bin total from 35.8% to 11.3%. Adopting such a constraint would change the role of the control regions and requires a collaboration decision.




Figure 12: The control regions with beam-on data: FHC (top) and RHC (bottom), data compared with the neutrino simulation + beam-off + dirt prediction in each region, normalised to the exposure of the run periods with beam-on data (Run 2 has none). The ratios per run period, the fitted normalisations and their correlations are in the control-sample note.

9.3 Validation of the background fit

The background fit was validated on simulation before it was applied to data (Fig. 13).

- **Toys.** Sets of 300 toys give pull means of 0.02–0.07, widths of 0.89–1.03 and coverage of 67–75%; ensembles of up to 10^5 toys per model are described below.
- **Injections** (Asimov data, two-parameter fit of the totals). Shifts of a class are recovered exactly; a 30% signal excess moves the normalisations by 0.13σ , an unfitted class at +30% by -0.07σ , a detector variation by at most 0.14σ , and flux and GENIE universes by -0.97 to $+0.84\sigma$, within the uncertainty they are part of.
- **FHC against RHC.** A likelihood-ratio test of common against beam-specific normalisations inside the joint covariance (a parameter goodness-of-fit test is not valid, since the modes share universes): a +15% shift of the RHC neutrino simulation alone gives $\Delta\chi^2 = 12.9/2$ ($p = 0.002$, totals) and $20.3/2$ ($p = 4 \times 10^{-5}$, two-dimensional).
- **Markov chain.** An adaptive Metropolis sampler with four chains and flat priors gives posterior means within 0.02 and widths within 2% of MIGRAD/MINOS on the Asimov dataset and on a $CC0\pi \times 1.25$ injection, with $\hat{R} \leq 1.004$. With a few parameters and a nearly quadratic likelihood the profile likelihood suffices, and the sampler is kept as a cross-check.




Figure 13: Validation of the background fit on simulation, for the totals-only fit (top) and the two-dimensional fit (bottom): pull distributions from toy ensembles, recovery of injected normalisation shifts, and the profile-likelihood interval against the Markov-chain posterior.

Convergence of the pull means. With 300 toys the uncertainty of a pull mean is 0.058, so a mean of +0.07 cannot be distinguished from zero. Ensembles of 10^5 toys per model (uncertainty 0.003) and three diagnostic variants of 5×10^4 toys each resolve this (Table 14, Fig. 14). With the combined Neyman–Pearson data variance, the pull mean of the two-dimensional fits does not converge to zero: +0.019 remains for $CC0\pi$ and +0.034 and +0.031 for $CC0\pi + \pi^0$, six to ten standard deviations of the ensemble. The offset persists with Gaussian instead of Poisson fluctuations (+0.031), so it does not come from Poisson counting in sparse bins. It disappears when the data variance is the prediction at the nominal parameters (+0.000, -0.004 , -0.001): the combined Neyman–Pearson variance depends on the fluctuating data, which correlates the weight of a bin with its own fluctuation. The bias is small, 0.3–0.7% on the normalisation against a 16–21% uncertainty, but avoidable, so the fit uses the Pearson form, which every data result uses. The pull widths are 1.00 and the coverage 68% in every variant.

Table 14: Pull means $\pm$ their ensemble error. “standard”: systematic throw from the covariance, Poisson, CNP data variance (10^5 toys); “Gaussian”: Gaussian instead of Poisson fluctuation; “Pearson”: standard toys fitted with the prediction as data variance (5×10^4 each).

Bins	Model (parameter)	standard, CNP	Gaussian, CNP	standard, Pearson
totals	CC0 π	$+0.003 \pm 0.003$	$+0.005 \pm 0.005$	-0.011 ± 0.005
totals	CC0 $\pi + \pi^0$ (CC0 π)	$+0.012 \pm 0.003$	$+0.014 \pm 0.005$	$+0.003 \pm 0.005$
totals	CC0 $\pi + \pi^0$ (π^0)	$+0.016 \pm 0.003$	$+0.014 \pm 0.005$	$+0.008 \pm 0.005$
2D	CC0 π	$+0.019 \pm 0.003$	$+0.031 \pm 0.005$	$+0.000 \pm 0.005$
2D	CC0 $\pi + \pi^0$ (CC0 π)	$+0.034 \pm 0.003$	$+0.031 \pm 0.005$	-0.004 ± 0.005
2D	CC0 $\pi + \pi^0$ (π^0)	$+0.031 \pm 0.003$	$+0.029 \pm 0.005$	-0.001 ± 0.005




Figure 14: Running pull mean against the ensemble size for each model and parameter; band: statistical error of the standard ensemble. The ensemble thrown with statistical fluctuations only, and no systematic variation, has narrower pulls as expected.

10 Cross-section extraction

10.1 Extraction chain

The extraction proceeds in three stages, shown for each configuration in the reconstructed-bin space in which the unfolding operates, with all analysis slices concatenated.

Selected events. The selected reconstructed spectrum is scaled to the exposure of each configuration with the per-run factors of §4, and the simulation is split into signal and background (Fig. 15). At the final cut the purity is 58–59% (Table 10), so about 40% of the selected events are background. The points are the fake data.

Background subtraction. The estimated background is subtracted bin by bin (Fig. 16). The neutrino background is taken from the simulation, the cosmic-ray background from the beam-off data and the background from interactions outside the cryostat from the dirt overlay; the control regions of §9 test this prediction. For observable x the background-subtracted spectrum in reconstructed bin i is

$$d_i^{\text{sig}} = d_i^{\text{data}} - b_i^{\nu\text{-bkg}} - b_i^{\text{EXT}} - b_i^{\text{dirt}}. \quad (9)$$

Unfolding. The subtracted spectrum still contains the detector response and the selection efficiency, described by the response matrix of §10.2. With up to half of the selected events of a bin coming from other true bins, a direct inversion is unstable (§10.3), and the spectrum is unfolded with the Wiener-SVD method. The regularisation introduces the additional-smearing matrix A_C : the unfolded result estimates A_C times the true signal, and every prediction is multiplied by A_C before it is compared with the result (§13.4).




Figure 15: Stage 1, FHC: selected reconstructed spectra scaled to the full exposure, with the MC split into signal and background. Points are the Poisson-thrown fake data. The RHC equivalent is Fig. 17 and the combined one Fig. 11 of the technical supplement.




Figure 16: Stage 2, FHC: background-subtracted data compared with the central-value MC signal; the input to the unfolding. RHC: Fig. 18; combined: Fig. 12 of the technical supplement.




Figure 17: Stage 1, RHC: selected reconstructed spectra scaled to the full RHC exposure, MC split into signal and background.




Figure 18: Stage 2, RHC: background-subtracted data compared with the central-value MC signal.

10.2 Response matrix

Two matrices are used, and they differ by the efficiency. The migration matrix

$$M_{ij} = P(\text{reco in bin } i \mid \text{true in bin } j, \text{ selected}), \quad \sum_i M_{ij} = 1, \quad (10)$$

describes how selected events move between bins. The smearceptance matrix also contains the probability to be selected,

$$S_{ij} = \varepsilon_j M_{ij}, \quad \sum_i S_{ij} = \varepsilon_j, \quad (11)$$

where ε_j is the ratio of selected to generated signal events in true bin j . The unfolding uses S , built from selected signal events with the same binning on both axes, and the efficiency of Fig. 23 is its column sum. The folding relation that is inverted is

$$n_i^{\text{reco}} = \sum_j S_{ij} N_j^{\text{true}} + b_i, \quad S_{ij} = \varepsilon_j M_{ij}, \quad (12)$$

so the unfolded $\hat{N}_j^{\text{true}}$ is already corrected for the efficiency, and no separate division by ε_j follows (Eq. (13)). Figure 22 shows S as passed to the unfolding, with each column summing to the efficiency of its true bin. The migration diagonals quoted in the text, for example 91%, 51% and 52% for the three p_π regions, are elements of the column-normalised M , obtained by dividing each column of S by its sum.

10.3 Regularised unfolding

Unfolding estimates the true signal spectrum $\vec{s}$ from the measured $\vec{m}$ through the folding relation $\vec{m} = A\vec{s} + \vec{b}$. The direct inverse $A^{-1}(\vec{m} - \vec{b})$ is unbiased, but bin migration makes A nearly singular, and the inverse turns statistical fluctuations into large anticorrelated oscillations between bins. Regularisation reduces this variance at the cost of a controlled bias. Two methods are used.

Wiener-SVD (primary). After a singular-value decomposition of the response, each mode k is weighted by the Wiener filter $W_k = \sigma_k^2 / (\sigma_k^2 + N_k)$, the weighting that minimises the mean squared error: it is close to one for well-measured modes and tends to zero for modes dominated by noise, and its strength is not tuned by hand. The resulting linear bias is known: the unfolded data estimate the smeared spectrum $A_C\vec{s}$, and the additional-smearing matrix A_C is provided with every result so that each prediction is smeared in the same way before comparison.

D’Agostini iterative Bayesian [27] (cross-check). Bayes’ theorem inverts the migration, $P(C_j|E_i) = R_{ij}\pi_j / \sum_l R_{il}\pi_l$, and the unfolded estimate is $\hat{s}_j = \varepsilon_j^{-1} \sum_i P(C_j|E_i)(m_i - b_i)$. The prior π_j , initially the simulation, is replaced by the result and the procedure iterated; the number of iterations is the regularisation parameter, with few iterations biased towards the prior and many iterations noisy. There is no A_C , and predictions are compared directly. The residual bias is towards the prior and can only be evaluated with fake data that differ from the prior. The uncertainty is the data covariance propagated through the iterations at fixed prior.

Table 15: Trade-offs of the two unfolding methods. Neither is “more correct”; they make complementary choices about how to trade bias for variance and how to expose the residual bias.

	Wiener-SVD	D’Agostini (iter. Bayes)
Regularisation	Wiener filter (automatic, from S/N)	Iteration count n (chosen)
Residual bias	Explicit A_C , published	Implicit; set by early stopping
Result integral	Suppressed, observable-dependent	Approximately preserved in nominal closure; prior-biased in general
Theory comparison	Smear theory by A_C	Raw theory
Positivity	Not guaranteed	Guaranteed for positive input
Uncertainties	Closed-form, linear	Correlated across iters.; needs care

Wiener-SVD provides the results, for its objective regularisation and its explicit bias A_C . D’Agostini serves as a cross-check (Sec. 6.4 of the technical supplement). Its integral is independent of the iteration count and of the observable, and lies 30–35% above the Wiener-SVD integrals averaged over the observables, depending on the configuration; this is consistent with the spread of the Wiener-SVD integrals between observables (Table 24) being caused by the A_C smearing, which the row sums of §13.4 show directly. With fake data equal in expectation to the central-value simulation that also serves as its prior, D’Agostini returns its prior, so the stability of its integral does not measure its bias for a different truth; that requires fake data from a different model (§12.3).

10.4 Wiener-SVD implementation and normalisation

The Wiener-SVD unfolding [7] of the XSecAnalyzer framework [18, 17] produces every result: the data covariance is whitened by a Cholesky decomposition, the regularisation is applied inside the SVD of RC^{-1} , and the Wiener filter of Eq. 3.24 of Ref. [7] is set by the data covariance. The regularisation matrix C is the second-derivative operator on the density, corrected for the non-uniform bin widths, with one-sided first-derivative rows at the ends. The width correction matters because the bin widths vary by up to a factor of 56 ($\cos\theta_\mu$; 14 for p_μ): a stencil for a uniform grid weights the curvature of the density wrongly and drives A_C towards rank one, so that every model is smeared onto one shape and the result carries no shape information (Sec. 8.1 of the technical supplement). A standalone replica of the unfolding is used for cross-checks between pipelines; it takes the covariance from the framework because it generates no universes, and with a statistical covariance alone it under-regularises and over-corrects the integral by about 26%, so it is not used for any result.

The differential cross section in true bin j is

$$\left.\frac{d\sigma}{dx}\right|_j = \frac{\hat{N}_j^{\text{true}}}{\Delta x_j \cdot \Phi_{\text{total}} \cdot N_{\text{Ar}}}, \quad (13)$$

where $\hat{N}_j^{\text{true}}$ is the Wiener-SVD estimate of the true signal count obtained by inverting Eq. (12), $\Phi_{\text{total}} = \Phi_{\text{per POT}} \times T_{\text{POT}}$ the integrated flux and N_{Ar} the number of argon nuclei in the FV. Since the unfolding uses $S = \varepsilon M$, the efficiency correction is part of the inversion and no factor $1/\varepsilon_j$ appears. The implementation forms each bin as $\hat{N}_j^{\text{true}}/(\text{conv} \cdot \Delta x_j)$ with the conversion $\text{conv} = N_{\text{Ar}} \Phi_{\text{total}}/10^{-38}$, whose inputs are:

Quantity	FHC	RHC	Combined
N_{Ar} (FV, unresponsive region excluded)		8.710×10^{29}	
Flux constant ϕ ($E_\nu > 60$ MeV; 10^{-10} cm $^{-2}$ /POT)	6.81159	6.44646	6.59691
Data POT (10^{20})	7.766	11.082	18.848
Integrated flux $\Phi_{\text{total}} = \phi \times \text{POT}$ (10^{11} cm $^{-2}$)	5.290	7.144	12.434
$\text{conv} = N_{\text{Ar}} \Phi_{\text{total}}/10^{-38}$ (events per 10^{-38} cm 2 /Ar)	4608	6223	10830

The integrated cross section of a differential result is the width-weighted sum

$$\sigma_{\text{int}} = \sum_j \left.\frac{d\sigma}{dx}\right|_j \Delta x_j. \quad (14)$$

Since the Wiener-SVD result estimates $A_C x_{\text{true}}$ and A_C does not preserve the normalisation, σ_{int} differs between observables (§13.4). The total cross section, independent of the observable, is the one-bin extraction of §13.6.

10.5 Reconstructed spectra, response and efficiency

Figures 19–21 show the reconstructed spectra of the five observables at the final cut, with the simulated signal and the background components.




Figure 19: Reconstructed-space spectra at the final cut for all five observables, FHC full exposure (Runs 1+2+4+5): Poisson-thrown central-value fake data (black points), stacked signal (blue), beam background (orange), dirt (green), and EXT (grey). The fake data is drawn from the central-value Monte Carlo and so agrees with the stacked prediction by construction (a reco-level self-closure); the differential results and their closure against the A_C -smeared truth are shown in Fig. 26.




Figure 20: Reconstructed-space spectra, RHC full exposure (Runs 1+2+3+4a+4b+4c); same components as Fig. 19.






Figure 21: Reconstructed-space spectra, combined FHC+RHC (per-mode normalised); same components as Fig. 19.

Figure 22 shows the smearacceptance matrices of the five observables. The column-normalised diagonals of the FHC responses are 0.86–0.91 in the best-resolved bin of each angular observable (the forward bin of $\cos \theta_\mu$ and the first bin of the others) and 0.66–0.88 elsewhere in $\cos \theta_\mu$, $\cos \theta_\pi$ and $\theta_{\mu\pi}$. For the six p_μ bins they are 0.75, 0.64, 0.58, 0.59, 0.58 and 0.57: the MCS estimator, which measures 63% of the selected muons, has a resolution of about 30%, comparable to the widths of the p_μ bins. Every binning meets the 0.50 criterion in FHC and in RHC (App. A), and the RHC diagonals differ from the FHC ones by at most 0.04.


At any fine binning the p_π response is far from diagonal, because the range estimator saturates near 0.34 GeV/c once pions interact before stopping; the fine-bin study is in Sec. 2.1 of the technical supplement. The three-region response has column-normalised diagonals of 91%, 51% and 52%. It separates the two intervals near threshold from the region above them; the upper region is an inclusive region rather than a resolved point of a spectrum (§13.3).




(a) p_μ




(b) p_π



(c) $\cos \theta_\mu$



(d) $\cos \theta_\pi$



(e) $\theta_{\mu\pi}$

Figure 22: Smear acceptance matrices $S_{ij} = \varepsilon_j M_{ij}$, FHC, exactly as passed to the unfoldr, at the binning of the results, including the three-region p_π response (91%, 51% and 52% column-normalised diagonals); the rejected five-bin p_π response, which shows the reconstructed events piling into the lowest bin from every true bin, is Fig. 13 of the technical supplement.

Figure 23 shows the selection efficiency in each true bin, the ratio of selected to generated signal events.




Figure 23: Selection efficiency as a function of true observable value, in the measured phase space ($p_\mu > 0.15$, $p_\pi > 0.175 \text{ GeV}/c$, $\theta_{\mu\pi} < 2.6 \text{ rad}$). Efficiency is the ratio of selected to generated signal MC ($\varepsilon_j = \sum_i S_{ij}$, summed over reconstructed bins). With the momentum thresholds and the 2.6 rad opening-angle cut the efficiency is flat at 13–20% across all observables; the low-momentum turn-on and the large-angle region dominated by cosmic rays lie outside the measured phase space. The lowest point is the smallest opening-angle bin (9% efficiency), which is kept because its purity is high (58%).

10.6 Reconstruction resolution

Figure 24 shows the resolution of each observable for selected signal events. For the momenta the relative residual $(\text{reco} - \text{true})/\text{true}$ is shown, for the angles the absolute residual $(\text{reco} - \text{true})$, since the relative residual diverges as $\cos \theta \rightarrow 0$ or $\theta_{\mu\pi} \rightarrow 0$. The red points are the mean residual in each true bin; the quoted numbers are the overall mean (bias) and RMS (resolution).




Figure 24: Reconstruction resolution vs. true value, selected signal. Momenta: relative ($\text{reco} - \text{true}$)/ true ; angles: absolute ($\text{reco} - \text{true}$). Red points: mean residual per true bin. The muon MCS momentum slightly over-estimates at low p_μ and increasingly under-estimates above ~ 1 GeV/ c ; the range estimator is badly biased for *exiting* muons (range = distance to wall), which is why the analysis uses range only for contained muons (range-or-MCS). The pion panel shows the estimator the analysis actually bins on, not the raw muon-hypothesis range momentum (Sec. 2.1 of the technical supplement): signed mean bias $b = -31\%$, RMS 27% . That average is dominated by high momenta; b is only -2% in the first true bin and reaches -65% above 0.8 GeV/ c , because the reconstructed value saturates near 0.34 GeV/ c once pions interact hadronically before stopping. The angular variables are essentially unbiased with narrow resolution.

11 Systematic uncertainties

Every result carries flux, interaction-model, hadron re-interaction, detector, POT, target-count and statistical terms. Each is a covariance matrix built from the universes or variation samples of its source, propagated through the unfolding with the error-propagation matrix and provided per component (§13.4); Table 16 lists the sources and their sizes. Four identified effects lie outside this configured covariance: the MCS momentum scale, evaluated in §11.3 and provided as a separate matrix for p_μ , the beamline-geometry flux term, the beam-off gate ratios and the Run-2 beam-off substitution (§11.4).

Table 16: Systematic sources and their bin-averaged fractional contribution, computed from the per-component covariance matrices of the results. Ranges span observables and configurations; the two families are given separately because the proton-tagged sample is both statistically sparser and more detector-sensitive.

Source	Branch / method	Inclusive (%)	Proton-tagged (%)
Flux (PPFX multisims)	weightsFlux	22.8–38.5	17.6–74.6
Detector response	Dedicated samples	16.6–34.2	9.9–74.5
Cross-section model	weightsGenie	8.4–16.6	6.7–53.8
Hadron re-interaction	weightsReint	3.5–13.3	3.2–33.5
POT counting	Beam toroids	2.8–4.5	2.1–8.7
Target count	FV geometry	1.4–2.2	1.1–4.4
MC statistics	universe spread	1.9–6.1	2.3–19.8
EXT statistics	beam-off sample	0.5–2.3	0.6–3.5
Data statistics	thrown fake data	6.6–19.3	7.8–60.1
Prediction total	quadrature sum	32.6–56.7	22.0–120.2
Total	incl. data stats	33.7–60.0	23.8–124.4

Flux and detector response are the largest terms in both families; which of the two dominates depends on the observable and configuration (Table 17; Table 21 of the technical supplement–Table 22 of the technical supplement).

11.1 Breakdown by source

Table 17 gives the bin-averaged fractional uncertainty of each source group for the FHC extractions; the RHC and combined tables, from the same dumps, are Table 21 of the technical supplement and Table 22 of the technical supplement, and the values agree with Table 21.

Table 17: FHC (Runs 1, 2, 4, 5): bin-averaged fractional uncertainty (%) by source group for the five extractions, from the systematics dumps of the extraction. “Prediction total” excludes the data statistics; the flux row is the PPFX multisim term. Every row is a bin average of the per-bin fractional uncertainties, so the rows do not combine in quadrature to the totals (the quadrature is taken per bin before averaging).

Source	p_μ	p_π	$\cos \theta_\mu$	$\cos \theta_\pi$	$\theta_{\mu\pi}$
Prediction total	39.3	56.7	39.6	39.0	37.1
Cross section (GENIE)	11.9	16.6	10.9	11.4	10.6
Flux (PPFX)	30.0	38.5	26.6	27.1	28.8
Detector	21.2	34.2	25.4	23.2	19.9
Reinteraction	4.2	13.3	3.5	5.1	4.1
MC stat	3.8	6.1	5.2	4.2	2.8
EXT stat	1.1	2.3	1.7	1.2	1.0
Data stat	12.1	19.3	16.6	13.6	8.9
POT + targets	3.8	5.0	3.5	3.4	3.6
Total (incl. data stat)	41.2	60.0	43.1	41.5	38.3

The NuMI flux uncertainty has a hadron-production part, the PPFX multisim term that dominates the budget, and a beamline-geometry part. The latter is built from the twelve geometry variations stored as weights in the overlay (horn current, horn positions, beam spot size and position, horn water layer and target position), each treated as a two-sided unisim and summed in quadrature.

Evaluated separately it is 1.6–3.3% across the observables. It is not part of the covariance, whose configuration omits the beamline block; added in quadrature to the 28% PPFX term it would raise the flux row by at most 0.2% absolute. Figure 25 shows the breakdown for the three configurations.




Figure 25: Bin-averaged fractional systematic uncertainty by source vs. observable, for FHC (left), RHC (centre) and combined (right); the figure form of Table 20 of the technical supplement–Table 22 of the technical supplement. Colours use the Okabe-Ito colourblind-safe palette with distinct markers per source: total (black), flux (blue), detector (green), cross section (vermillion), reinteraction (purple), data stat (orange).

11.2 Flux uncertainty and background subtraction

The NuMI flux normalisation is uncertain by about 17%, while the flux term of the one-bin total cross section is 28–30% (Table 28) and that of the differential results 23–30% (Tables 21–23; 28–39% for the three-region p_π). The difference comes from the background subtraction, not from the unfolding: the one-bin total has no regularisation and carries the same factor. A flux universe scales the whole reconstructed prediction, signal and neutrino background together, and the covariance is that of the measured quantity $(N_{\text{pred}} - B)/(\varepsilon \Phi N_{\text{Ar}})$, with B and Φ moving together in every universe. Relative to the signal, the flux variation is therefore

$$\frac{\delta\sigma}{\sigma} = \frac{\delta\Phi}{\Phi} \left(1 + \frac{B}{S}\right), \quad (15)$$

with B the neutrino background; the beam-off and dirt contributions carry no flux weight. For FHC, $B/S = 560/833$ (Table 10) and a reconstructed-level flux term of 17.7% give 29.6%, against 29.9% for the one-bin extraction. For RHC, $15\% \times (1 + 720/1015) = 25.6\%$ against 27.7%, and for the proton-tagged FHC total, a sample of different purity, $17.7\% \times (1 + 298/331) = 33.6\%$ against 35.5%. The differential flux terms show the same effect bin by bin; the three-region p_π , whose response has the lowest diagonals, carries the largest (39% in FHC).

The flux term is thus fully correlated between the normalisation and the subtracted background, which is correct when both come from the same flux prediction: 17% is the normalisation part, and the remainder is the flux dependence of the background model. It can be reduced by a higher purity or by a background constraint from the control regions (§9), not by the regularisation. The reconstructed-level flux term is 16.4–17.7% bin-averaged and flat across observables, as expected of a normalisation term; the 600 PPFX universes reweight the same events, so the term carries no simulation-statistics component.

11.3 Muon momentum scale

The momentum of 63% of the selected muons is measured by multiple Coulomb scattering, whose scale and resolution are not varied by the detector-variation samples. The momentum of a selected muon is exactly the range value when the track is contained and exactly the MCS value when it exits, so the MCS component can be scaled on its own. A momentum-scale error is a difference between data and simulation, so the variation is applied to the data side: the fake data are thrown again with every MCS-measured muon momentum multiplied by 1 ± 0.05 , while

the simulation, the response and the selection stay nominal, and each variation is unfolded with the nominal response. The half-difference of the two results is taken as a fully correlated term. The size of the variation, $\pm 5\%$, is to be confirmed by a data-driven comparison of the MCS and range momenta of contained muons (§14.1).

For FHC $d\sigma/dp_\mu$ the half-difference is 5–8% of the bin value in the five lowest bins, $0.11\text{--}0.19\sigma$ of the per-bin uncertainty, and 17% with the opposite sign in the open highest bin, 0.56σ : an upward scale moves exiting muons into that bin, whose spectrum falls steeply. The integral moves by -2.3% and -3.4% for the upward and downward scale. The combined configuration, in which FHC and RHC are scaled together as one common MCS scale, behaves like FHC: 7–8% in the five lowest bins and 15% in the highest, at most 0.42σ , with the integral moving by -2.1% and -0.7% . In RHC the five lowest bins move by 10–13% ($0.23\text{--}0.30\sigma$) and the highest by 10% (0.21σ); the RHC integral moves by -8.2% and $+0.9\%$. Adding the term raises the per-bin total uncertainty from 31–47% to 35–47% (FHC, with 31% to 35% in the highest bin), from 34–43% to 37–44% (combined) and from 41–49% to 43–50% (RHC). The term is provided for the three p_μ extractions as `cov_MCSscale.txt`, together with `cov_total_plusMCS.txt`, and is not yet included in the covariance computed by the framework.

The selection applies no requirement on the reconstructed muon momentum, and no other inclusive observable depends on it. The scaled fake data therefore give results identical to the nominal ones for $\cos\theta_\mu$, $\cos\theta_\pi$, $\theta_{\mu\pi}$, p_π and the total cross section, as verified by extraction for FHC $\cos\theta_\mu$, $\cos\theta_\pi$, p_π and the total at the earlier binnings. The reconstructed bins of these observables do not use the muon momentum, so this holds at any binning, and the term is zero for these results. The comparison of the control regions with data is also unaffected, since it is made at reconstruction level with the same estimator in data and simulation.

11.4 Effects not included in the covariance

Besides the MCS scale, three effects are evaluated or bounded but not included in the covariance. The beamline-geometry flux term (1.6–3.3%) is discussed above. The covariance contains the statistical uncertainty of the beam-off samples but no uncertainty on the per-run gate ratios, nor on the use of the Run-1 beam-off sample for Run 2, which is 16% of the FHC and 23% of the RHC exposure. With the beam-off background at 1.6% of the selected sample, a 20% uncertainty on either would add 0.3% to the total uncertainty.

12 Validation

The validation uses simulation, where the true answer is known, and has three parts.

1. **Closure under the nominal model.** Fake data thrown from the simulation that built the response are unfolded, and the truth must be recovered. This tests the implementation: the matrices, the binning, the normalisation and the bookkeeping. It is reported with every extraction as the closure χ^2 and is passed.
2. **Statistical coverage.** Many independent Poisson realisations are carried through the whole chain, and the spread of the results is compared with the uncertainty each reports. This is passed for six extractions covering the three configurations.
3. **Robustness to a different model.** Fake data whose truth differs from the model of the response are unfolded, which measures the bias from the regularisation and the nominal response. With reweighted GENIE fake data the largest bias is 1.18σ , in one sparse bin (§12.3); a test with an independent generator has not yet been possible.

The first two parts test the procedure; the third measures how much of a result comes from the data and how much from the model used to unfold it.

12.1 Closure under the nominal model

Unfolding a spectrum thrown from the simulation that built the response detects an inverted matrix, a misaligned binning or a normalisation applied twice. It is reported with every extraction as the χ^2 against the A_C -smeared truth of the throw, and the ratio of unfolded to true integral averages 0.996 over the fifteen inclusive extractions (0.946–1.047). The test cannot detect a wrong model, since the model is the same on both sides, nor a wrongly sized uncertainty, since a single throw does not sample the distribution of throws. The closure χ^2 uses the full covariance, whose systematic terms are common to the fake data and the reference and cancel in the residual while still entering the covariance; its p -values therefore cluster near unity and are not a sensitive test. The statistical coverage is tested with ensembles.

12.2 Ensemble coverage

One hundred independent Poisson realisations of the data of each configuration were each carried through the whole chain, with the universes rebuilt from that realisation, and the pull in each bin, $p_i = (\hat{x}_i - (A_C x_{CV})_i) / \sigma_i$, was formed against the central-value truth smeared by the A_C of that realisation. Six extractions were chosen to span the conditioning and the configurations: p_μ and $\cos \theta_\mu$ in FHC (the A_C row sums of the latter fall to 0.22), $\cos \theta_\mu$ in RHC (0.05), the two-region p_π in the combined configuration, and the one-bin totals of FHC and RHC, which have no regularisation.

Reference: statistical covariance	pull mean	pull width	68% / 95% coverage
$d\sigma/dp_\mu$, FHC	−0.08	0.98	67.9% / 96.4%
$d\sigma/d \cos \theta_\mu$, FHC	−0.08	1.01	68.0% / 95.2%
$d\sigma/d \cos \theta_\mu$, RHC	−0.07	1.07	66.6% / 93.6%
$d\sigma/dp_\pi$ (two regions), combined	−0.09	0.92	74.0% / 95.5%
one-bin total, FHC	−0.10	0.97	67.0% / 95.0%
one-bin total, RHC	−0.21	0.99	64.0% / 94.0%
expected	0	1	68% / 95%

With 100 members each width is determined to about $\pm 13\%$. The widths are consistent with one and the coverage with its nominal value, so the statistical uncertainties are correctly sized in the three configurations, and the strongly suppressed observables behave like the well-conditioned ones. Against the full covariance the pull widths are 0.22–0.35 (0.11 for the totals) with essentially complete coverage, because the ensemble varies only the data statistics while the full uncertainty also contains the flux, detector and cross-section terms, which are the same in every member; this is the origin of the closure p -values near unity. The ensemble means lie below the reference by 0.4–1.1%, at 0.5–1.4 standard deviations of the mean (0.8% at 2.0 for the RHC total), and the per-bin pull means are −0.07 to −0.10 (−0.21 for the RHC total). This common negative offset of about 1% is well inside the statistical uncertainty of a single measurement and is not corrected. The proton-tagged extractions have not yet been tested this way. The construction of the reference and the full-covariance pulls are described in Sec. 14.6 of the technical supplement.

12.3 Robustness to a different interaction model

Here the fake data are thrown from a reweighted interaction model, while the response, the efficiency and the background prediction remain nominal, so the bias that appears is the combined effect of the regularisation and of a response built from the wrong model. The fake-data universe of the framework is filled from the thrown files at truth and reconstruction level, so the truth used for comparison is that of the alternative model, and the reweighting applies to signal and background alike. The fake data are reconstructed GENIE events reweighted event by event to an alternative parameter set, with the same detector simulation, unfolded with the nominal

response and compared using the nominal covariance. This is a model-distortion test; a test with a different generator and its own detector simulation would be stronger and is pending.

Two variations are used, in FHC and RHC. The first is GENIE multisim universe 545, a simultaneous variation of all interaction parameters, chosen because it changes the shapes more than the rate: the signal rate changes by 4% while the true distributions change by up to 29% in $\cos \theta_\pi$ and 24% in p_π . The second is the $\Delta \rightarrow N\pi$ decay angular unisim, a targeted change of the pion angular distribution in resonance events. Both changes are larger than the differences between the tune and the standalone generators.

Table 18: Unfolding pseudo-data from an alternative interaction model with the nominal response, FHC (upper block) and RHC (lower). “Ratio” is the mean of unfolded over the alternative model’s own truth, over the bins holding at least 10% of the largest truth; “max dev” is the largest per-bin departure from unity among those bins, and “bias” the largest per-bin departure over every bin, expressed in units of the uncertainty that bin carries.

Observable	Alternative model	mean ratio	max dev.	largest bias
p_μ	GENIE multisim u545	0.920	19.6%	-0.47σ
	$\Delta \rightarrow N\pi$ angular	0.992	32.0%	$+0.83\sigma$
$\cos \theta_\mu$	GENIE multisim u545	0.942	43.8%	-0.87σ
	$\Delta \rightarrow N\pi$ angular	0.816	57.4%	-1.18σ
$\cos \theta_\pi$	GENIE multisim u545	0.980	14.7%	$+0.40\sigma$
	$\Delta \rightarrow N\pi$ angular	0.944	22.6%	-0.59σ
p_π (three regions)	GENIE multisim u545	0.750	43.4%	-0.70σ
	$\Delta \rightarrow N\pi$ angular	1.084	16.0%	$+0.34\sigma$
<i>RHC</i>				
p_μ	GENIE multisim u545	0.783	33.6%	-0.83σ
	$\Delta \rightarrow N\pi$ angular	1.027	21.0%	$+0.56\sigma$
$\cos \theta_\mu$	GENIE multisim u545	1.059	19.7%	$+0.78\sigma$
	$\Delta \rightarrow N\pi$ angular	1.045	11.4%	-0.47σ
$\cos \theta_\pi$	GENIE multisim u545	0.852	23.4%	-0.64σ
	$\Delta \rightarrow N\pi$ angular	0.987	7.7%	$+0.20\sigma$
p_π (three regions)	GENIE multisim u545	0.852	19.9%	-0.36σ
	$\Delta \rightarrow N\pi$ angular	1.158	24.6%	$+0.55\sigma$

One bin deviates by more than one standard deviation of its uncertainty: the largest bias is 1.18σ , in the third FHC $\cos \theta_\mu$ bin under the angular variation. In the well-populated bins the extraction follows the alternative truth: the five lower p_μ bins return 0.80–0.96 of it under both variations, with biases below 0.5σ , against truth-level changes of up to 4% in FHC and 9% in RHC.

The bins that do not follow it are sparse ones. The two near-threshold p_π regions return 0.60 and 0.57 of the alternative truth under the multisim variation, which lowers that truth by 20%; the open highest p_μ bin returns 1.32 under the angular variation; and the first and third FHC $\cos \theta_\mu$ bins ($\cos \theta_\mu < 0.4$ and 0.625 – 0.75) return 0.59 and 0.43 under the angular variation and 0.78 and 0.56 under the multisim one. These are the weakest bins of the measurement: the first two because the range-based pion momentum saturates, the others because they lie in the sparse tails or, for the third $\cos \theta_\mu$ bin, in a narrow bin whose A_C row sum is the smallest of the eight. Their biases stay near or below 1σ partly because their uncertainties are large, 106–110%, 29% and 92–114%. At the present exposure the model bias is therefore inside the uncertainties, but it would become significant if the uncertainties were substantially reduced, and the test should be repeated for a larger exposure. In RHC no bin exceeds 0.83σ : the multisim variation returns 0.66–0.88 of the alternative truth in every p_μ bin (-0.3 to -0.8σ) and 0.78–0.85 on average in p_μ , $\cos \theta_\pi$ and p_π , while $\cos \theta_\mu$ comes back 6% high; under the angular variation the RHC ratios are 0.99–1.16. The proton-tagged extractions are more sensitive: their low-imbalance δp_T and p_n

bins return 35% below the alternative truth under the multisim variation, -1.1 to -1.2σ (Sec. 4 of the proton-tagged note).

Independent-generator closure. The one existing NuWro overlay, for Run-1 FHC, is not used for this test. Its normalisation disagrees with the standalone NuWro prediction for the same flux: its true signal rate per POT is 1.76 times that of GENIE, against 1.6 expected from the standalone NuWro cross section and the flux. Its production could not be traced, so a closure test with it would test the sample rather than the extraction. A validated NuWro or GiBUU overlay is required (§14.1).

13 Results

13.1 Result set

Table 19 lists every extraction of the analysis. The primary results are the fifteen inclusive differential extractions, five observables in three configurations, and the three one-bin totals. Each carries the covariance of the configured uncertainty model, detector variations included; the effects outside it are listed in §11.4 and §11.3. The detector term of the inclusive results is 17–27% (34% for FHC p_π). The proton-tagged extractions are reported in the companion note.

Table 19: Every extraction of the analysis, generated from the result index (`report/tools/inventory.py`). The extraction count is one per observable and configuration. The differential $W_{\pi p}$ files remain in the result directory but, like the five-bin p_π , are not reported (Appendix C).

Family	Observables	Extractions	Status
Inclusive differential	5	15	primary
Inclusive one-bin total	1	3	primary
Proton-tagged differential	12	36	secondary
Proton-tagged one-bin total	1	3	secondary
Proton-tagged angular ($\theta_p, \theta_{\pi p}$)	2	6	proposed secondary ¹
Two-dimensional pairs	8	24	two proposed, combined only ¹
<i>Of the above, not reported (Appendix C)</i>			
Differential $W_{\pi p}$	1	3	not reported
Five-bin $d\sigma/dp_\pi$	1	n/a	not reported

¹ Extracted separately from the indexed result set and not listed in its index.

Table 20 gives, for the FHC results, the number of bins, the smallest migration diagonal, the range of the A_C row sums, the trace of A_C , which is the number of degrees of freedom that the data determine, the number of components measured with a signal-to-noise ratio of at least one, and the bin-averaged total uncertainty. A row sum well below one means that the corresponding bin is determined largely by the regularisation.

Table 20: Information content of the primary FHC results. The reported points are the Wiener-SVD estimator $A_C x_{\text{true}}$, not an unsmeared estimate of x_{true} ; an A_C row sum well below one means that bin is set largely by the regularisation. The trace of A_C is the number of degrees of freedom the data determine, and the next column the number of components measured with signal-to-noise of at least one (Wiener filter factors ≥ 0.5); the uncertainty is the bin-averaged total including data statistics. Migration diagonals are the smallest over bins.

Observable	bins	min. diagonal	A_C row sums	trace A_C	factors ≥ 0.5	unc. [%]
p_μ	6	0.57	0.45–0.98	2.09	2	41.2
p_π (three regions)	3	0.51	0.52–0.91	1.57	2	60.0
$\cos \theta_\mu$	8	0.66	0.49–0.70	4.00	4	43.1
$\cos \theta_\pi$	5	0.74	0.68–0.78	3.11	4	41.5
$\theta_{\mu\pi}$	5	0.79	0.52–1.04	2.35	3	38.3

13.2 Differential cross sections

Tables 21–23 summarise the differential results of the three configurations, and Figs. 26–28 show them. All are extracted from fake data, and all recover the A_C -smeared truth of the fake data in every observable.

Table 21: FHC (Runs 1, 2, 4, 5; data POT 7.766×10^{20}): integrated cross section, the two largest systematic terms, the total and the closure against the A_C -smeared truth of the fake data, from the covariance and result files. The flux term is discussed in §11.2.

Observable	σ_{int}	Flux	Detector	Total	Closure χ^2/ndf
p_μ	0.814	30.0%	21.2%	39.3%	0.82/6 ($p = 0.99$)
p_π	0.830	38.5%	34.2%	56.7%	1.01/3 ($p = 0.80$)
$\cos \theta_\mu$	0.713	26.6%	25.4%	39.6%	2.39/8 ($p = 0.97$)
$\cos \theta_\pi$	0.668	27.1%	23.2%	39.0%	1.21/5 ($p = 0.94$)
$\theta_{\mu\pi}$	0.796	28.8%	19.9%	37.1%	0.39/5 ($p = 1.00$)

The RHC flux is 54.6% $\bar{\nu}_\mu$, against 35% in FHC (Table 5). The selected RHC events are nevertheless 76% ν_μ and 20% $\bar{\nu}_\mu$, since the ν_μ cross section is two to three times larger. The RHC overlays (Runs 1, 2, 3, 4a, 4b and 4c) were checked to be ν_μ overlays, with a truth composition of 100% $\nu_\mu + \bar{\nu}_\mu$ and the same number of events per POT in every file (2.62×10^{-16} with the PPF weights); together they contain 4.05×10^6 events and 1.546×10^{22} simulated POT, the largest period being Run 3 (5.52×10^{21} POT, 1.44×10^6 events). The RHC detector uncertainty uses the native Run-4 RHC detector-variation samples (light yield down, Rayleigh scattering, recombination, space charge and four wire-response variations), with 3.80×10^{-16} events per POT in every variation. No RHC dirt sample exists, and the FHC dirt overlay scaled to the RHC exposure is used. The generator predictions are averaged over the RHC flux composition.

The combined measurement uses the summed response of the fourteen overlay files, 5.6×10^6 events, with each horn mode normalised to its own exposure (1.885×10^{21} POT in total). The prediction thus weights the two modes by their recorded exposure, 0.70:1, rather than by their simulated exposure, 0.62:1, and the fake data are thrown at the same rates. The combined generator predictions are the fluence-weighted average of the two modes, 0.425 FHC and 0.575 RHC, each mode weighted by $\Phi_{\text{tot}} \times \text{POT}$.

The flux uncertainty of the combination needs no separately constructed correlation between the horn modes. All fourteen overlay files carry the same 600 PPF multisim universes, and the universes are accumulated index by index across both modes, so universe u is the same hadron-production variation in FHC and RHC. Summing per universe propagates each variation through both configurations and yields the physical correlation between the modes. The detector

variations are carried per mode, with the Run-4 FHC and Run-4 RHC samples of each variation accumulated under one label, so that the combined detector term is weighted by exposure like the central value.

The combined and RHC closures pass in every observable (Tables 22 and 23): the RHC truth is recovered with $\chi^2 \leq 0.98$ on three to eight degrees of freedom and $p \geq 0.99$. Excluding p_π , the unfolded flux term is 22.8–27.9% combined, 26.6–30.0% in FHC and 24.6–29.4% in RHC (31.0%, 38.5% and 27.7% for p_π), as expected for a common flux model propagated through similarly conditioned responses. The reconstructed-level flux term is about 15% in RHC and 17% in FHC, the $\bar{\nu}_\mu$ -dominated flux having a smaller hadron-production uncertainty, but the unfolded terms are similar because the background subtraction multiplies both by $1 + B/S$ (§11.2).

Table 22: RHC (Runs 1, 2, 3, 4; data POT 1.1082×10^{21}): integrated cross section, the two largest systematic terms, the total and the closure against the A_C -smeared truth of the fake data, from the covariance and result files. The flux term is discussed in §11.2.

Observable	σ_{int}	Flux	Detector	Total	Closure χ^2/ndf
p_μ	0.644	29.4%	23.8%	40.6%	0.44/6 ($p = 1.00$)
p_π	0.727	27.7%	24.0%	41.1%	0.01/3 ($p = 1.00$)
$\cos \theta_\mu$	0.694	24.6%	23.5%	36.9%	0.98/8 ($p = 1.00$)
$\cos \theta_\pi$	0.615	25.2%	25.5%	38.5%	0.49/5 ($p = 0.99$)
$\theta_{\mu\pi}$	0.831	25.2%	16.6%	32.6%	0.18/5 ($p = 1.00$)

Table 23: Combined (data POT 1.8848×10^{21}): integrated cross section, the two largest systematic terms, the total and the closure against the A_C -smeared truth of the fake data, from the covariance and result files. The flux term is discussed in §11.2.

Observable	σ_{int}	Flux	Detector	Total	Closure χ^2/ndf
p_μ	0.742	27.9%	23.7%	39.0%	0.08/6 ($p = 1.00$)
p_π	0.762	31.0%	26.9%	45.9%	0.21/3 ($p = 0.98$)
$\cos \theta_\mu$	0.744	22.8%	21.2%	33.1%	1.29/8 ($p = 1.00$)
$\cos \theta_\pi$	0.649	25.7%	23.0%	37.5%	0.58/5 ($p = 0.99$)
$\theta_{\mu\pi}$	0.860	25.8%	18.0%	33.6%	0.21/5 ($p = 1.00$)




Figure 26: Differential cross sections $d\sigma/dx$ of the five distributions, FHC full exposure, extracted from fake data. The points are the Wiener-SVD estimator $A_C x_{\text{true}}$; shown are the unfolded fake data (black points), the A_C -smearred truth of the fake data (grey dashed), the MicroBooNE tune (black solid) and the four generator predictions (GENIE, GiBUU, NEUT, NuWro), all smeared with A_C . The unfolded points recover the injected truth in every observable, the closure summarised in Table 21.




Figure 27: RHC differential cross sections from fake data, with the generator predictions averaged over the RHC flux. Same layout as Fig. 26.




Figure 28: Combined FHC+RHC differential cross sections from fake data normalised per horn mode, with the fluence-weighted combination of the generator predictions. Same layout as Fig. 26.

Table 24: Integrated cross section σ_{int} of each unfolded differential result over the measured phase space ($10^{-38} \text{ cm}^2/\text{Ar}$), from fake data, with the flux of each configuration. The value differs between observables because A_C does not preserve the normalisation; the D’Agostini cross-check (Sec. 6.4 of the technical supplement) gives a total 12–55% higher, independent of the observable to 1% and equal to the one-bin total, and the A_C row sums of §13.4 show the effect directly. The total cross section is the one-bin extraction of Table 28.

Observable	FHC	RHC	Combined
p_μ	0.814	0.644	0.742
p_π	0.830	0.727	0.762
$\cos \theta_\mu$	0.713	0.694	0.744
$\cos \theta_\pi$	0.668	0.615	0.649
$\theta_{\mu\pi}$	0.796	0.831	0.860

Uncertainties shown. The error bars of the unfolded points are the square roots of the diagonal of the total covariance propagated through the unfolding: flux (PPFX multisims), cross section (GENIE multisims and unisims), detector variations [15], hadron re-interaction [25], POT and target number, and the data, simulation and beam-off statistical terms. For FHC p_μ the bin-averaged total is 41.2% (39.3% without the data statistics, Table 17), with flux 30.0%, detector 21.2%, data statistics 12.1% and cross section 11.9%. The one-bin total (§13.6) has its own breakdown from the same universes, systematic $\oplus$ statistical = 38.0% $\oplus$ 4.5% in FHC; without regularisation its flux and detector terms are not redistributed between bins.

13.3 Pion momentum in three regions

The upper pion-momentum bin is open-ended, so the p_π result is reported as three bin-integrated cross sections (Table 25): two intervals near threshold, 0.175 to 0.205 and 0.205 to 0.26 GeV/c, and everything above. These are the quantities to compare with a prediction. The differential form, with a conventional width for the upper bin (Appendix A), is used only to draw the result

together with the other observables.

Table 25: The p_π result as bin integrals (10^{-38} cm²/Ar), from fake data. The upper region is open-ended and has no width. The uncertainties are the total of the covariance. All three bins are A_C -smeared, so a prediction is compared after multiplication by A_C .

Region	FHC	RHC	Combined
$\sigma(0.175 < p_\pi < 0.205 \text{ GeV}/c)$	0.048 ± 0.036	0.064 ± 0.029	0.059 ± 0.032
$\sigma(0.205 < p_\pi < 0.26 \text{ GeV}/c)$	0.086 ± 0.061	0.110 ± 0.049	0.102 ± 0.053
$\sigma(p_\pi > 0.26 \text{ GeV}/c)$	0.695 ± 0.236	0.553 ± 0.222	0.601 ± 0.216

13.4 Comparing a prediction with the results

The Wiener-SVD result estimates $A_C x_{\text{true}}$, and A_C does not preserve the normalisation, so the sum of a differential result over its bins is not the total cross section. The row sums of A_C (Table 26) describe the effect: a row summing to one passes the normalisation of a flat bin-integrated truth unchanged, and a row summing to 0.2 passes one fifth of it. For a non-flat prediction, bin i receives $\sum_j (A_C)_{ij} t_j$, which depends on the individual elements and on the shape of t . For FHC p_μ the ratio of the smeared integral to the true total is 0.78, 0.78, 0.77 and 0.79 for GENIE, GiBUU, NEUT and NuWro. The true totals of the same generators span 0.82 to 1.32×10^{-38} cm²/Ar (Table 19 of the technical supplement), a factor 1.6, so a common factor near 0.78 does not produce the differences between the generators; they are model differences.

Table 26: Range of the row sums of A_C for each observable and configuration, from the result files.

Observable	FHC	RHC	Combined
p_μ	0.45–0.98	0.35–0.72	0.45–0.73
p_π	0.52–0.91	0.45–0.89	0.45–0.89
$\cos \theta_\mu$	0.49–0.70	0.26–0.88	0.47–0.67
$\cos \theta_\pi$	0.68–0.78	0.70–0.84	0.70–0.83
$\theta_{\mu\pi}$	0.52–1.04	0.58–1.15	0.63–1.20

The suppression is large, down to a row sum of 0.26 for $\cos \theta_\mu$ in RHC, and differs strongly between observables. This is why the integrated cross section of a differential result varies between observables while the physical total cannot; the D’Agostini cross-check, whose integral does not depend on the observable, agrees with this interpretation (Sec. 6.4 of the technical supplement). The row sums of the proton-tagged results, and the one matrix whose least-determined row sums to a negative value (RHC $W_{\pi p}$, -0.14), are discussed in Sec. 14.5 of the technical supplement; a negative row is allowed for a regularisation operator, but that row calls for caution.

Numerical results. All fifty-one differential extractions, inclusive and proton-tagged, are provided in `report/data_release/`, indexed by `index_extractions.tsv`: the additional-smearing matrix, the covariance of every component in cross-section units, and `curves_*.tsv` with the unfolded data and every model curve shown here. The six one-bin totals, with their uncertainty breakdown and the realised truth, are in `total_xsec.tsv`. A prediction given as a cross section p_j in true bin j is compared with the results through

$$\left(\frac{d\sigma}{dx}\right)_i^{\text{smeared}} = \frac{1}{w_i} \sum_j (A_C)_{ij} p_j,$$

with w_i the width of bin i . Applied to the true generator predictions, this reproduces the generator curves of the result files to a largest relative difference of 4×10^{-8} per bin, the precision

of the text files, so a user can check an implementation against any column of `curves_*.tsv` before applying it to a new model. The check detects a prediction smeared twice, which would miss the curves by up to 55% and which a closure test cannot see. The physical total cross section is the one-bin extraction, which is not subject to A_C ; the sum of a Wiener-SVD differential result is a quantity in smeared space.

13.5 Model comparisons

Table 27 gives the χ^2 of the unfolded fake data against each prediction, for every observable and configuration. All predictions are smeared with A_C , so the comparison is made in the space of the measurement, with the full covariance. The truth column is the closure test.

Every prediction is smeared with A_C exactly once, by the extraction, and the figures use the smeared curves without smearing again. The prescription above, applied to the true predictions, reproduces the smeared curves to a relative 4×10^{-8} (Sec. 15 of the technical supplement). A second application would over-smooth the sharply peaked generators and create a spurious separation in the muon-angle observables. With a single application NEUT and NuWro give 3.93/8 and 4.88/8 in FHC $\cos \theta_\mu$, in line with the other observables.

Table 27: χ^2/nfd of the unfolded fake data against each A_C -smeared prediction, for the inclusive selection. “truth” is the fake-data truth (the closure test).

Beam	Observable	truth	uB tune	GENIE	GiBUU	NEUT	NuWro
FHC	p_μ	0.82/6	1.52/6	2.73/6	1.47/6	2.65/6	1.41/6
	p_π	1.01/3	1.20/3	1.56/3	4.13/3	1.83/3	2.26/3
	$\cos \theta_\mu$	2.39/8	1.75/8	2.37/8	2.57/8	3.93/8	4.88/8
	$\cos \theta_\pi$	1.21/5	1.21/5	1.92/5	2.56/5	2.26/5	3.09/5
	$\theta_{\mu\pi}$	0.39/5	0.47/5	0.82/5	0.67/5	2.71/5	3.16/5
RHC	p_μ	0.44/6	0.76/6	0.83/6	0.57/6	0.90/6	1.25/6
	p_π	0.01/3	0.01/3	0.58/3	0.71/3	0.26/3	0.27/3
	$\cos \theta_\mu$	0.98/8	1.27/8	1.84/8	0.71/8	6.36/8	10.39/8
	$\cos \theta_\pi$	0.49/5	0.43/5	1.27/5	2.86/5	1.79/5	1.08/5
	$\theta_{\mu\pi}$	0.19/5	0.34/5	2.52/5	3.54/5	0.46/5	1.08/5
comb	p_μ	0.08/6	0.23/6	1.36/6	0.31/6	1.04/6	0.63/6
	p_π	0.21/3	0.26/3	0.89/3	2.66/3	0.70/3	1.19/3
	$\cos \theta_\mu$	1.29/8	1.17/8	2.17/8	1.30/8	5.24/8	9.52/8
	$\cos \theta_\pi$	0.58/5	0.52/5	1.65/5	2.70/5	1.63/5	2.06/5
	$\theta_{\mu\pi}$	0.21/5	0.38/5	1.69/5	2.21/5	1.93/5	2.88/5

13.6 Total cross section

The total flux-averaged cross section in the measured phase space is extracted with the same machinery as a one-bin distribution: a single true bin that contains every signal event, defined on $\cos \theta_\mu$ over its full range $[-1, 1]$ so that no overflow exists, and a single reconstructed bin over the same range. With one bin the response $S = \epsilon M$ reduces to the efficiency, nothing is regularised, the Wiener filter is switched off, and the unfolding is the inversion

$$\sigma = \frac{N_{\text{sel}} - N_{\text{bkg}}}{\epsilon \Phi N_{\text{Ar}}}, \quad (16)$$

with $A_C \equiv 1$ exactly. Here N_{sel} is the selected count, N_{bkg} the predicted background (neutrino background, beam-off and dirt), $\epsilon = N_{\text{sel}}^{\text{sig}}/N_{\text{gen}}^{\text{sig}}$ the selection efficiency in the measured phase space (17.5% in every configuration), Φ the integrated flux ($\Phi = \text{POT} \times \Phi_{\text{POT}}$, Table 5) and $N_{\text{Ar}} = 8.71 \times 10^{29}$ the number of argon nuclei in the fiducial volume, the same normalisation

$N_{\text{Ar}} \Phi / 10^{-38}$ as for the differential results. Every term of Eq. (16) is evaluated universe by universe, so the covariance contains the flux, cross-section, re-interaction and detector variations, the POT and target terms and the simulation, beam-off and data statistical terms, propagated as for the differential results. The flux term is 28–30%, the flux normalisation times $(1 + B/S)$ with B the neutrino background (§11.2); the detector term is 20–22% and the data-statistical term 3–5%. The differential results carry the same terms bin by bin, with the regularisation redistributing them between bins. Without additional smearing the one-bin result is the physical total. It lies above the A_C -suppressed σ_{int} of Table 24, most visibly in RHC, and agrees with the D’Agostini cross-check (Sec. 6.4 of the technical supplement).

A closed-form evaluation of Eq. (16) on the central-value prediction checks the normalisation chain: it reproduces the one-bin extraction of the central-value prediction to four digits (1.0315, 0.9285 and 0.9723 inclusive; 0.6269, 0.5506 and 0.5831 proton-tagged) and provides the N_{sel} , N_{bkg} and ϵ columns of Table 28. The counts carry the full central-value weight ($\text{tune} \times \text{PPFX} \times \text{normalisation}$, §4); the PPFX weight averages 0.944 on selected signal, so a count without it would lie 3–6% above the prediction used in the unfolding.

Table 28 gives the total cross sections of the inclusive selection and of the proton-tagged subsample. The σ column is the one-bin extraction from fake data and fluctuates around the tune by the Poisson throw; the realised truth of the throw is listed, and the ratio to it is 0.952–0.998 with $\chi^2/1 \ll 1$ in all six cases. The uncertainty is that of the one-bin extraction itself. The ratio RHC/FHC of the tune is 0.900 (0.878 proton-tagged); the extracted ratio from the fake data is 0.901 (0.873). Since the extraction does not depend on an observable, this is also the proton-tagged total to quote; the integrals of the proton-tagged differential results (Table 1 of the proton-tagged note–3) are A_C -suppressed and depend on the observable.

Table 28: Headline total flux-averaged cross sections ($10^{-38} \text{ cm}^2/\text{Ar}$) in the measured phase space, from the one-bin extraction (§13.6) on the Poisson-thrown central-value fake data. N_{sel} and N_{bkg} are the predicted selected count and background (ν -background + EXT + dirt) of the central-value prediction and ϵ the selection efficiency, from the cut-and-count check; σ is the one-bin extraction with its full universe-propagated uncertainty, “stat” its data-statistical term, “syst” the prediction total (flux, cross section, re-interaction, detector, POT, targets, MC and EXT statistics) and “total” their quadrature sum; “truth” is the realised truth of the throw and $\chi^2/1$ the closure against it. MC counts carry the framework’s central-value weight ($\text{tune} \times \text{PPFX} \times \text{normalisation}$), the same as Table 10; the two differ by $\lesssim 1\%$ because the yield comparison uses the Run-1 dirt overlay with per-mode scale factors rather than the per-run dirt files.

Config	N_{sel}	N_{bkg}	ϵ [%]	σ	stat [%]	syst [%]	total [%]	truth	σ/truth	$\chi^2/1$
<i>Inclusive CC $1\mu 1\pi X p$</i>										
FHC	1 417	584	17.5	1.034 ± 0.396	4.5	38.0	38.3	1.044	0.991	0.001
RHC	1 759	744	17.6	0.932 ± 0.350	4.1	37.4	37.6	0.941	0.990	0.001
Combined	3 176	1 328	17.5	0.975 ± 0.366	3.0	37.5	37.6	0.985	0.990	0.001
<i>Proton-tagged CC $1\mu 1\pi 1p$ ($p_p > 0.3 \text{ GeV}/c$)</i>										
FHC	637	305	11.5	0.608 ± 0.256	7.8	41.3	42.0	0.639	0.952	0.014
RHC	786	395	11.4	0.531 ± 0.251	7.4	46.6	47.2	0.532	0.998	0.000
Combined	1 423	700	11.5	0.564 ± 0.239	5.4	42.1	42.5	0.578	0.976	0.003

The flux-averaged generator totals in the inclusive phase space are 0.82 (GENIE), 1.13 (GiBUU), 1.32 (NEUT) and $1.25 \times 10^{-38} \text{ cm}^2/\text{Ar}$ (NuWro) for FHC (Table 19 of the technical supplement for all three configurations). The one-bin total lies $0.5\text{--}0.8\sigma$ above GENIE and within 0.8σ of the other three; the spread of the generators, a factor 1.6, is larger than any of these differences. In the proton-tagged phase space the FHC generator totals are 0.50 (GENIE), 0.65 (NuWro), 0.74 (GiBUU) and 0.78 (NEUT), around the injected 0.63.

13.7 Proton-tagged subsample

The proton-tagged subsample, the inclusive selection with a reconstructed proton candidate above $0.3\text{ GeV}/c$, is measured separately [53]: a one-bin total, the proton momentum, the calorimetric hadronic mass, four transverse-kinematic-imbalance observables, two proton angles and two-dimensional pairs. It has its own signal definition, proton-assignment requirements and validation, and it is not part of the approval request. Its one-bin totals, from the same fake data, are 0.608 ± 0.256 (FHC), 0.531 ± 0.251 (RHC) and $0.564 \pm 0.239 \times 10^{-38} \text{ cm}^2/\text{Ar}$ (combined).

14 Limitations and requirements for unblinding

Blind signal region. The signal region has not been examined in beam-on data, and every result is extracted from fake data. The results therefore demonstrate the extraction and its uncertainties, not the cross section. The background model enters strongly at 58–59% purity, with a proton reconstructed as the pion as the largest neutrino background, and closure tests cannot probe it because the same simulation provides the fake data and the background estimate. It has been tested at reconstruction level in the four control regions (§9.2), which require no correction; the offset of FHC Run 1 relative to the other run periods is not yet explained.

Pion momentum. Charged pions interact before they stop, so the reconstructed momentum saturates near $0.34\text{ GeV}/c$ and carries little information above about $0.4\text{ GeV}/c$. Golden-pion tagging, multiple Coulomb scattering, calorimetry, multivariate regression, recovery of the daughter energy and the proton-tagged selection were all studied, and a truth-level golden selection performs no better than the trained classifier, which places the limit in the physics of pion interactions rather than in the estimator. For this reason p_π is reported in three regions.

Model dependence. The closure tests use fake data from the model of the response. The test with reweighted interaction models bounds the model bias at 1.18σ (§12.3), but a test with an independent generator and its own detector simulation has not been made. The model dependence of the efficiency matters most for the one-bin total, a single efficiency averaged over the phase space, and for the open-ended $p_\pi > 0.26\text{ GeV}/c$ region, inside which the efficiency varies strongly with the true momentum.

Ensemble offset. The ensembles place the mean result 0.4–1.1% below the reference, at 0.5–1.4 standard deviations of the mean (0.8% at 2.0 for the RHC total), with the same sign in all six extractions (§12.2). No correction or uncertainty is assigned; the offset is well inside the statistical uncertainty of a single measurement.

Muon momentum scale. The MCS momentum-scale term (§11.3) is provided as a separate covariance for p_μ and is not yet part of the covariance computed by the framework. Its size, $\pm 5\%$, rests on an assumption that a data-driven comparison of MCS and range momenta on contained muons should confirm.

Two totals. The Wiener-SVD differential result estimates $A_C x_{\text{true}}$, with row sums as low as 0.26, so the sum of a differential result is not the physical total cross section. The physical total is the one-bin extraction.

Proton-tagged results. The limitations specific to the proton-tagged subsample (the proton mis-assignment, which no control region yet constrains, and the transverse-kinematic prescription) are discussed in the companion note [53].

14.1 Requirements before the signal region is opened

The prerequisites are listed as a checklist in Sec. 12 of the technical supplement. The main ones are:

1. written confirmation from the NuMI flux contact that the integrated flux constants of Table 5 are PPFX-corrected central values (§3.4);
2. an understanding of the 13% offset of FHC Run 1 relative to the other run periods in the control regions: its exposure bookkeeping, the completeness of its beam-on sample and the normalisation of its overlay (§9.2);
3. a closure test with a validated independent-generator overlay (§12.3);
4. the data-driven check of the $\pm 5\%$ MCS scale and the inclusion of the term in the framework covariance (§11.3);
5. the beam-on POT and gate counts of the unblinded files, checked against the run lists of those files, in place of the fake-data values; the Run-1 beam-off sample remains the Run-2 substitute (§4);
6. a final reproduction of the result set from empty caches.

The control regions were compared with beam-on data under the frozen procedure of Sec. 13.5 of the technical supplement, and any correction to the background model follows that procedure. No switch is needed for the beam-off treatment at unblinding: the framework recognises fake data by the presence of simulated truth, and it refuses a beam-on file of the signal region unless an explicit unblinding switch is set. Five automated checks protect the result set against regressions: the normalisation manifest, the assertion coupling the systematics configuration to the file lists, the staleness check, the table-consistency check and the blinding guard (Sec. 15 of the technical supplement).

15 Summary

We have developed a measurement of $\nu_\mu + \bar{\nu}_\mu$ charged-current single charged-pion production on argon with the full MicroBooNE NuMI exposure: forward horn current (7.766×10^{20} POT), reverse horn current (1.1082×10^{21} POT) and their combination (1.8848×10^{21} POT). The signal region is blind, and all results are extracted from fake data thrown from the central-value simulation; no statement about a generator is yet a statement about nature.

Phase space and selection. The signal is a ν_μ or $\bar{\nu}_\mu$ charged-current interaction with its vertex in the fiducial volume and exactly one muon and one charged pion, no neutral pion or kaon and any number of nucleons, with $p_\mu > 0.15 \text{ GeV}/c$, $p_\pi > 0.175 \text{ GeV}/c$ and $\theta_{\mu\pi} < 2.6 \text{ rad}$. The selection reaches 17.5% efficiency at 58–59% purity. Angles are measured about the neutrino direction of each event, 28.6° from the detector axis.

Results. In each configuration the total flux-averaged cross section and five differential distributions are measured: $d\sigma/dp_\mu$, $d\sigma/d\cos\theta_\mu$, $d\sigma/d\cos\theta_\pi$, $d\sigma/d\theta_{\mu\pi}$, and p_π in three regions. The binnings meet a 0.50 migration criterion (App. A). The total cross section carries a 38% uncertainty, with flux 28–30% and detector 20–22%; the flux term is the 17% flux normalisation multiplied by $1 + B/S$ through the background subtraction. The differential results carry 30–46% bin-averaged uncertainty without the data statistics (67% for FHC p_π). Each differential result is provided with its additional-smearing matrix A_C , which must be applied to a prediction before comparison, and with its covariance by source.

Background. The four control regions, disjoint from the signal region, have been compared with beam-on data. The global test passes ($\chi^2 = 8.1/8$), no class-specific correction of the background is supported, and the data lie 12–20% above the prediction, within the flux and cross-section uncertainties, in every run period except FHC Run 1, whose offset of 13% relative to the others is under investigation.

Validation. Closure recovers the input in every extraction, with a mean ratio of 0.996 over the fifteen inclusive extractions. Ensembles of 100 pseudo-experiments in six extractions give pull widths of 0.92–1.07 and correct coverage in all three configurations. Fake data from reweighted interaction models bias no bin by more than 1.18σ . The identity $UR = A_C$ holds to 6.1×10^{-5} (App. 9 of the technical supplement), and the A_C row sums show that the difference between the integrals of the differential results is the additional smearing. The result set is rebuilt from the processed ntuples with empty caches in a single pass; an independent rebuild into a separate tree reproduced every bin of the fifty-one extractions to 10^{-9} (Sec. 15 of the technical supplement), and the generator curves are reproduced from the provided A_C to a relative 4×10^{-8} , the precision of the text files.

Limitations. The saturation of the range-based pion momentum near 0.34 GeV/ c limits p_π to three regions and is the main physics limitation of the selection. Before the signal region is opened, the flux normalisation must be confirmed, the FHC Run-1 offset understood and a closure test with an independent generator made (§14.1).

Glossary

Signal definition The set of particles leaving the nucleus that counts as the process being measured, with momentum thresholds and a fiducial requirement. It makes no statement about the interaction mechanism, because final-state interactions mix mechanisms and only the final state is observable.

Fake data / pseudo-data A simulated dataset used in place of beam-on data while the analysis is blind. Here it is a Poisson throw: each simulated event is replicated a Poisson-distributed number of times with mean equal to its weight at data exposure, so the sample has the statistical fluctuations of a real dataset.

Blind The signal region has not been examined in beam-on data. Every plotted “data” point of the signal region is fake data.

Beam-on, beam-off (EXT), dirt Beam-on data are recorded in the beam spill; beam-off data are recorded between spills and measure cosmic-ray activity; “dirt” is simulated neutrino interactions outside the cryostat whose products enter the detector. The prediction is neutrino simulation + beam-off + dirt.

Gate ratio The factor scaling beam-off data to the beam-on exposure: the ratio of the number of beam-on to beam-off trigger gates in a running period.

Tune, central value (CV) The MicroBooNE tune is the collaboration’s adjusted GENIE model; the central value is the nominal prediction from it. “CV weight” is the per-event weight that turns the raw simulation into the tuned prediction.

Safe-weight rule Any event weight that is non-finite, negative or above 30 is replaced by 1. One rule, applied everywhere, so that the prediction and every pseudo-data throw agree.

Universe One alternative version of the prediction with a systematic parameter varied. A *multisim* is a set of hundreds of universes drawn randomly from the parameter uncertainties;

a *unisim* is a single $\pm 1\sigma$ variation of one parameter. PPFX universes vary the flux; GENIE universes the interaction model; reinteraction universes the hadron transport.

Detector variations (detVar) Dedicated simulation samples with one aspect of the detector response changed (light yield, recombination, space charge or wire response), used to estimate the detector uncertainty.

Migration matrix, efficiency, smearceptance The migration matrix M_{ij} is the probability that an event in true bin j is reconstructed in bin i ; the efficiency ε_j is the fraction of true bin j that is selected at all. Their product $S_{ij} = \varepsilon_j M_{ij}$ is the response the unfolding inverts; we call it the smearceptance matrix.

Migration diagonal The fraction of the selected events of a true bin that are reconstructed in the same bin (the diagonal of the column-normalised M). A binning is accepted only if every bin's diagonal exceeds 0.68.

Unfolding, Wiener-SVD Solving for the true spectrum given the reconstructed one and the response. A plain inversion amplifies statistical noise; Wiener-SVD regularises it with a filter built from the simulation, trading a small bias for a large reduction in variance.

Additional-smearing matrix A_C The price of regularisation, made explicit: the unfolded result estimates $A_C T$ rather than the truth T . A prediction must be multiplied by A_C before it is compared with the result. A_C is released with every extraction.

Not normalisation-preserving The rows of A_C do not sum to one, so the integral of an unfolded differential result is not the total cross section and differs between observables of the same sample. The total is measured separately.

One-bin extraction The total cross section, extracted with a single true bin covering the whole phase space. Nothing is regularised, $A_C = 1$, and the result is the physical total.

Conditioning, $s_{\min}/s_1$, rank-deficient The ratio of the smallest to the largest singular value of the regularised response. Near zero, some direction in the true spectrum cannot be recovered from the data at all (rank-deficient); the result in that direction is the prior. A binning can pass the migration-diagonal criterion and still be unusable here.

Closure Unfolding pseudo-data and checking that the truth comes back. Closure against the A_C -smeared truth tests the implementation; it cannot test whether the model is right, because the same model is on both sides.

Covariance The matrix of uncertainties and their correlations between bins, one per systematic source, propagated through the unfolding. “Nominal configured covariance” means the sum over the sources implemented in the analysis; the sources known to be missing are listed in the limitations.

Flux-averaged cross section The cross section averaged over the neutrino energy spectrum of the beam, per argon nucleus, in the stated phase space. It depends on the flux, which is why each horn configuration is a separate measurement.

FHC, RHC Forward and reverse horn current: the two polarities of the NuMI focusing horns, giving a neutrino-dominated and an antineutrino-enriched beam.

Beam frame All angles are measured about the neutrino direction, which at MicroBooNE's off-axis position is 28° from the detector z axis.

Control region, sideband A sample selected by inverting one requirement of the signal selection, so that it is disjoint from the signal region and enriched in one background. It can be inspected on beam-on data while the signal region stays blind.

Transverse kinematic imbalance (TKI) Variables built from the momenta transverse to the neutrino direction of the lepton and the hadronic system: δp_T , $\delta \alpha_T$, $\delta \phi_T$ and the inferred nucleon momentum p_n . They isolate nuclear effects from the unknown neutrino energy.

LLR particle identification The log-likelihood-ratio score built from track dE/dx that separates protons from muons and pions.

A Observable binning

Table 29 is generated from the `curves_*.tsv` result files, so the edges are those of the results. The binnings are the same in the three configurations and meet the 0.50 migration criterion in FHC and RHC. They were selected in the binning review of Sec. 5 of the technical supplement: among the binnings that meet the criterion and hold at least 70 selected FHC signal events in every true bin (50 in the proton-tagged family), the one with the fewest bins that reaches 90% of the largest trace of A_C , the number of degrees of freedom that the data determine. $\cos \theta_\pi$ and $\theta_{\mu\pi}$ keep their earlier five bins, which pass both criteria (the seven and six bins the rule selects would add 0.45 and 0.1 degrees of freedom), $\delta \alpha_T$ keeps its split at 120° , and the proton-tagged p_π keeps two regions, because three fail the criterion in that sample. The generator predictions are computed at these edges from the event-level files.

Table 29: Bin edges of every observable, from the result files. The binnings meet the 0.50 migration criterion (Sec. 5 of the technical supplement) and are the same in the three configurations.

Observable	Bins	Edges
<i>Inclusive</i>		
p_μ (GeV/c)	6	0.15, 0.3, 0.425, 0.575, 0.825, 1.25, 3*
p_π (GeV/c)	3	0.175, 0.205, 0.26, 1*
$\cos \theta_\mu$	8	-1, 0.4, 0.625, 0.75, 0.825, 0.9, 0.95, 0.975, 1
$\cos \theta_\pi$	5	-1, -0.1, 0.35, 0.6, 0.8, 1
$\theta_{\mu\pi}$ (rad)	5	0, 0.6, 0.85, 1.3, 1.85, 2.6
<i>Proton-tagged</i>		
p_μ (GeV/c)	6	0.15, 0.3, 0.425, 0.575, 0.825, 1.25, 3*
p_π (GeV/c)	2	0.175, 0.205, 1*
$\cos \theta_\mu$	8	-1, 0.4, 0.625, 0.75, 0.825, 0.9, 0.95, 0.975, 1
$\cos \theta_\pi$	4	-1, -0.1, 0.35, 0.75, 1
$\theta_{\mu\pi}$ (rad)	5	0, 0.6, 0.85, 1.3, 1.85, 2.6
$W_{\pi p}$ (GeV/c ²)	6	1.08, 1.19, 1.23, 1.27, 1.34, 1.47, 2.9*
W_{had} (GeV/c ²)	2	0, 1.18, 2.74*
δp_T (GeV/c)	2	0, 0.3, 2.5*
$\delta \phi_T$ (deg)	3	0, 22.5, 75, 180
$\delta \alpha_T$ (deg)	2	0, 120, 180
p_n (GeV/c)	2	0, 0.375, 2*
p_p (GeV/c)	3	0.3, 0.42, 0.6, 1.5*
θ_p (rad) [†]	4	0, 0.44, 0.785, 1.162, π
$\theta_{\pi p}$ (rad) [†]	4	0, 0.911, 1.414, 1.916, π

[†] Proposed secondary scope: the proton polar angle and the pion–proton opening angle (Sec. 3.3 of the proton-tagged note).

* Numerical plotting boundary only: the upper truth bin is open-ended, and a density drawn over the conventional width represents the bin integral.

The upper p_π region is defined at truth level as $p_\pi > 0.26$ GeV/c with no upper edge. An open-ended bin has no width and hence no differential density: the 1 GeV/c in Table 29 is a

convention that assigns a width of 0.74 GeV/ c so that the region can be drawn on the same axis as the other two, not a kinematic limit of the sample. A height read from that figure reflects the convention. The three region integrals are unambiguous, and the result is quoted and compared as those integrals (Table 25). The other open-ended bins of the table follow the same convention.

B Alternative binnings

The binnings used until 2026-09-26 came from the optimisation of Sec. 4 of the technical supplement. The review on the corrected exposure (Sec. 5 of the technical supplement) found that the seven-bin p_μ failed both the 0.68 and the 0.50 criterion (worst diagonal 0.33), as did the six-bin W_{had} (0.09), and that the five-bin $\cos \theta_\mu$ and θ_μ failed 0.68 narrowly (0.67 and 0.65). Under the 0.68 criterion the same rule selects four p_μ bins, two p_π regions, seven $\cos \theta_\mu$ bins, two W_{had} regions, two p_p bins and no $\delta\alpha_T$ shape; Tables 8 and 9 of the technical supplement compare the two criteria observable by observable. Table 30 lists the alternatives examined before the review, with the conditioning s_{min}/s_1 on which they were judged then; the review replaces that measure, which most binnings do not meet either.

Table 30: Alternative binnings examined before the review of Sec. 5 of the technical supplement, as judged then, and their status after it. s_{min}/s_1 is the conditioning of A_C (FHC/RHC unless stated), which is set by the weakest component of the regularised response alone.

Observable	Alternative	s_{min}/s_1	Status
$\cos \theta_\mu$	6 bins	0.11/0.09	superseded by the eight bins of the review
	7 bins	0.02/0.05	superseded by the eight bins of the review
θ_μ	7, 8 bins	0.00/0.00	θ_μ no longer reported
p_π	3 bins	0.05/0.01	three regions adopted in the inclusive family
δp_T	3 bins	0.00/0.00	not adopted: lower trace than two bins
$\delta\alpha_T$	3 bins	0.03/0.01	not adopted: the third bin adds nothing
$W_{\pi p}$	2 bins	0.26/0.15	sound response, no shape (App. C)
W_{had}	2 bins	0.58/0.54	adopted
$\cos \theta_p$	4 bins	0.45/0.02	superseded by θ_p
	6 bins	0.00/0.00	superseded by θ_p
θ_p	5, 7 bins	≤ 0.03	superseded by four bins at maximin edges
	3 bins (RHC edges)	0.12/0.05; COMB 0.49	superseded by four bins at maximin edges
$\cos \theta_\pi$	4 bins	0.27/0.24	superseded by the five-bin scheme

The two-bin W_{had} had been set aside because two bins discard most of the shape for which W_{had} is measured. The six bins used until the review did not measure that shape either: every one failed 0.68 and five failed 0.50, and two regions split at 1.18 GeV/ c^2 are the only W_{had} binning that meets either criterion. The earlier four-bin θ_p held 21 selected events in its last bin, below the minimum of 50 for the proton-tagged family; the maximin edges now used meet it.

C Distributions not reported

Three distributions were extracted and examined but are not reported, two because their responses do not support them and one because it duplicates another. Their files remain in the result directory and are not to be quoted.

C.1 The differential $W_{\pi p}$ shape

The pion–proton invariant mass would show the $\Delta(1232)$ directly and motivated the proton-tagged subsample. At the present exposure its shape cannot be measured with any binning tried. The six-bin scheme that resolves the resonance region has a highest bin with a 2% diagonal, so nearly every event reconstructed there comes from another true bin. Two bins restore the response but leave no shape: the best split, at $1.20 \text{ GeV}/c^2$, passes the 0.50 criterion by 0.003, and its second filter factor is 0.29 in FHC and 0.27 in RHC. Closure against the A_C -smeared truth does not help: it tests whether the response was inverted correctly and cannot restore information the response does not contain. The proton-tagged sample supports the integrated cross section and the calorimetric hadronic mass W_{had} in two regions split at $1.18 \text{ GeV}/c^2$, the reported binning.

C.2 The five-bin $d\sigma/dp_\pi$

Charged pions interact before they stop, so the reconstructed momentum saturates near $0.34 \text{ GeV}/c$ and the migration matrix loses its diagonal above that value: four of five bins fall below the 0.68 criterion, the highest reaching 6%. The pion momentum is therefore reported in three regions, two near threshold and one broad inclusive region above $0.26 \text{ GeV}/c$. The limitation is the estimator, since the range is a valid momentum estimate only for a pion that stops, and no rebinning of a saturated variable recovers the information.

C.3 The muon polar angle θ_μ

$d\sigma/d\theta_\mu$ was reported alongside $d\sigma/d\cos\theta_\mu$ until 2026-09-26. A partition of the muon direction has the same migration diagonals in either variable, but the regularisation acts on the density per unit of the variable, and the cosine retains up to one degree of freedom more for the same partition (Sec. 5 of the technical supplement). The muon angle is therefore reported as $\cos\theta_\mu$ only; the earlier comparison of the two parameterisations is recorded in Sec. 8.1 of the technical supplement.

D Selection cut parameters

Table 31: All selection threshold constants.

Parameter	Value	Ntuple variable
Topological score	> 0.67	Pandora score
Muon track score	> 0.80	Pandora track score
Muon vertex distance	≤ 4.0 cm	
Muon track length	> 10.0 cm	
Muon LLR PID	> 0.20	LLR PID
Muon MIP BDT	≥ -0.10	TMVA BDT
Pion track length	> 20.0 cm	
Pion vertex distance	< 4.0 cm	
Pion LLR PID	> 0.10	LLR PID
Pion MIP BDT	> -0.10	TMVA BDT
Pion pion-BDT	> -0.10	TMVA BDT
Pion Bragg score	≥ 0.08	Bragg likelihood
Max. opening angle	< 2.6 rad	
Shower Y-plane hits	< 50 (if 1 shower)	
Shower vertex distance	> 10 cm (if 1 shower)	
Non-proton multiplicity	< 4	LLR PID > 0.1
Primary track multiplicity	< 5	
Cosmic BDT	> 0.0 (disabled)	TMVA BDT
Software trigger	$= 1$	
<i>Proton-tagged selection only (Sec. 2 of the proton-tagged note)</i>		
Leading-proton LLR PID	< 0.05	most proton-like non- μ/π primary
Proton end point	in FV	
Proton momentum	≥ 0.3 GeV/c	proton-hypothesis range KE

E The MicroBooNE detector

E.1 Operating Principle

A liquid argon time projection chamber records three-dimensional ionisation images of charged-particle tracks by exploiting two properties of ultra-pure liquid argon:

1. Ionisation electrons produced by a traversing charged particle are not captured by argon atoms and can drift centimetres or metres under an applied electric field without significant recombination, provided the purity is sufficient (O_2 -equivalent $\lesssim 100$ ppt).
2. Argon scintillates at 127 nm (vacuum ultraviolet), producing $\sim 24\,000$ photons per MeV of deposited energy, enabling nanosecond-scale timing.

An applied electric field drifts the ionisation electrons to a set of sense-wire planes at the anode, which record three two-dimensional projections of each track. The third coordinate (along the drift direction) is reconstructed from the known drift velocity and the timing of the signals relative to the beam-gate start, combining the two-dimensional wire images into a full 3D reconstruction.

MicroBooNE operated from October 2015 to July 2021 at the Fermilab Accelerator Laboratory, collecting data in both the Booster Neutrino Beam (BNB) and, as a parasitic experiment, the NuMI beam [19].

E.2 Active Volume and Electric Field

The MicroBooNE active volume has rectangular geometry: 256 cm (drift, x) $\times$ 233 cm (height, y) $\times$ 1037 cm (beam, z), containing ≈ 87 tonnes of liquid argon at ≈ 87 K. The cathode is a stainless-steel plane at $x = 256$ cm held at -70 kV; the anode wire planes are at $x = 0$. The resulting 273 V/cm drift field gives an electron drift velocity of ≈ 1.11 mm/ μ s, so the maximum drift time across the full gap is approximately 2.3 ms.

Liquid argon purity is maintained by continuous recirculation through oxygen and water absorbers, keeping the electron lifetime well above 5 ms (corresponding to O₂-equivalent contamination below ~ 100 ppt). At this purity level, charge attenuation over the full 2.56 m drift is less than 10%.

E.3 Wire Readout and Signal Formation

Three wire planes at the anode sample the drifting charge:

- **U plane** (first induction): 2400 wires at $+60^\circ$ from the vertical, 3 mm pitch.
- **V plane** (second induction): 2400 wires at -60° from the vertical, 3 mm pitch.
- **Y plane** (collection): 3456 vertical wires, 3 mm pitch.

Induction-plane wires are biased to be transparent to the drifting electrons, producing bipolar (derivative-like) induced signals; collection-plane wires absorb the charge and produce unipolar signals. The combination of two stereo induction views at $\pm 60^\circ$ and a vertical collection view provides sufficient 2D information in each plane to resolve ambiguities during 3D reconstruction.

Front-end electronics (application-specific integrated circuits, ASICs) are mounted inside the cryostat on the wire planes and cold-multiply each channel at ≈ 87 K, reducing thermal noise. The analogue signals are digitised at 16 MHz inside the cryostat by 12-bit ADCs and transmitted over ~ 110 m of cold cables to warm readout electronics. The data stream is sampled into a circular buffer; beam triggers and cosmic triggers select 4.8 ms readout windows (corresponding to $\approx$ two full drift windows) for permanent storage.

A dead-wire region at $z \in [675.1, 775.1]$ cm interrupts the collection plane over a 100 cm span due to a shorted wire bundle. Tracks crossing this region have degraded hit-counting and range-based momentum estimation; the fiducial volume (Section 6) excludes this region.

E.4 Photon Detection System

Thirty-two 8-inch Hamamatsu R5912-MOD photomultiplier tubes (PMTs) are installed behind the collection-plane wire frame facing the active volume. The PMTs detect the 127 nm argon scintillation light via a tetraphenyl-butadiene (TPB) wavelength-shifting coating that converts VUV photons to visible (~ 430 nm) light. The system serves two purposes in this analysis:

1. **Beam-gate synchronisation:** the PMT flash time identifies which 1.6 μ s NuMI spill window produced the neutrino interaction.
2. **Flash-matching for cosmic rejection:** the Pandora reconstruction associates PFParticle hypotheses with the PMT flash closest in time and position, suppressing out-of-time cosmic-ray muons that overlap with beam-induced events in the long 2.3 ms drift window.

E.5 Calibration and Corrections

Three independent calibration inputs are used:

- **UV laser system** [11]: two pulsed laser beams produce straight ionisation tracks at precisely known locations. Deviations of the reconstructed track positions from the known trajectories are used to build a three-dimensional space-charge distortion map [10], which is applied to all reconstructed hits as a position correction. Distortions reach ~ 10 cm near the anode and cathode edges.

- **Stopping cosmic-ray muons:** the Bragg peak of cosmic muons stopping inside the active volume calibrates the absolute energy scale and the effective recombination factor as a function of local electric field and drift distance. Through-going muons calibrate relative wire-by-wire gains and the electron lifetime.
- **Michel electrons:** $\mu^- \rightarrow e^- \bar{\nu}_e \nu_\mu$ decay products from stopped muons provide a check of the calorimetric reconstruction at $E_e \lesssim 55$ MeV.

E.6 Cosmic Ray Environment and Reconstruction

Operating at the surface with minimal overburden, MicroBooNE is exposed to a cosmic-ray muon rate of ≈ 5 kHz through the active volume. Each 4.8 ms readout window contains on average ~ 10 cosmic-ray tracks unrelated to the beam interaction. The $1.6 \mu\text{s}$ NuMI spill itself contains $\lesssim 0.1$ expected cosmics on average, but the drift-window pile-up is the dominant background at early selection stages (Section 7).

Cosmic rejection proceeds in stages:

1. Flash-matching (PMT timing): candidate neutrino slices must have a consistent optical flash within the beam window.
2. Topological score: the Pandora-based topological neutrino score discriminates between neutrino-like vertex structures and through-going/stopping cosmic topologies.
3. Dedicated BDT scores: the MIP and cosmic scores trained on track-level and event-level features further suppress residual cosmics that passed the topological cut.

The CRT (Cosmic Ray Tagger) sub-detector [34], consisting of scintillator panels surrounding the cryostat, provides an additional cosmic veto but is not used as a cut in the present analysis: the CRT branches are carried through the event processing as pass-through quantities only, they are populated only for the CRT-era data (Runs 3–5) and zero for the pre-CRT Run-1 exposure, and a cut would require a data/MC tagging systematic; the residual beam-off background is subtracted instead (§4).

Reconstruction uses the Pandora [6, 12] multi-algorithm pattern recognition framework, which groups hits into PFParticles organised in a hierarchy with the neutrino interaction as the top-level particle. Track candidates are fitted with a Kalman-filter-based track fitter to obtain momentum from track curvature (negligible for < 300 MeV/ c muons at 273 V/cm) and from range for contained tracks. Log-likelihood ratio particle identification (LLR PID) [13] using the dE/dx profile along the track is the primary tool for muon/pion/proton separation (Section 7).

References

- [1] MicroBooNE Collaboration, “Measurement of the flux-averaged differential ν_e CC1 π cross section in the NuMI beam,” Internal Note v3.2 (2025).
- [2] K. Mistry, “First Measurement of the Flux-Averaged Differential Charged-Current Electron-Neutrino and Antineutrino Cross Section on Argon with the MicroBooNE Detector,” PhD thesis (2021).
- [3] NOvA Collaboration, M. A. Acero *et al.*, “New constraints on oscillation parameters from ν_e appearance and ν_μ disappearance in the NOvA experiment,” *Phys. Rev. D* **98**, 032012 (2018).
- [4] T2K Collaboration, K. Abe *et al.*, “Constraint on the matter–antimatter symmetry-violating phase in neutrino oscillations,” *Nature* **580**, 339–344 (2020).
- [5] DUNE Collaboration, B. Abi *et al.*, “Deep Underground Neutrino Experiment (DUNE), Far Detector Technical Design Report, Volume II: DUNE Physics,” arXiv:2002.03005 [hep-ex] (2020).

- [6] J. S. Marshall and M. A. Thomson, “The Pandora Software Development Kit for Pattern Recognition,” *Eur. Phys. J. C* **75**, 439 (2015).
- [7] W. Tang, X. Li, X. Qian, H. Wei, and C. Zhang, “Data Unfolding with Wiener-SVD Method,” *JINST* **12**, P10002 (2017), arXiv:1705.03568 [physics.data-an].
- [8] L. Aliaga *et al.* (MINERvA Collaboration), “Neutrino Flux Predictions for the NuMI Beam,” *Phys. Rev. D* **94**, 092005 (2016).
- [9] C. Andreopoulos *et al.*, “The GENIE Neutrino Monte Carlo Generator,” *Nucl. Instrum. Methods A* **614**, 87–104 (2010).
- [10] MicroBooNE Collaboration, P. Abratenko *et al.*, “Measurement of space charge effects in the MicroBooNE LArTPC using cosmic muons,” *JINST* **15**, P12037 (2020).
- [11] MicroBooNE Collaboration, C. Adams *et al.*, “A Method to Determine the Electric Field of Liquid Argon Time Projection Chambers Using a UV Laser System and Its Application in MicroBooNE,” *JINST* **15**, P07010 (2020).
- [12] MicroBooNE Collaboration, R. Acciarri *et al.*, “The Pandora multi-algorithm approach to automated pattern recognition of cosmic-ray muon and neutrino events in the MicroBooNE detector,” *Eur. Phys. J. C* **78**, 82 (2018).
- [13] MicroBooNE Collaboration, P. Abratenko *et al.*, “Calorimetric classification of track-like signatures in liquid argon TPCs using MicroBooNE data,” *JHEP* **12**, 153 (2021).
- [14] MicroBooNE Collaboration, P. Abratenko *et al.*, “Determination of muon momentum in the MicroBooNE LArTPC using an improved model of multiple Coulomb scattering,” *JINST* **12**, P10010 (2017).
- [15] MicroBooNE Collaboration, P. Abratenko *et al.*, “Novel approach for evaluating detector-related uncertainties in a LArTPC using MicroBooNE data,” *Eur. Phys. J. C* **82**, 454 (2022).
- [16] MicroBooNE Collaboration, P. Abratenko *et al.*, “New CC0 π GENIE model tune for MicroBooNE,” *Phys. Rev. D* **105**, 072001 (2022), arXiv:2110.14028 [hep-ex].
- [17] MicroBooNE Collaboration, P. Abratenko *et al.*, “First double-differential measurement of kinematic imbalance in neutrino interactions with the MicroBooNE detector,” *Phys. Rev. Lett.* **131**, 101802 (2023), arXiv:2301.03700 [hep-ex].
- [18] S. Gardiner *et al.*, “XSecAnalyzer: a cross-section analysis framework for MicroBooNE,” https://github.com/uboone/xsec_analyzer (accessed 2026).
- [19] P. Adamson *et al.*, “The NuMI Neutrino Beam,” *Nucl. Instrum. Methods A* **806**, 279–306 (2016).
- [20] L. Alvarez-Ruso *et al.* (GENIE Collaboration), “Recent highlights from GENIE v3,” *Eur. Phys. J. ST* **230**, 4449–4467 (2021).
- [21] T. Golan, J. T. Sobczyk, and J. Żmuda, “NuWro: the Wrocław Monte Carlo generator of neutrino interactions,” *Nucl. Phys. B Proc. Suppl.* **229–232**, 499 (2012).
- [22] O. Buss *et al.*, “Transport-theoretical description of nuclear reactions,” *Phys. Rept.* **512**, 1–124 (2012); U. Mosel, “Neutrino event generators: foundation, status and future,” *J. Phys. G* **46**, 113001 (2019).

- [23] Y. Hayato and L. Pickering, “The NEUT neutrino interaction simulation program library,” *Eur. Phys. J. ST* **230**, 4469–4481 (2021).
- [24] S. Agostinelli *et al.* (GEANT4 Collaboration), “GEANT4: a simulation toolkit,” *Nucl. Instrum. Methods A* **506**, 250–303 (2003).
- [25] J. Calcutt, C. Thorpe, K. Mahn, and L. Fields, “Geant4Reweight: a framework for evaluating and propagating hadronic interaction uncertainties in Geant4,” *JINST* **16**, P08042 (2021).
- [26] A. Hoecker *et al.*, “TMVA: Toolkit for Multivariate Data Analysis,” PoS **ACAT**, 040 (2007), arXiv:physics/0703039.
- [27] G. D’Agostini, “A multidimensional unfolding method based on Bayes’ theorem,” *Nucl. Instrum. Methods A* **362**, 487–498 (1995).
- [28] X.-G. Lu, L. Pickering, S. Dolan, G. Barr, D. Coplowe, Y. Uchida, D. Wark, M. O. Wascko, A. Weber, and T. Yuan, “Measurement of nuclear effects in neutrino interactions with minimal dependence on neutrino energy,” *Phys. Rev. C* **94**, 015503 (2016).
- [29] A. P. Furmanski and J. T. Sobczyk, “Neutrino energy reconstruction from one-muon and one-proton events,” *Phys. Rev. C* **95**, 065501 (2017).
- [30] D. Rein and L. M. Sehgal, “Coherent π^0 production in neutrino reactions,” *Nucl. Phys. B* **223**, 29–44 (1983).
- [31] C. Berger and L. M. Sehgal, “Partially conserved axial vector current and coherent pion production by low energy neutrinos,” *Phys. Rev. D* **79**, 053003 (2009).
- [32] MicroBooNE Collaboration, P. Abratenko *et al.*, “First Measurement of Energy-Dependent Inclusive Muon-Neutrino Charged-Current Cross Sections on Argon with the MicroBooNE Detector,” *Phys. Rev. Lett.* **128**, 151801 (2022), arXiv:2110.14023 [hep-ex].
- [33] MicroBooNE Collaboration, R. Acciarri *et al.*, “Design and Construction of the MicroBooNE Detector,” *JINST* **12**, P02017 (2017).
- [34] MicroBooNE Collaboration, C. Adams *et al.*, “Design and construction of the MicroBooNE Cosmic Ray Tagger system,” *JINST* **14**, P04004 (2019).
- [35] D. Rein and L. M. Sehgal, “Neutrino Excitation of Baryon Resonances and Single Pion Production,” *Ann. Phys.* **133**, 79–153 (1981).
- [36] K. S. Kuzmin, V. V. Lyubushkin, and V. A. Naumov, “Lepton polarization in neutrino nucleon interactions,” *Mod. Phys. Lett. A* **19**, 2815–2829 (2004).
- [37] K. M. Graczyk and J. T. Sobczyk, “Form factors in the quark resonance model,” *Phys. Rev. D* **77**, 053001 (2008) [Erratum: *Phys. Rev. D* **79**, 079903 (2009)].
- [38] C. Berger and L. Sehgal, “Lepton mass effects in single pion production by neutrinos,” *Phys. Rev. D* **76**, 113004 (2007).
- [39] J. A. Nowak (MiniBooNE Collaboration), “Four-momentum transfer discrepancy in the charged current π^+ production in the MiniBooNE: Data vs. theory,” *AIP Conf. Proc.* **1189**, 243–248 (2009).
- [40] Y. Hayato, “A neutrino interaction simulation program library NEUT,” *Acta Phys. Polon. B* **40**, 2477–2489 (2009).

- [41] G. M. Radecky *et al.* (ANL Collaboration), “Study of single pion production by weak charged currents in low-energy νd interactions,” *Phys. Rev. D* **25**, 1161–1173 (1982).
- [42] T. Kitagaki *et al.* (BNL Collaboration), “Charged-current exclusive pion production in neutrino–deuterium interactions,” *Phys. Rev. D* **34**, 2554–2565 (1986).
- [43] A. Rodriguez *et al.* (K2K Collaboration), “Measurement of single charged pion production in the charged-current interactions of neutrinos in a 1.3 GeV wide band beam,” *Phys. Rev. D* **78**, 032003 (2008).
- [44] A. A. Aguilar-Arevalo *et al.* (MiniBooNE Collaboration), “Measurement of ν_μ -induced charged-current single positive pion production cross sections on mineral oil at $E_\nu \in 0.5$ –2.0 GeV,” *Phys. Rev. D* **83**, 052007 (2011).
- [45] X.-G. Lu *et al.* (MINERvA Collaboration), “Measurement of final-state correlations in neutrino muon-proton mesonless production on hydrocarbon at $\langle E_\nu \rangle = 3$ GeV,” *Phys. Rev. Lett.* **121**, 022504 (2018).
- [46] B. Eberly *et al.* (MINERvA Collaboration), “Charged pion production in ν_μ interactions on hydrocarbon at $\langle E_\nu \rangle = 4.0$ GeV,” *Phys. Rev. D* **92**, 092008 (2015).
- [47] C. L. McGivern *et al.* (MINERvA Collaboration), “Cross sections for ν_μ and $\bar{\nu}_\mu$ induced pion production on hydrocarbon in the few-GeV region using MINERvA,” *Phys. Rev. D* **94**, 052005 (2016).
- [48] MINERvA Collaboration, J. Wolcott *et al.*, “Evidence for Quasielastic Charge Current Antineutrino Scattering Off Protons in Hydrocarbon at $\langle E_{\bar{\nu}} \rangle = 3.5$ GeV,” *Phys. Rev. Lett.* **116**, 081802 (2016).
- [49] K. Abe *et al.* (T2K Collaboration), “Measurement of the muon neutrino charged-current single π^+ production on hydrocarbon using the T2K off-axis near detector ND280,” *Phys. Rev. D* **101**, 012007 (2020), arXiv:1909.03936 [hep-ex].
- [50] R. Acciarri *et al.* (ArgoNEUT Collaboration), “First measurement of the cross section for ν_μ and $\bar{\nu}_\mu$ induced single charged pion production on argon using ArgoNEUT,” *Phys. Rev. D* **98**, 052002 (2018), arXiv:1804.10294 [hep-ex].
- [51] MicroBooNE Collaboration, P. Abratenko *et al.*, “First measurement of ν_e and $\bar{\nu}_e$ charged-current single charged-pion production differential cross sections on argon using the MicroBooNE detector,” *Phys. Rev. Lett.* **135**, 061802 (2025), arXiv:2503.23384 [hep-ex].
- [52] MicroBooNE Collaboration, P. Abratenko *et al.*, “Measurement of single charged pion production in charged-current ν_μ -Ar interactions with the MicroBooNE detector,” *Phys. Rev. D* **113**, 032007 (2026), arXiv:2509.03628 [hep-ex].
- [53] I. Pophale and J. Nowak, “Measurement of proton-tagged charged-current single-charged-pion production on argon in the NuMI beam with MicroBooNE,” companion analysis note (secondary analysis under validation).